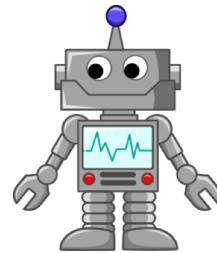


Sistemi Digitali M

07 – Embedded computer vision



Introduction

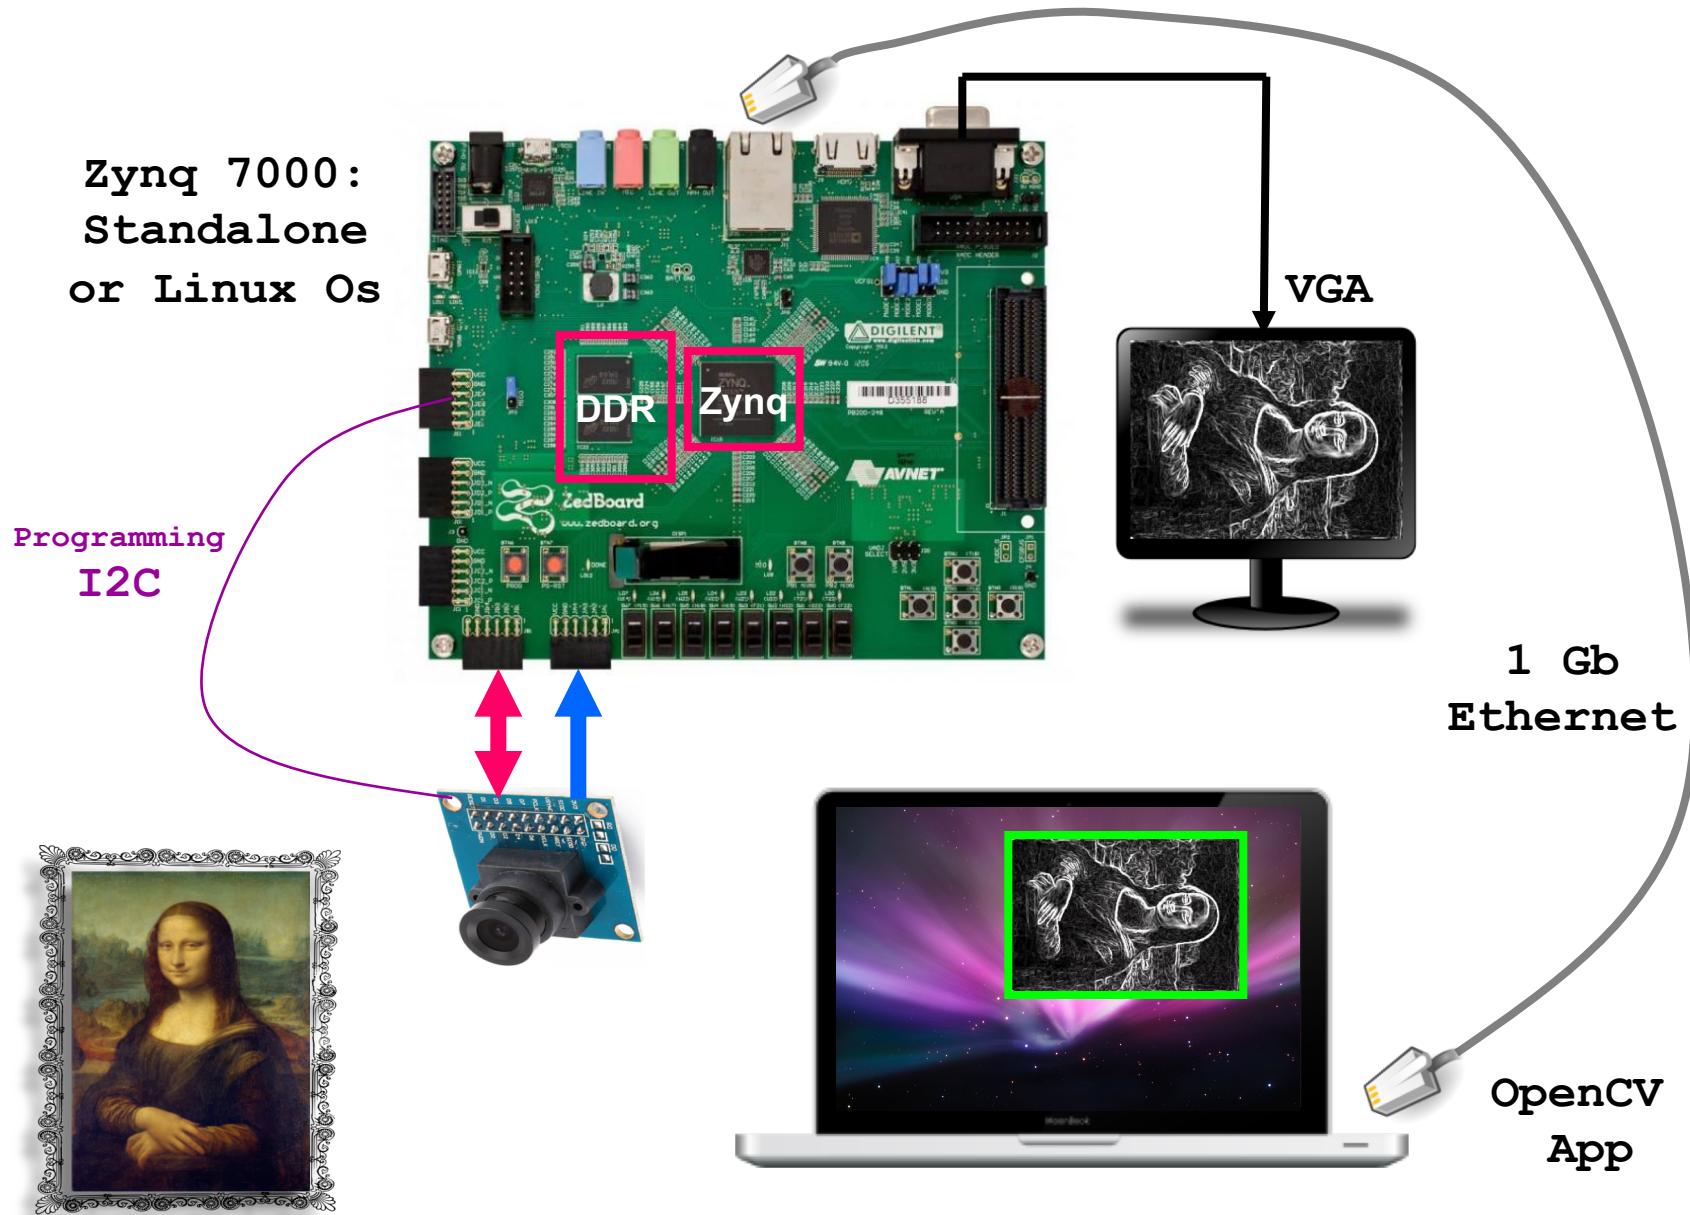
In many application domains, such as those involving autonomous navigation of vehicles or drones, it is crucial to infer cues (eg, depth or flow) from images deploying embedded devices with a limited power budget

We will describe next a whole computer vision (CV) pipeline, implemented from scratch and totally with high-level synthesis tools, on a Zynq evaluation board

Specifically, we will introduce:

- Imaging device technologies
- A simple, yet effective, CV processing system
- Strategies to implement CV modules on FPGA
- OpenCV app for remote image visualization

Overview 1/3



Overview 2/3

- Whole processing pipeline mapped on the FPGA
- Image sensor OV7670
- Circular frame buffer (DDR memory)
- VGA output
- Optional ethernet streaming, interrupt-driven, from ARM to remote host with OpenCV visualization

The whole project is available on Github at this link:

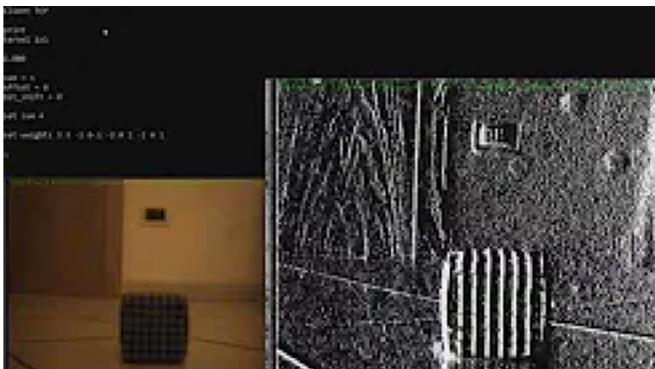
<https://github.com/smatt-github/SmartCamera>

Overview 3/3

Videos of the overall system



<https://www.youtube.com/watch?v=6HEgdZEHpsM>



<https://www.youtube.com/watch?v=QFxjXWgBBcc>

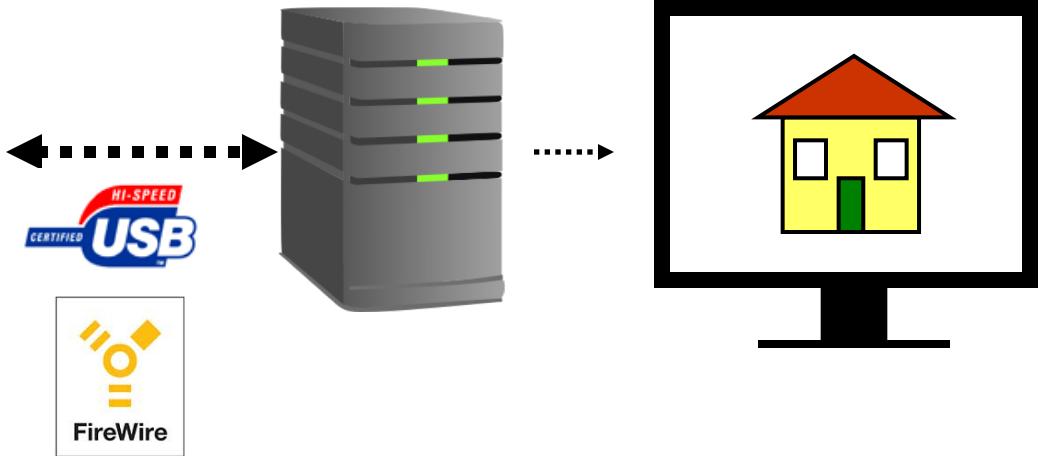
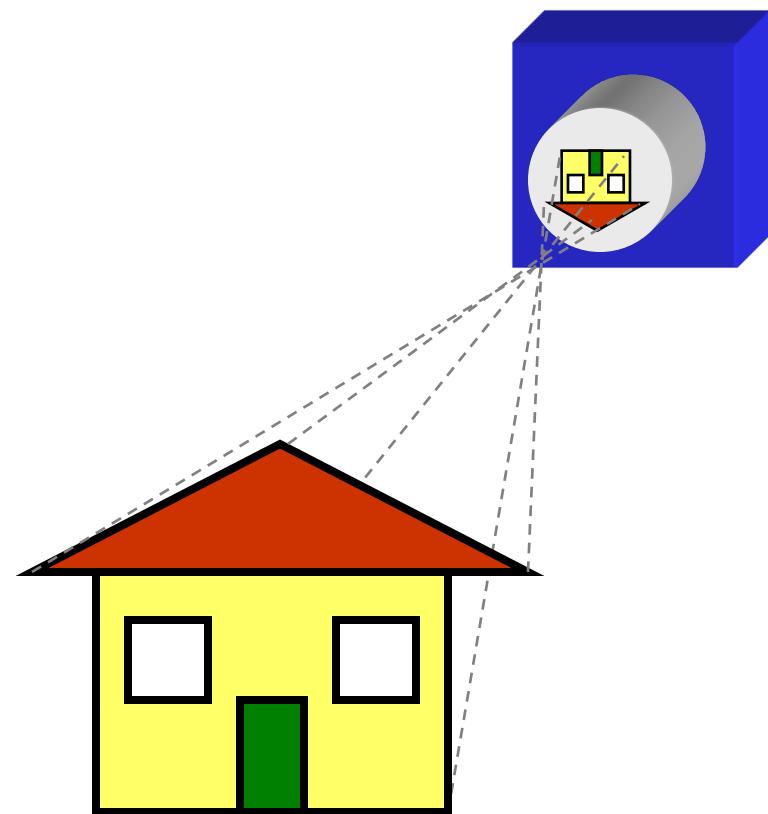


<https://www.youtube.com/watch?v=EG3NYqMJvZI>

Imaging sensors

Consumer cameras, webcams, industrial cameras, etc

??



GIGE™
VISION

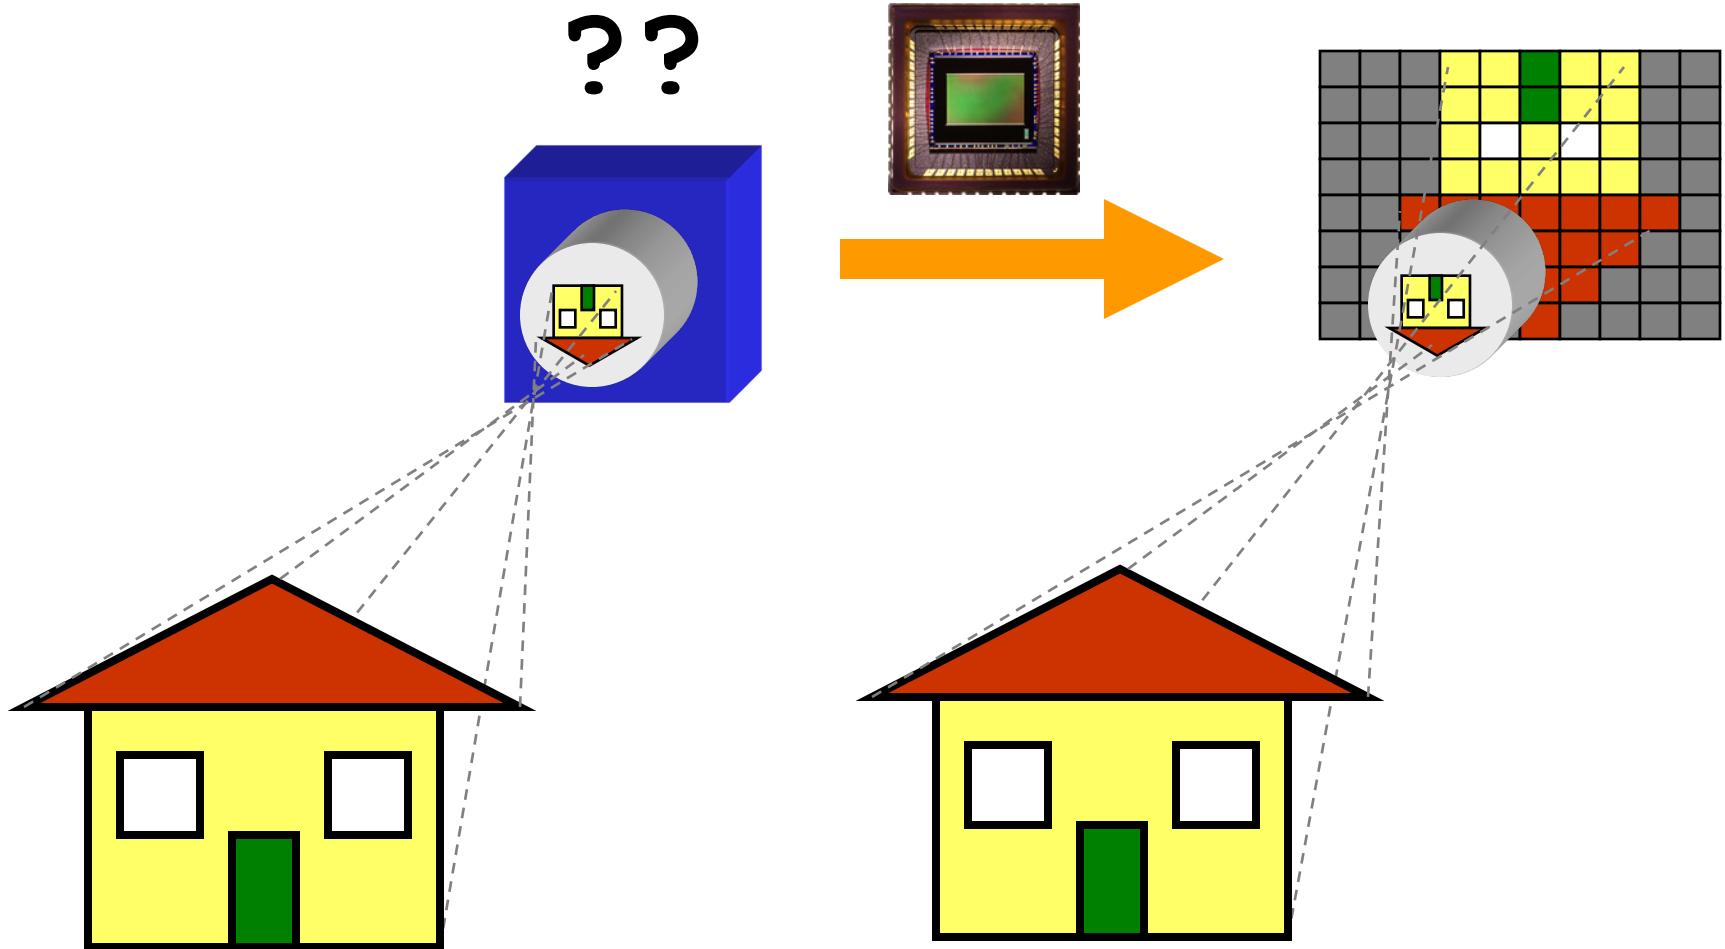


CAMERA
Link

- Conventional optics
- Digital signals

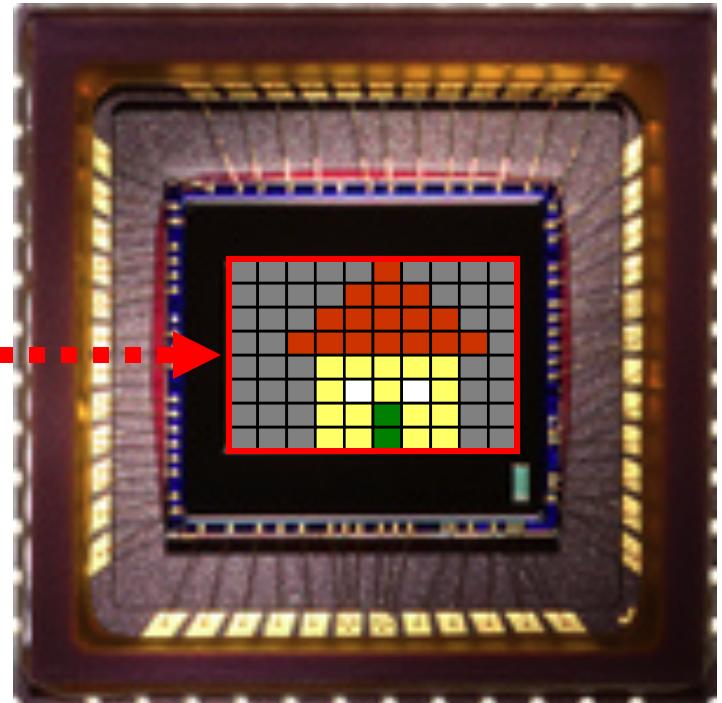
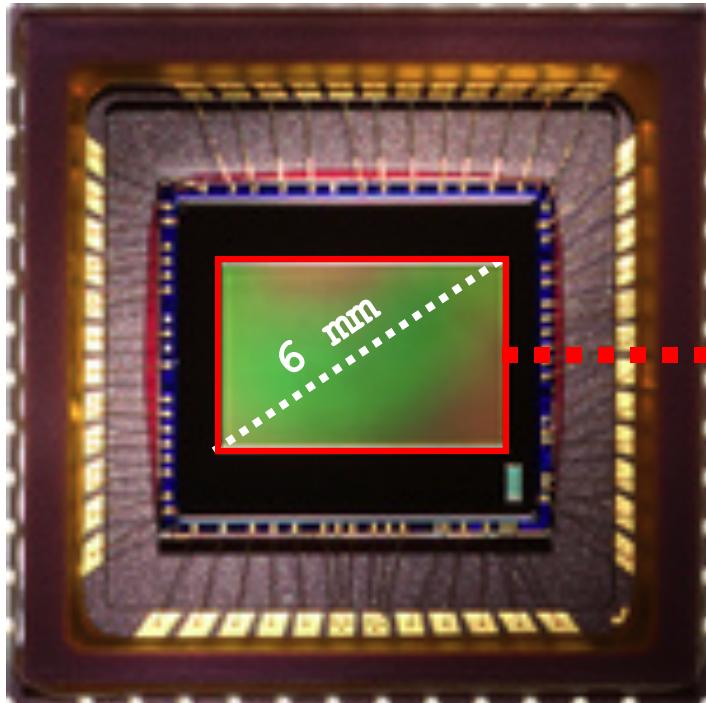
Acquisition process

Array of photosensitive elements (pixels)



Imaging sensor

Example: Aptina/ON global-shutter 1/3" (i.e. diagonal is about ≈ 6 mm) .



Technologies 1/2

There are two main technologies: **CMOS** (mainstream) and **CCD** (for specific application domains)

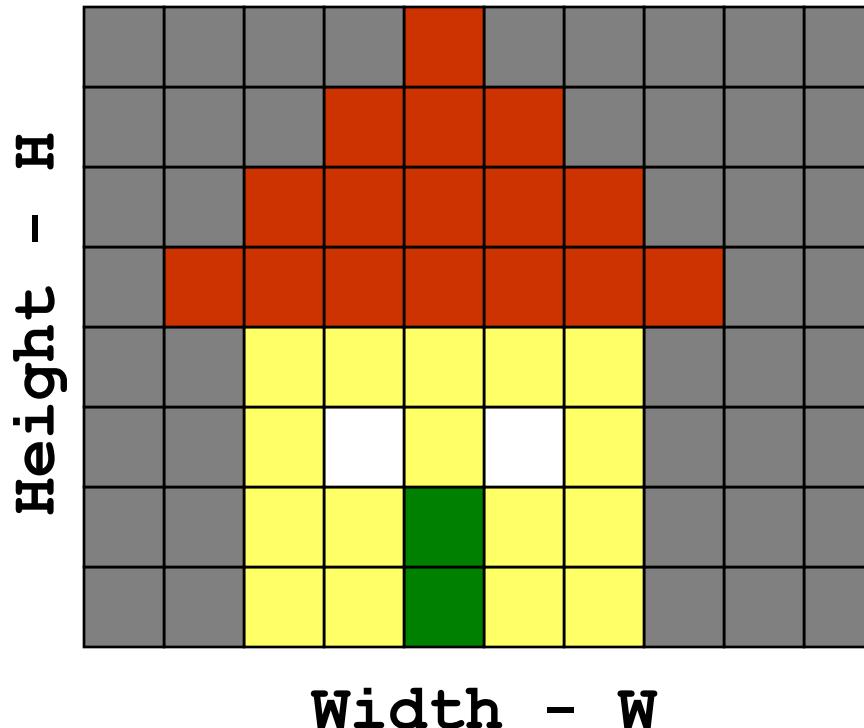


Image resolution: from 0.x MP to several MP

Widely adopted resolutions:

VGA : 640x480 (0.3 MP)

Wide VGA : 752x480 (0.35 MP)

Technologies 2/2

- *rolling shutter*
- *global shutter*

rolling shutter



www.ptgrey.com/

global shutter



www.ptgrey.com/

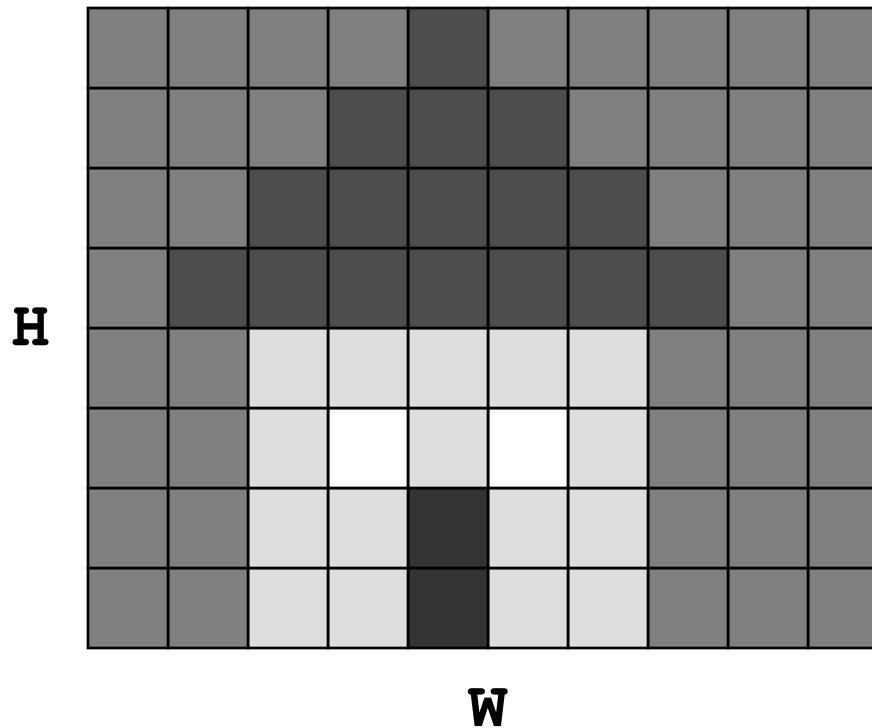
Global shutter is better...



<http://scorpionvision.co.uk/>

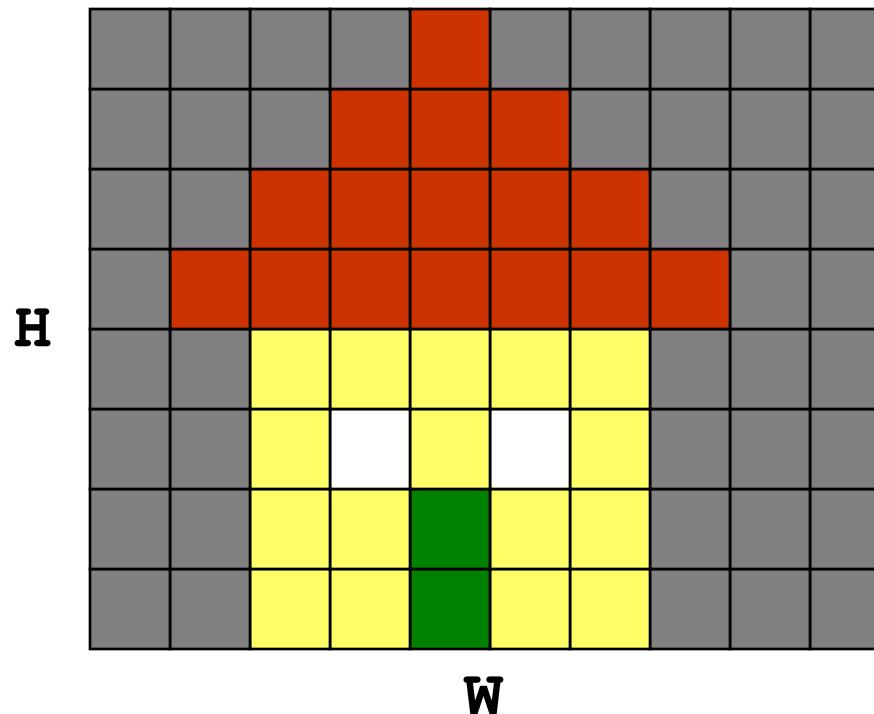
Grayscale (monochrome)

Every pixel is encoded with 8 or 10 bit (in most cases 8 bit). With 8 bit, the range is [0,255].



Color 1/4

Similarly to grayscale, most color sensors encode each pixel with 8 (or 10 bit) or 16 bit. There are several technologies/strategies but in most cases each pixel is encoded, as for greyscale, with 1 byte or 2 bytes.



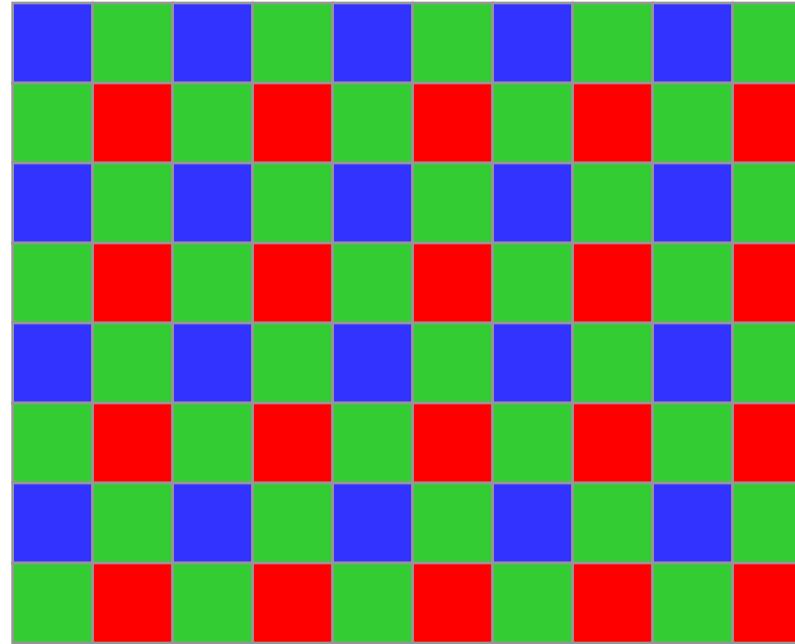
How to encode/obtain colors from 8 bit?

- Bayer pattern
- YUV/YCbCr (later)

Color 2/4 (Bayer 8 bit)

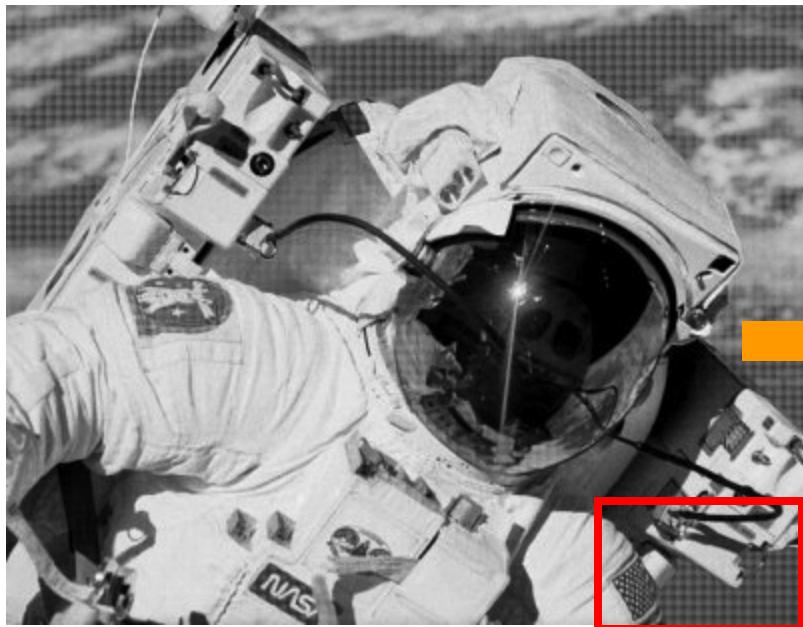
There is a filter for each pixel: **Blue**, **Red** and **Green** according to the followinng scheme (Bayer pattern)

Bayer
pattern



For each pixel, 2 out of 3 color channels are obtained by mean of interpolation. This fact implies a reduced spatial resolution with respect to grayscale. It is worth to note that, with this encoding scheme, the bandwidth is 1/3 (wrt to RGB encoding).

Color 3/4 (Bayer 8 bit)

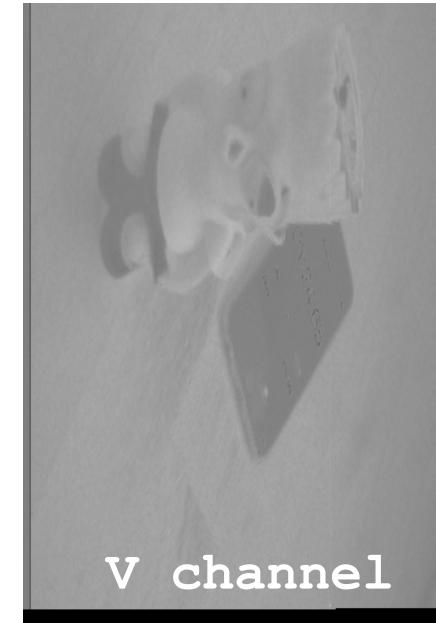


Example: *raw 8 bit image* with Bayer pattern (top left)

Color RGB image (top right)

Detail of the raw bayer pattern (on the left)

Color 4/4 (YUV or YCbCr 16 bit)



RGB image: more details in the next slides

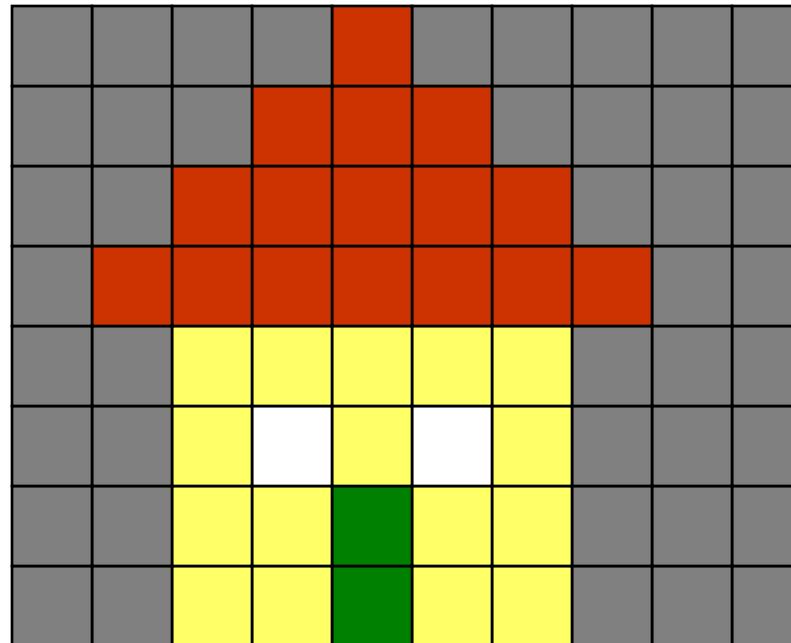


Frame rate 1/2

The pixels are transmitted in *raster scan order* (from top-left to bottom right).

Streaming, at `PIXEL_CLOCK`, can't be paused!

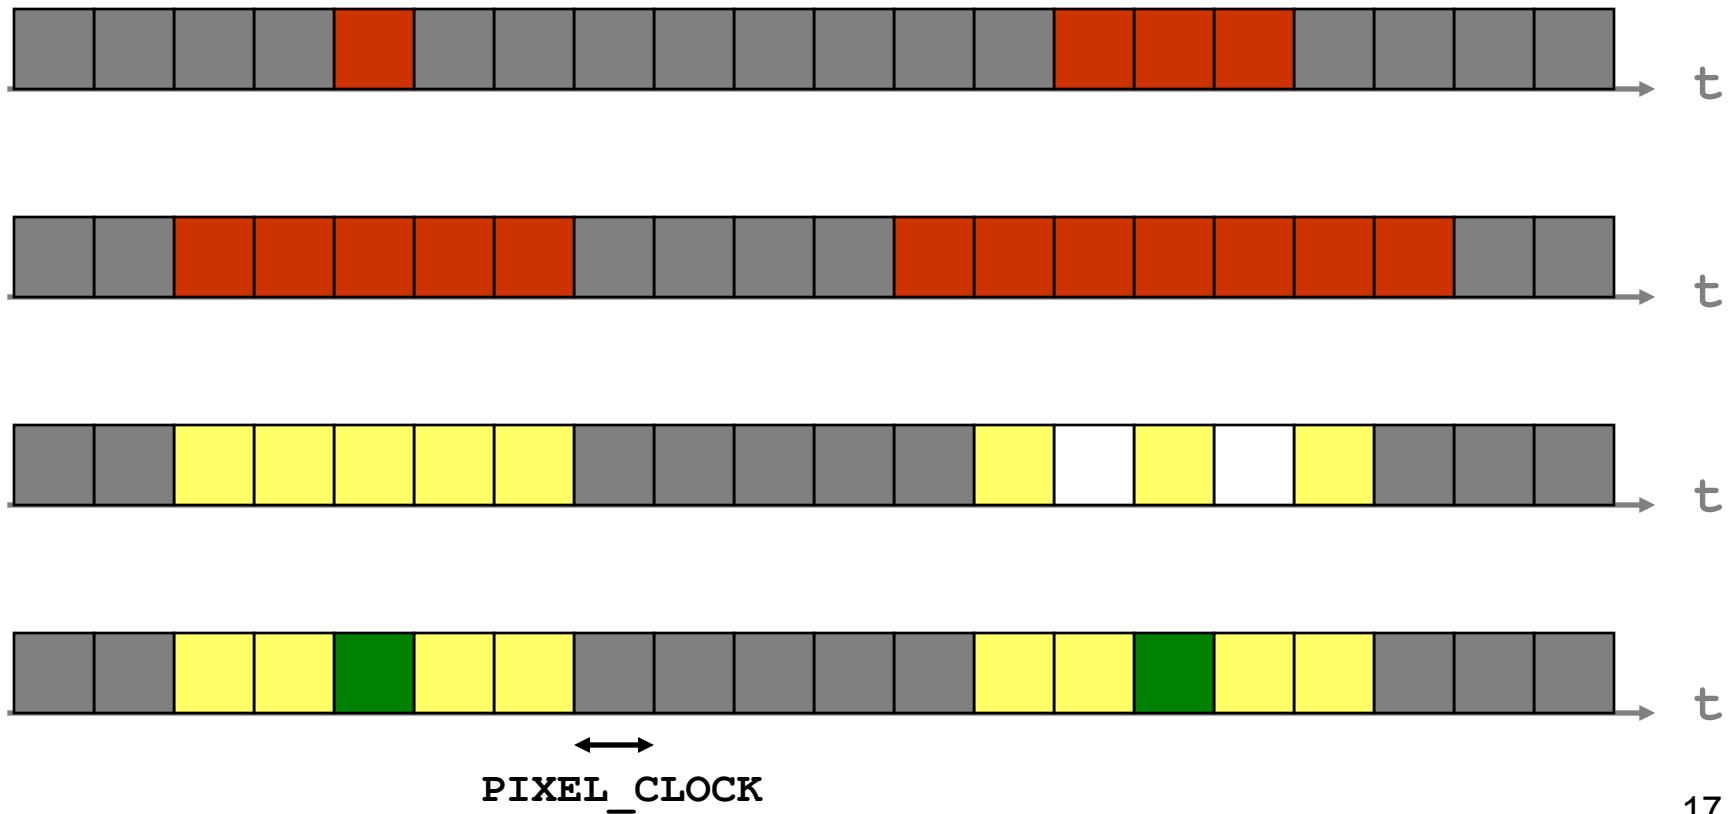
frame per second (fps): images/second



Frame rate 2/2

In most cases pixel are processed at `PIXEL_CLOCK`

With the same frame rate, increasing the resolution yields to higher frequency (`PIXEL_CLOCK` gets smaller)



A common image streaming protocol with embedded control codes

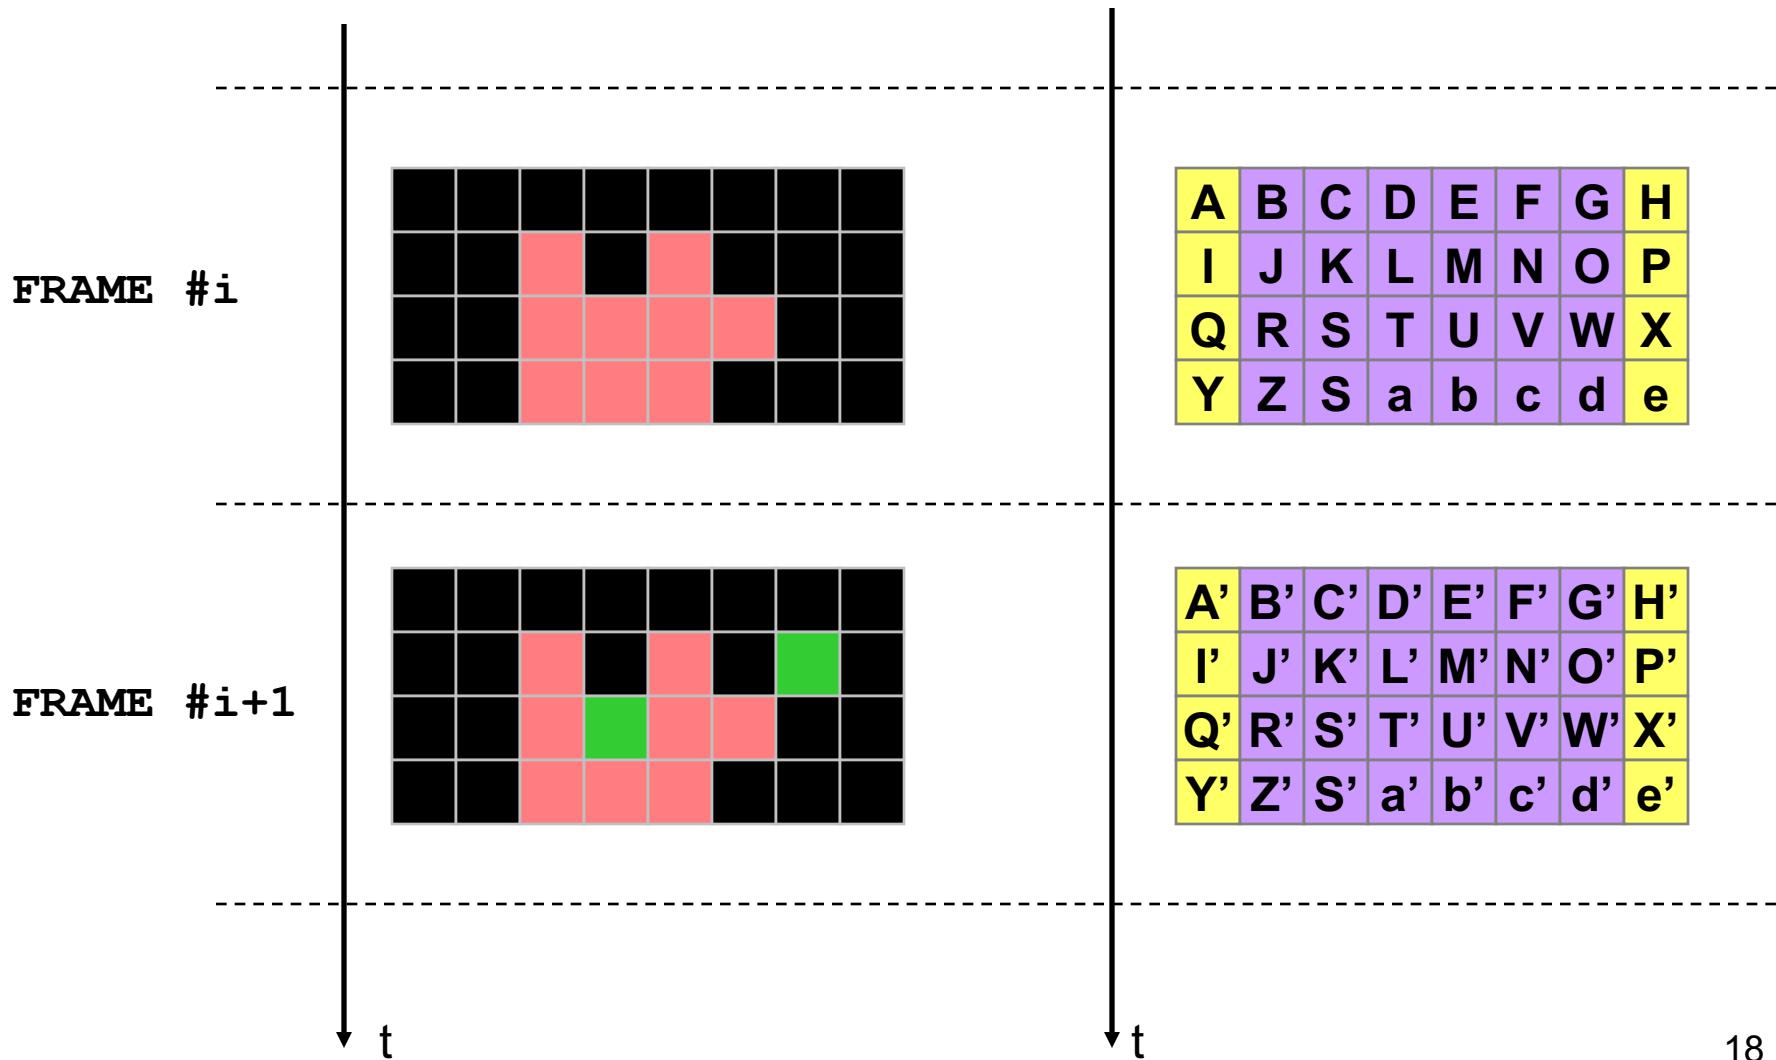


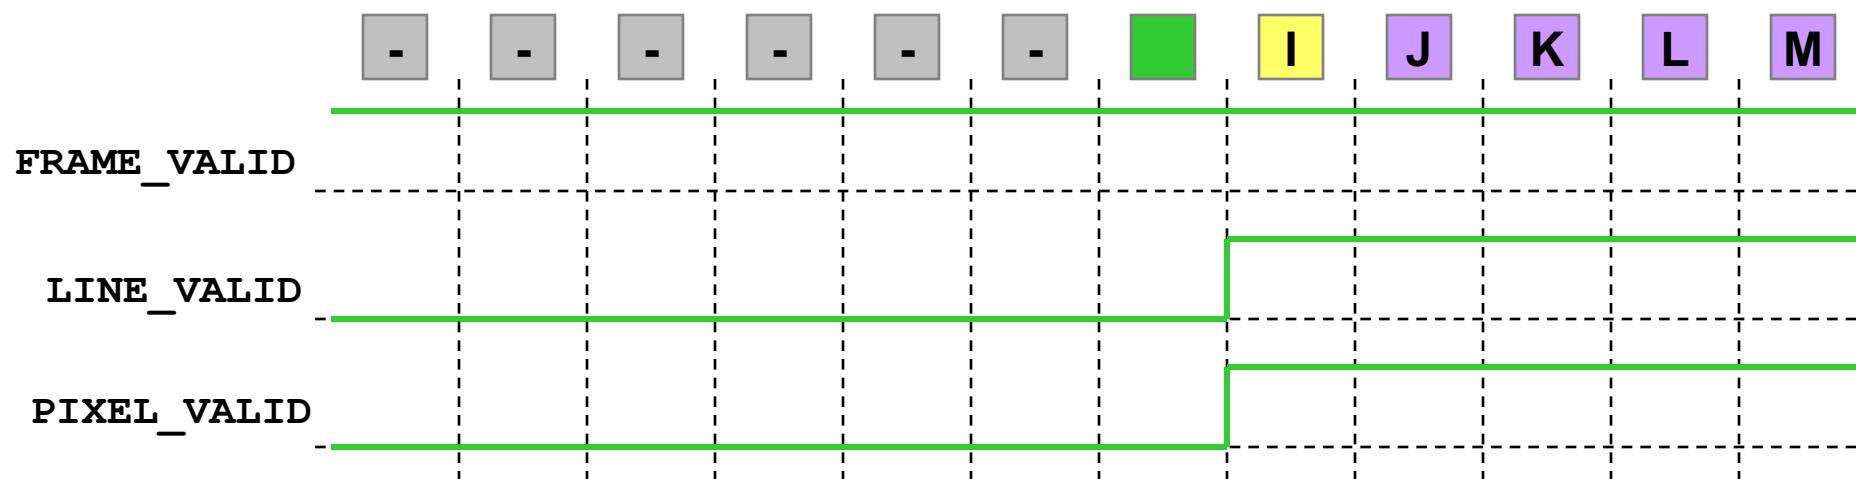
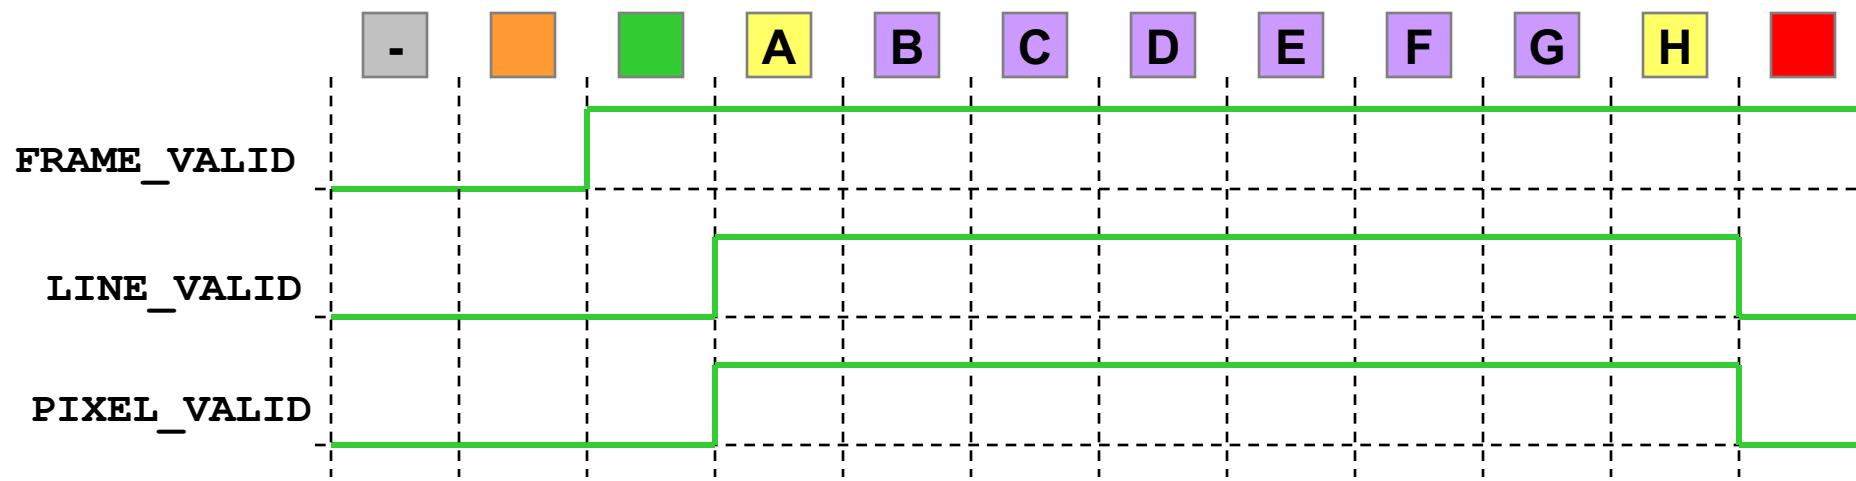
Image streaming contains embedded and reserved codes (**bytecode 0,1,2 e 3**) and a dummy pixel dummy. From this bytecodes, appropriate modules can infer the following signals: **FRAME_VALID**, **LINE_VALID** and also **PIXEL_VALID**

The diagram shows a 10x10 grid of colored squares representing image streaming signals. The grid is divided into two main sections by a vertical red line at column 8. The first section (columns 1-7) contains labels A through H in the top row and A' through H' in the bottom row. The second section (columns 8-10) contains a series of dashes. The colors used are orange, green, red, blue, and grey. A legend on the right side maps these colors to specific signals:

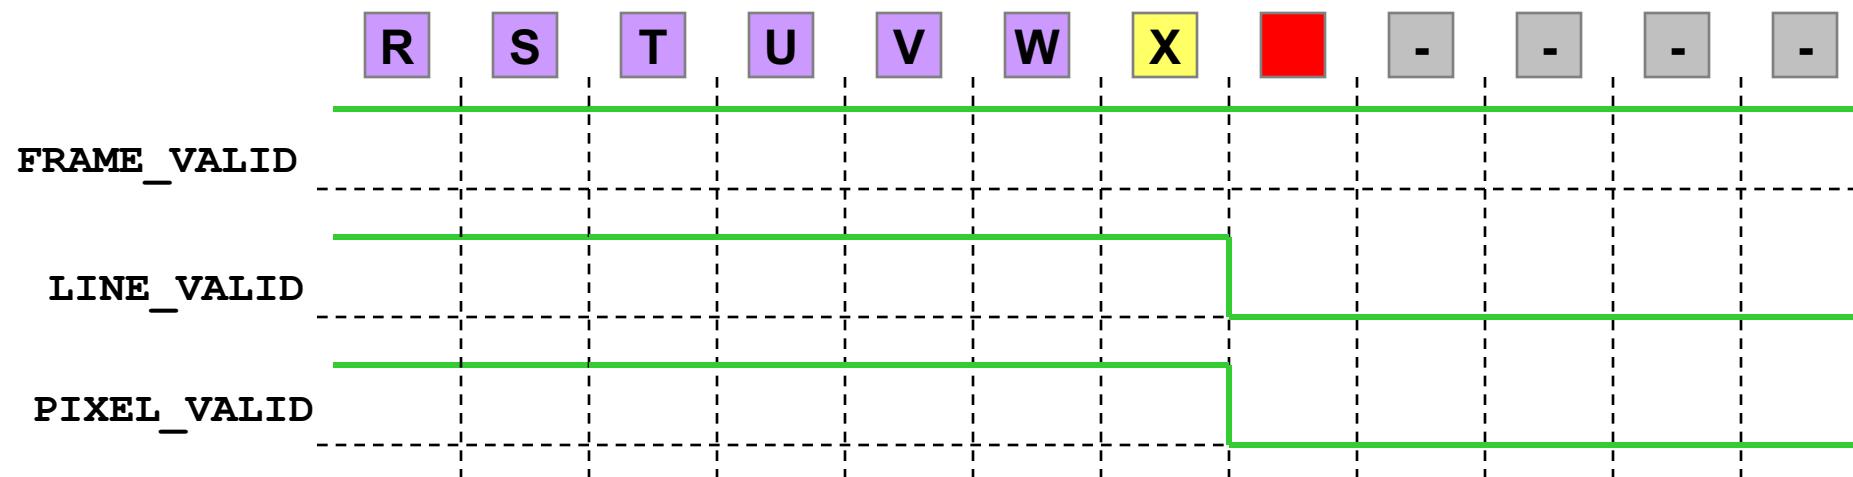
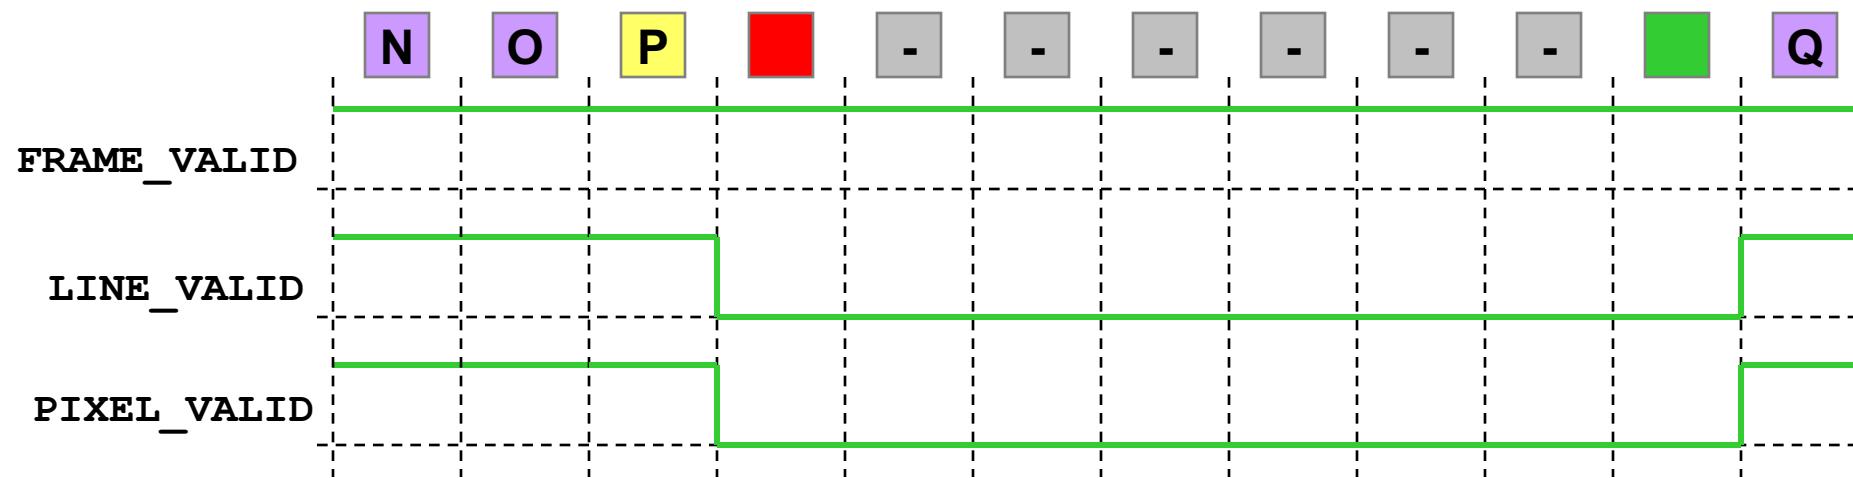
- CODE_START_FRAME (Orange)
- CODE_START_LINE (Green)
- CODE_END_LINE (Red)
- CODE_END_FRAME (Blue)
- PIXEL not valid (Grey)

-	-	-	-	-	-	-	-	-	-
-	-	-	-	-	-	-	-	-	-
A	B	C	D	E	F	G	H	-	-
I	J	K	L	M	N	O	P	-	-
Q	R	S	T	U	V	W	X	-	-
Y	Z	S	a	b	c	d	e	-	-
-	-	-	-	-	-	-	-	-	-
-	-	-	-	-	-	-	-	-	-
A'	B'	C'	D'	E'	F'	G'	H'	-	-
I'	J'	K'	L'	M'	N'	O'	P'	-	-
Q'	R'	S'	T'	U'	V'	W'	X'	-	-
Y'	Z'	S'	a'	b'	c'	d'	e'	-	-
-	-	-	-	-	-	-	-	-	-

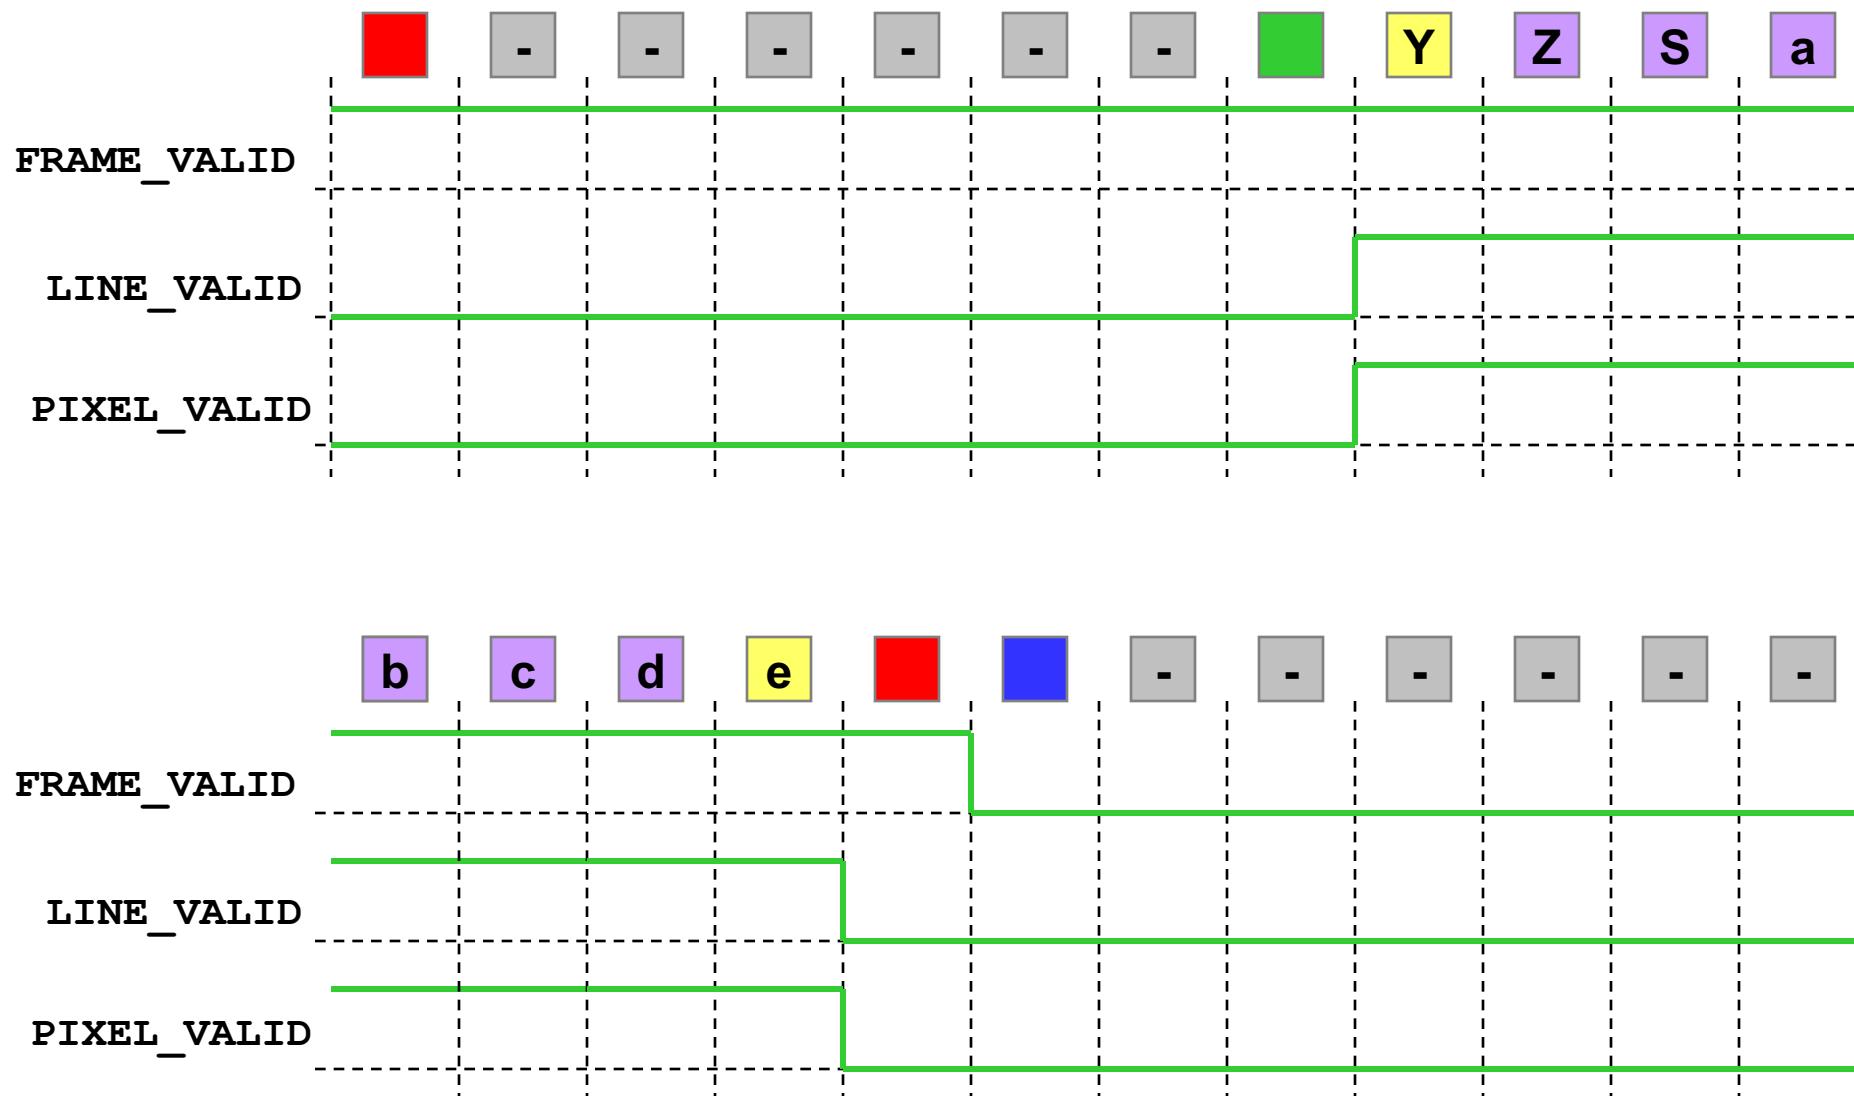
(Almost) standard control signals 1/3



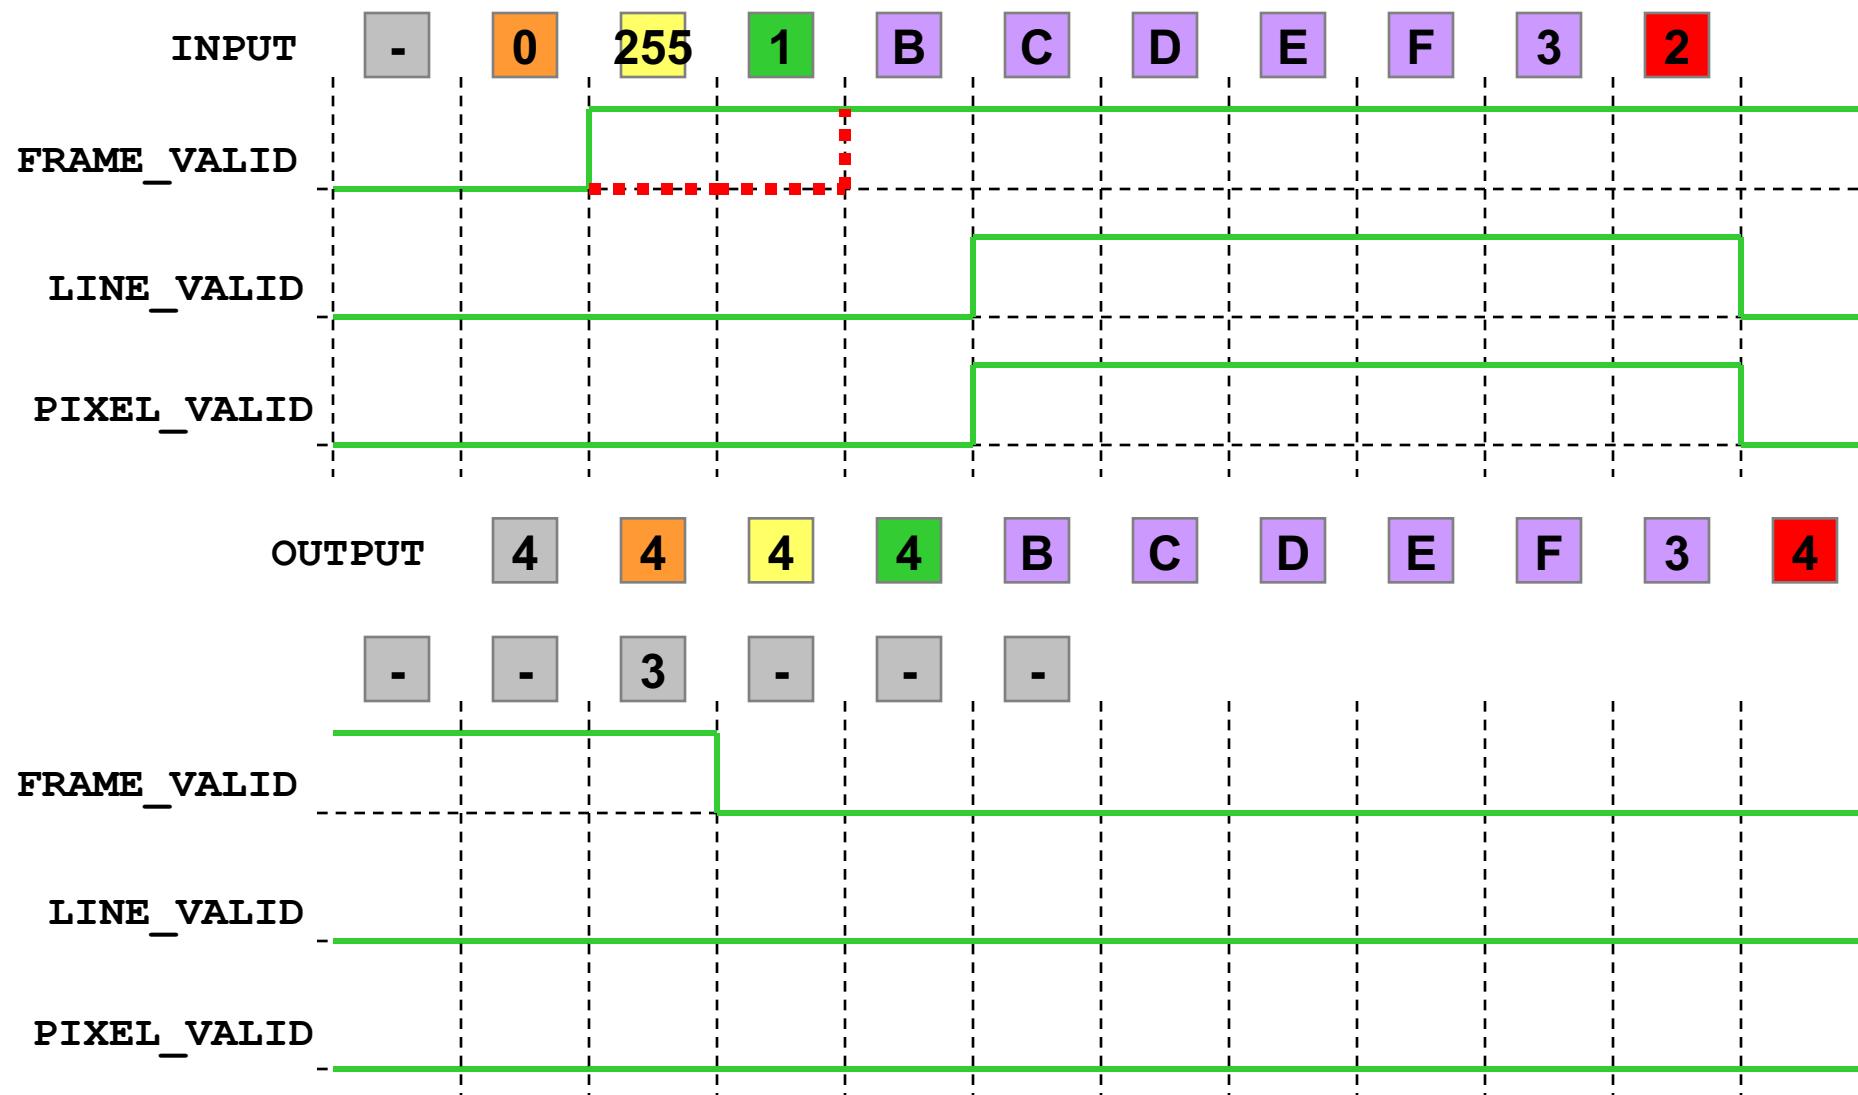
(Almost) standard control signals 2/3



(Almost) standard control signals 3/3



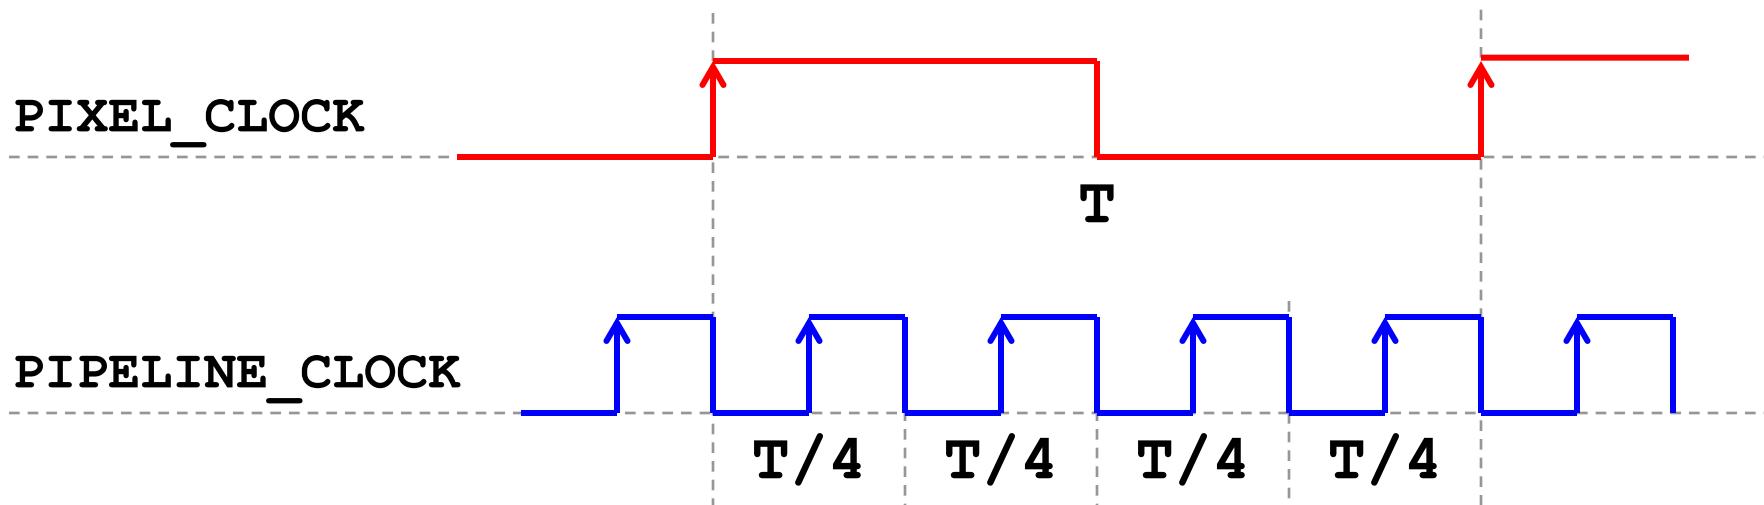
(Almost) standard control signals (1*/3)



Pixel clock e pipeline clock

In many cases, the image processing pipeline works at higher frequency, often referred to as **PIPELINE_CLOCK**, wrt **PIXEL_CLOCK**.

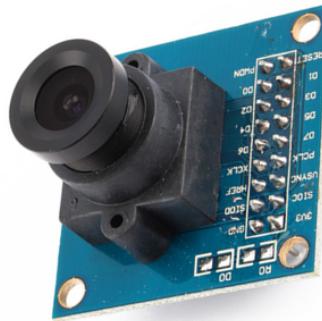
The two clocks are not always aligned (in general a simple FIFO with two clock domains would solve this issue) .



OV 7670: specifications

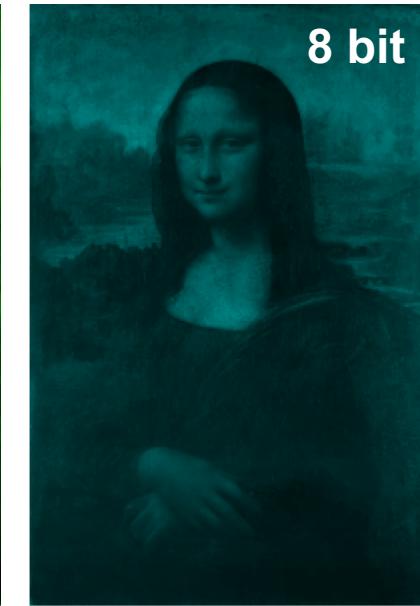
Digital imaging sensor manufactured by Omnivision:

- Pixel size : 3.6 um x 3.6 um
- Resolution : 640x480 (VGA) , 320x240 (QVGA) , etc
- Frame rate : 30 fps
- Color format : RGB, YUV (4:2:2) and YCbCr (4:2:2)
- Scan mode : rolling
- Output : parallel (16 bit for YCbCr)
- Power supply : max 3.0 V
- Programming : I2C (7 bit address 0x21)



YUV or YCbCr (4:2:2) color encoding

- YcbCr and YUV enables a more compact encoding wrt RGB
- YUV/YCbCr vs RGB: 8+8 bit vs 8+8+8 per pixel
- Grayscale component is Luma (Y)
- Color is subsampled (horizontally in 4:2:2)
- CbCr and UV components: resolution is halved (W/2xH)
- In YCbCr or YUV the grayscale component is Luma (Y)
- Grayscale is extracted ignoring chroma components
- With YCbCr or YUV, with 8 bit datapath, each pixel requires more than 1 clock (e.g., 2 clocks for 4:2:2)
- In our setup, the OV7670 is configured for YUV 4:2:2
- Other sensors (e.g. Aptina/ONSEMI MT9v034C), use a different color encoding scheme based on Bayer pattern

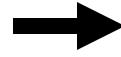


Original

Red

Green

Blue



Original

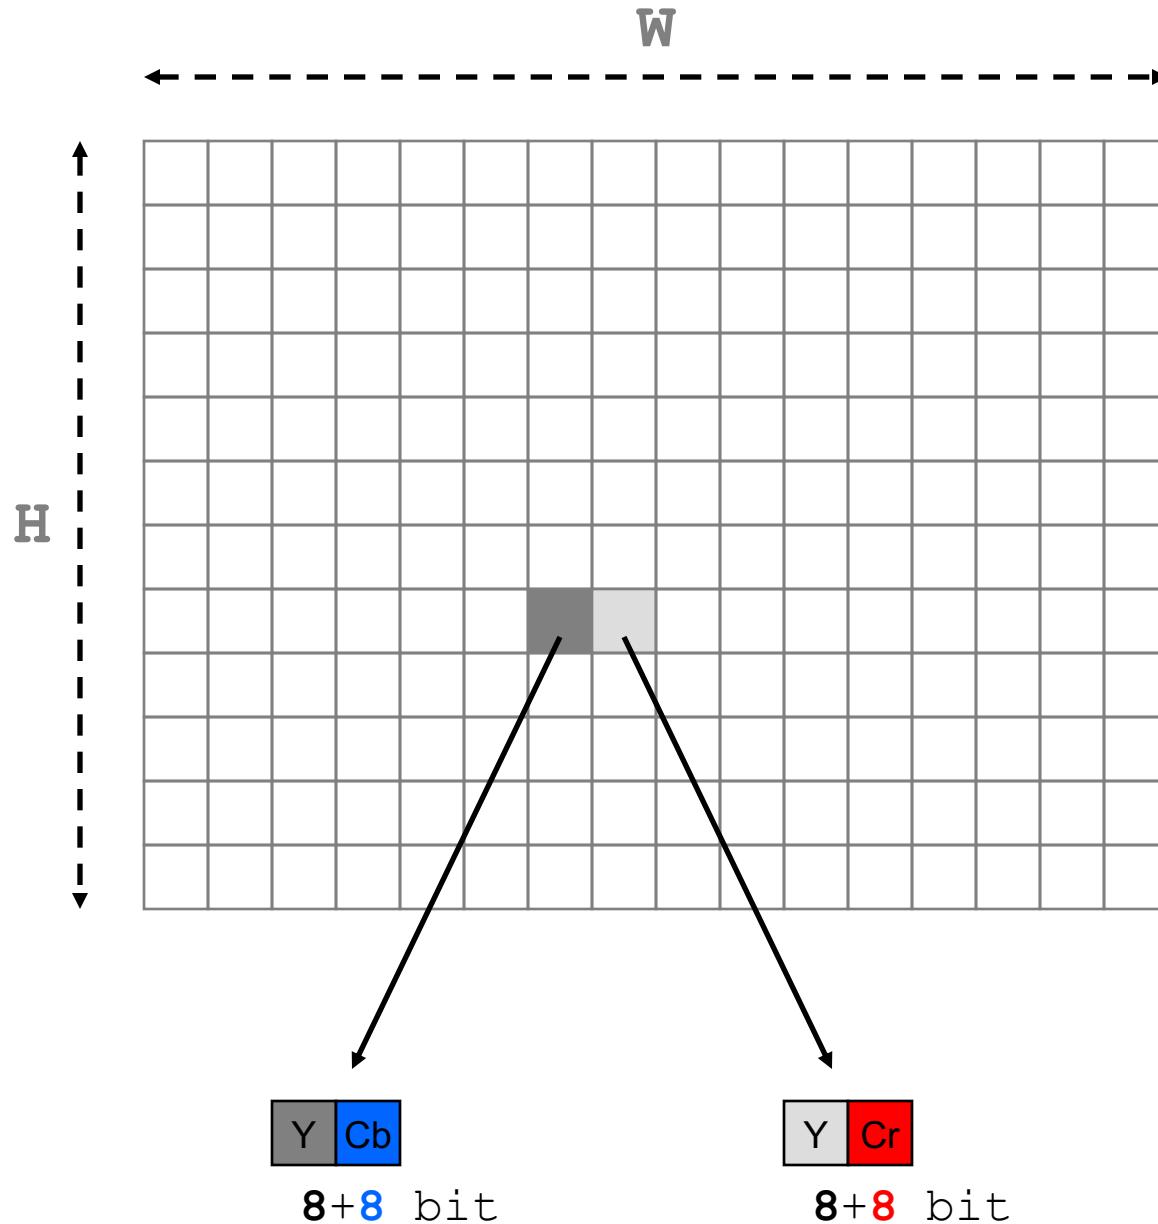
Luma Y

Chroma Cb Chroma Cr

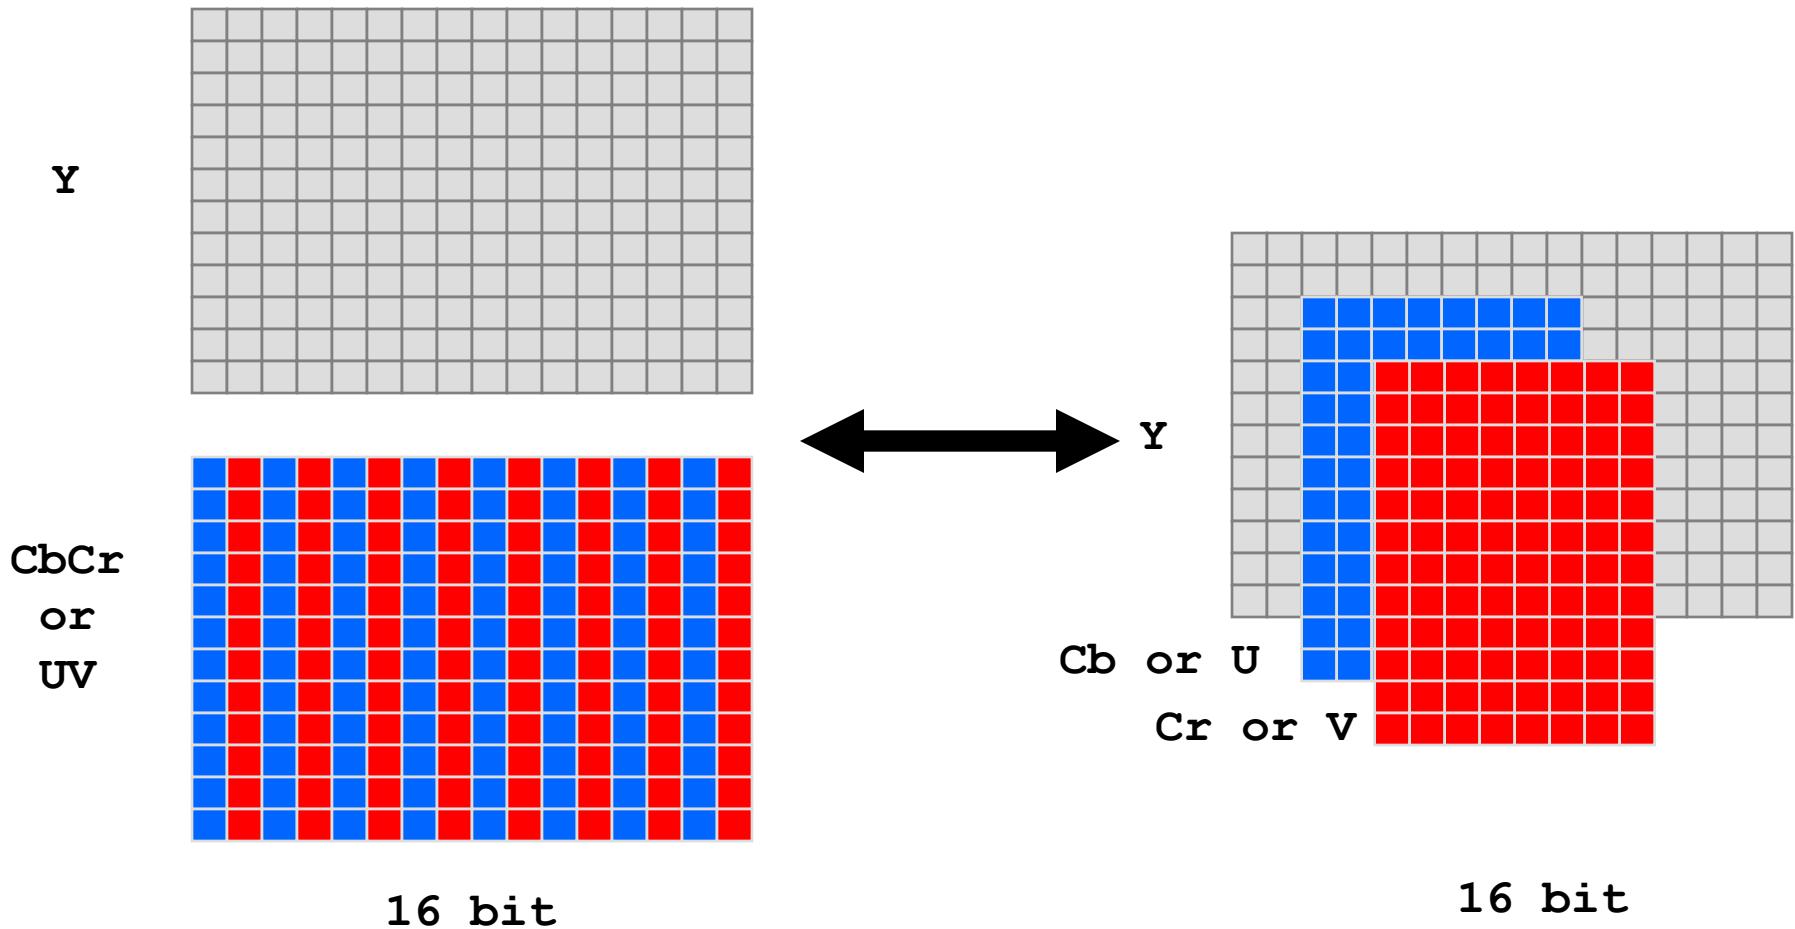
Alternative chroma components encoding for YCbCr or YUV:

- **4:4:4 (no subsampling)** – 24 bits
- **4:4:0 (vertical subsampling)** – 16 bits
- **4:2:0 (vertical and horizontal subsampling)** 12 bits

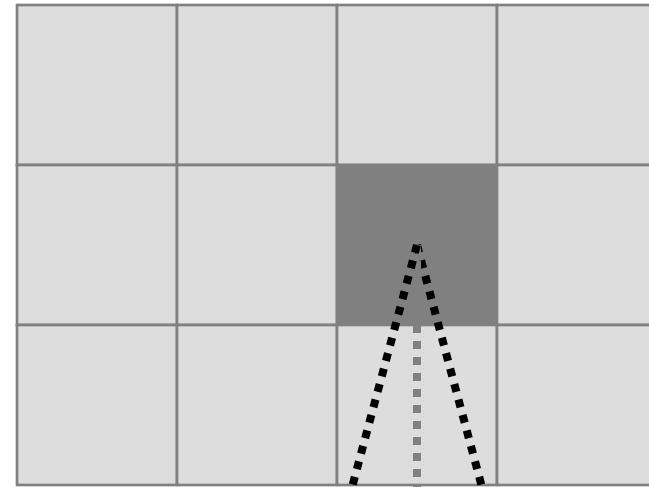
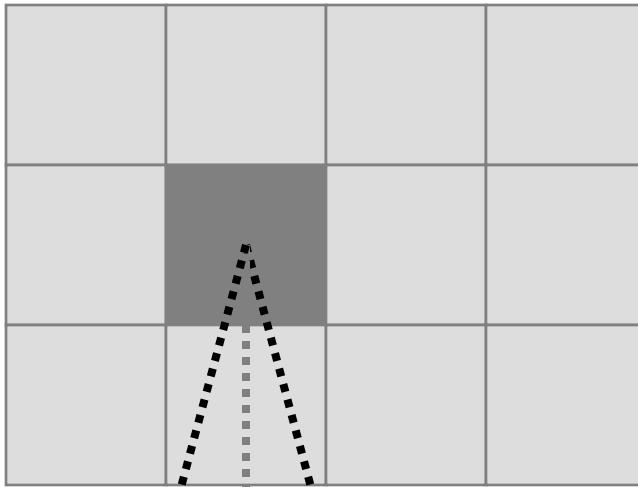
Subsampling mode	Horizontal downscale	Vertical downscale	Bits per pixel*
4:4:4	1x	1x	24
4:2:2	2x	1x	16
4:4:0	1x	2x	16
4:2:0	2x	2x	12



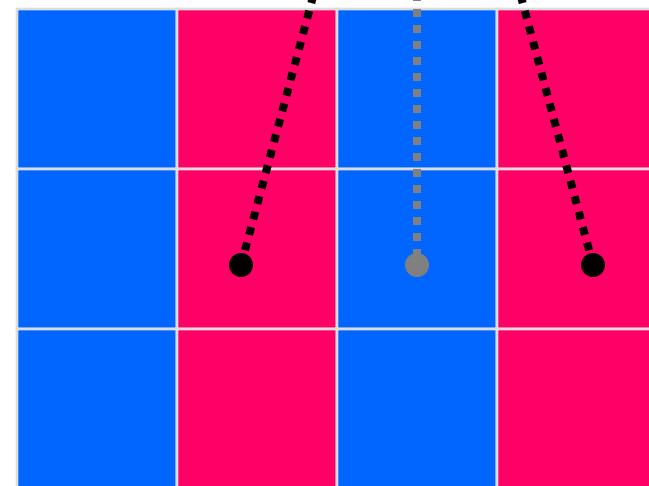
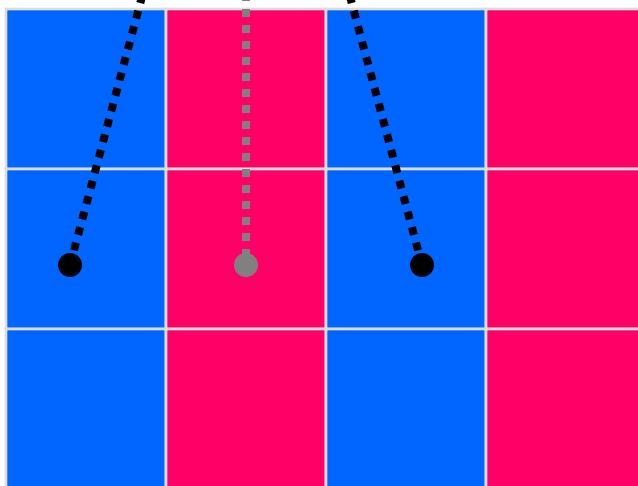
YCbCr or YUV 4:2:2

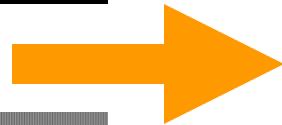


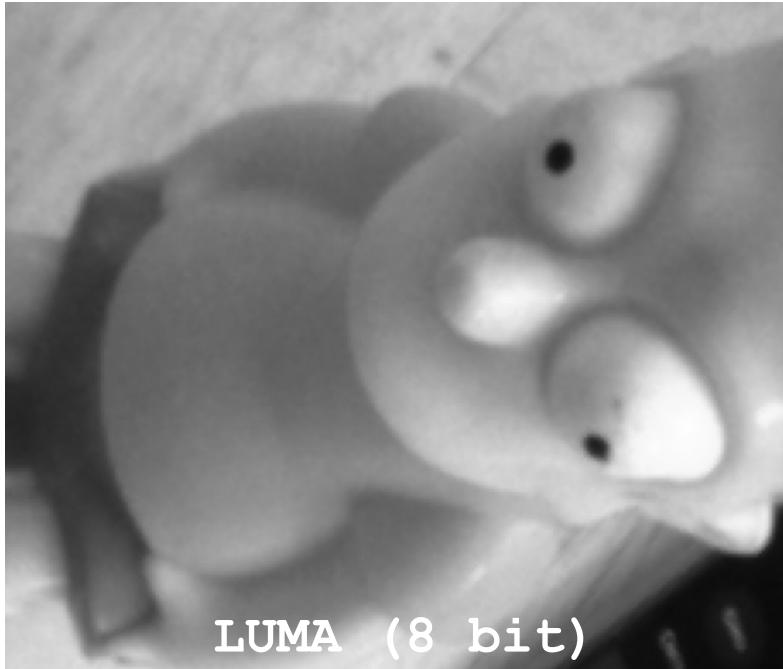
YCbCr or YUV 4:2:2



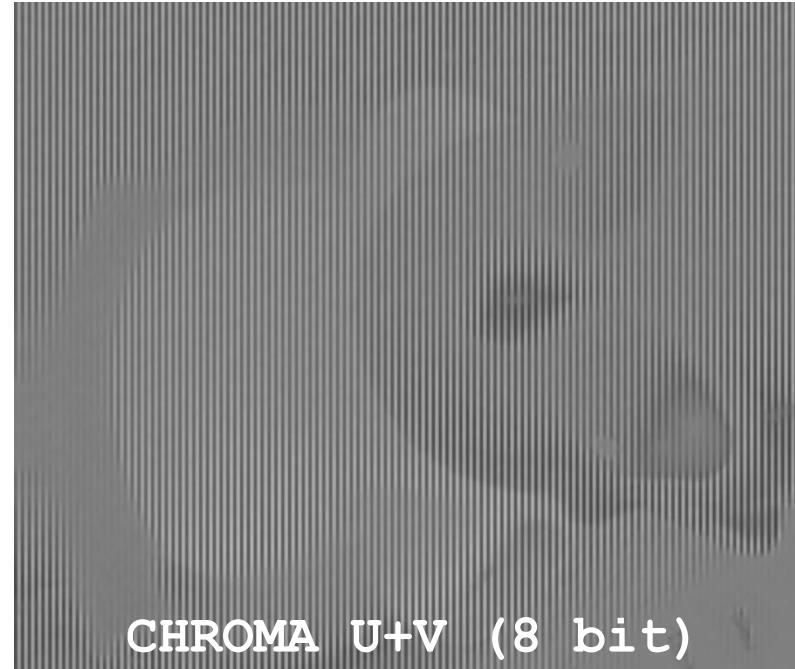
Average chromas?



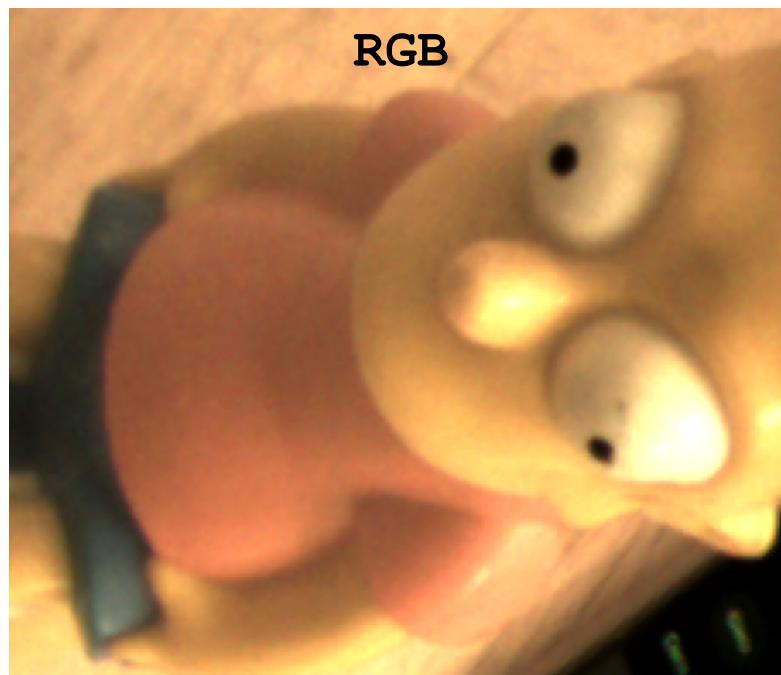




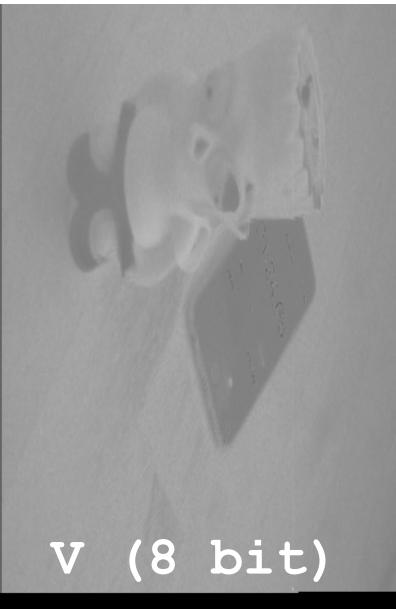
LUMA (8 bit)



CHROMA U+V (8 bit)



RGB

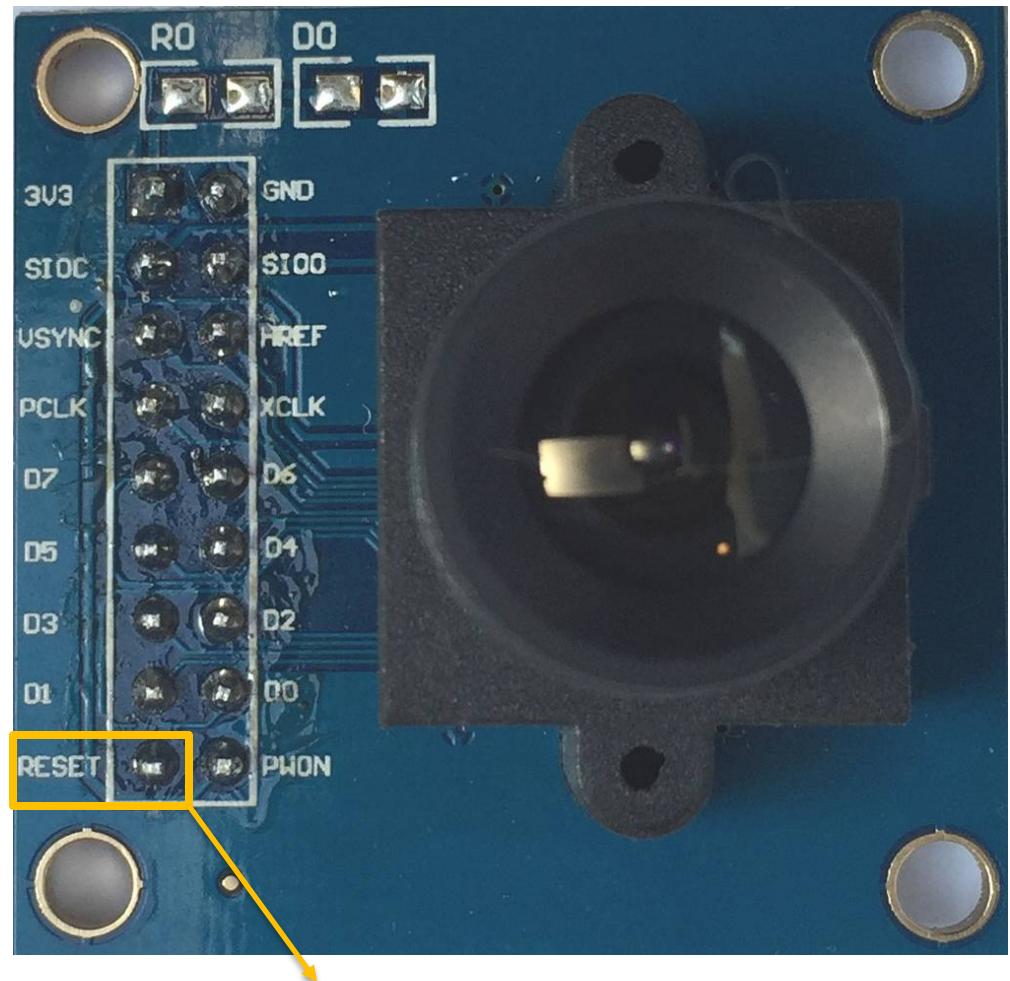


OV 7670: pinout

3V3	GND
SIOC	SIOD
VSYNC	HREF
PCLK	XCLK
D7	D6
D5	D4
D3	D2
D1	D0
RESET*	PWDN

RESET is active low

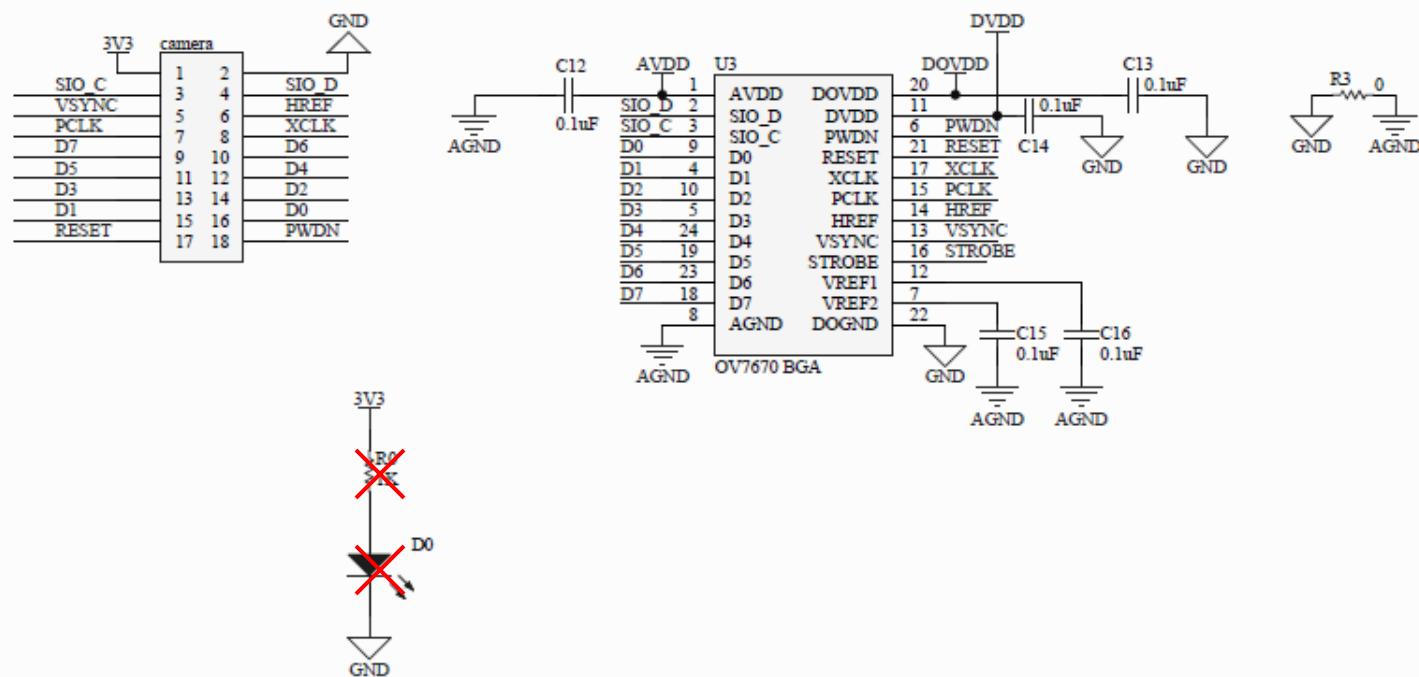
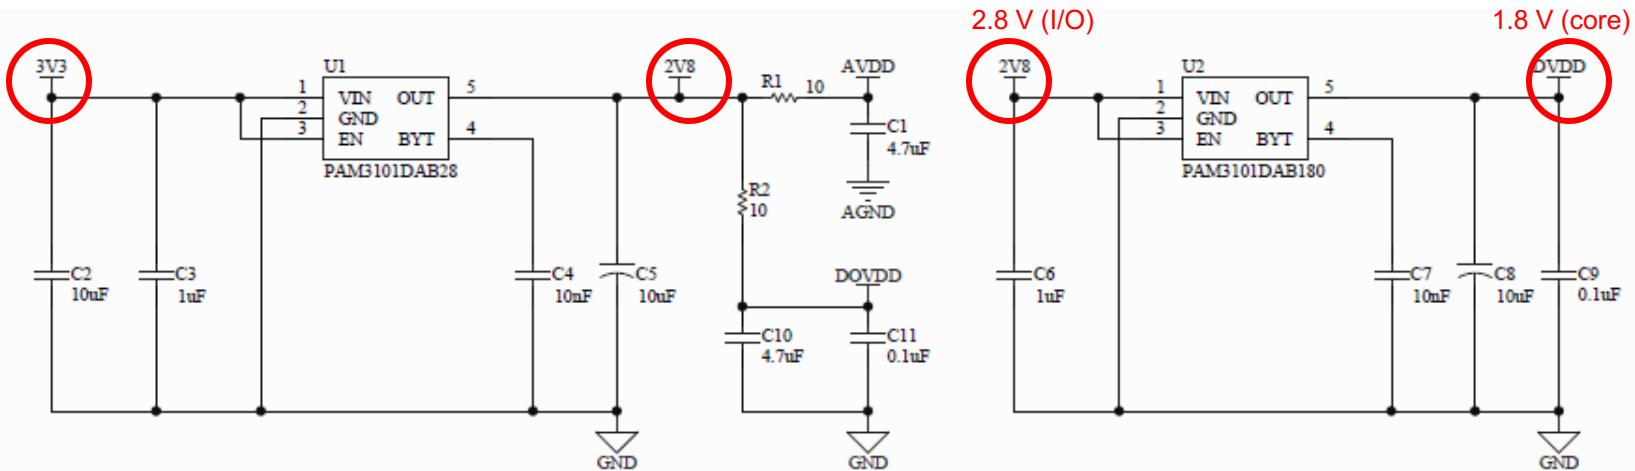
*Internal pull-down resistor
PWDN is active high*



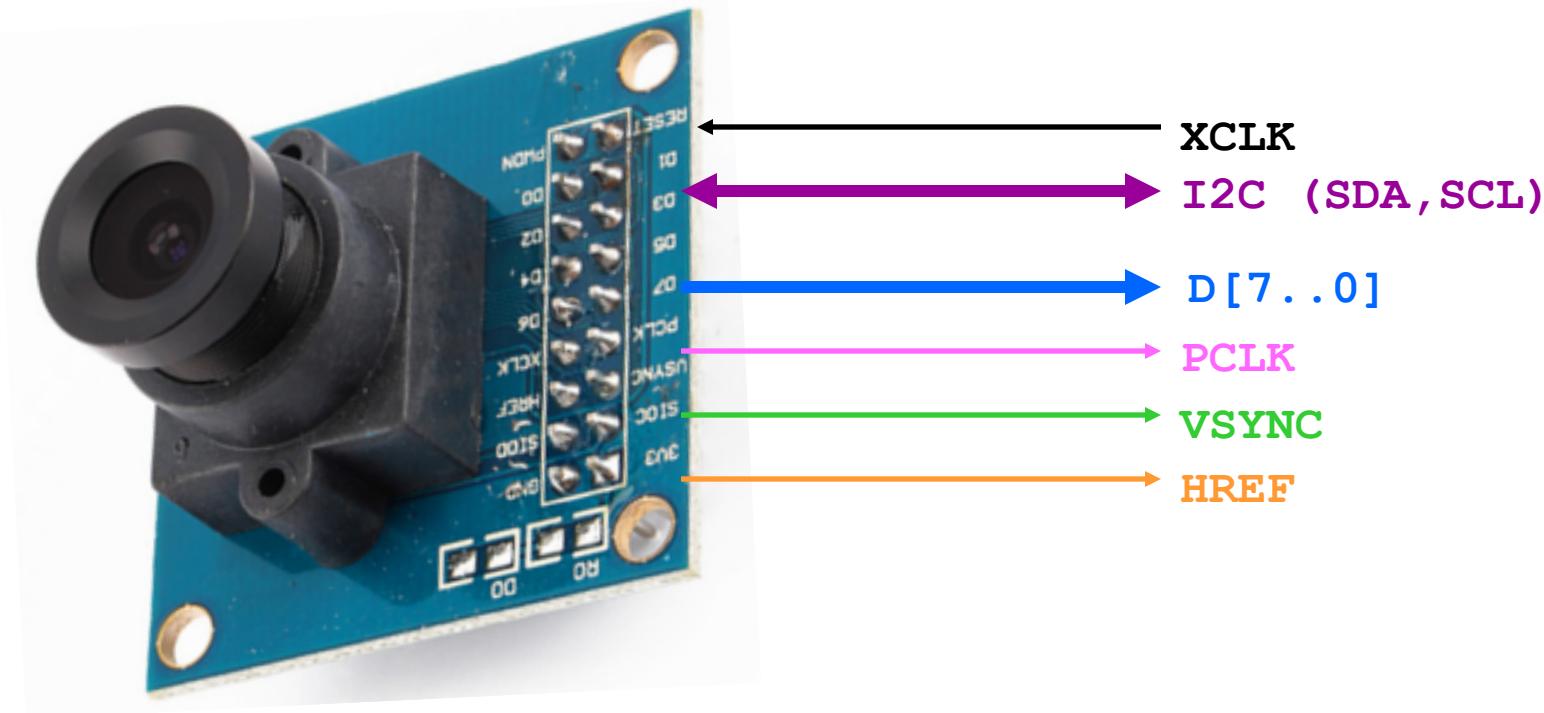
*Since it is active low,
it is renamed **RESET****

OV 7670: pin description

Pin	Type	Description
VDD/3V3	Supply	Power supply 3.3V
GND	Supply	Ground level
SIOC/SCL	Input	I2C clock
SIOD/SDA	Input/Output	I2C data
VSYNC	Output	Vertical synchronization
HREF	Output	Horizontal synchronization
PCLK	Output	Pixel clock
XCLK	Input	System clock
D0-D7	Output	Video parallel output
RESET*	Input	Reset (Active low)
PWDN	Input	Power down (Active high)

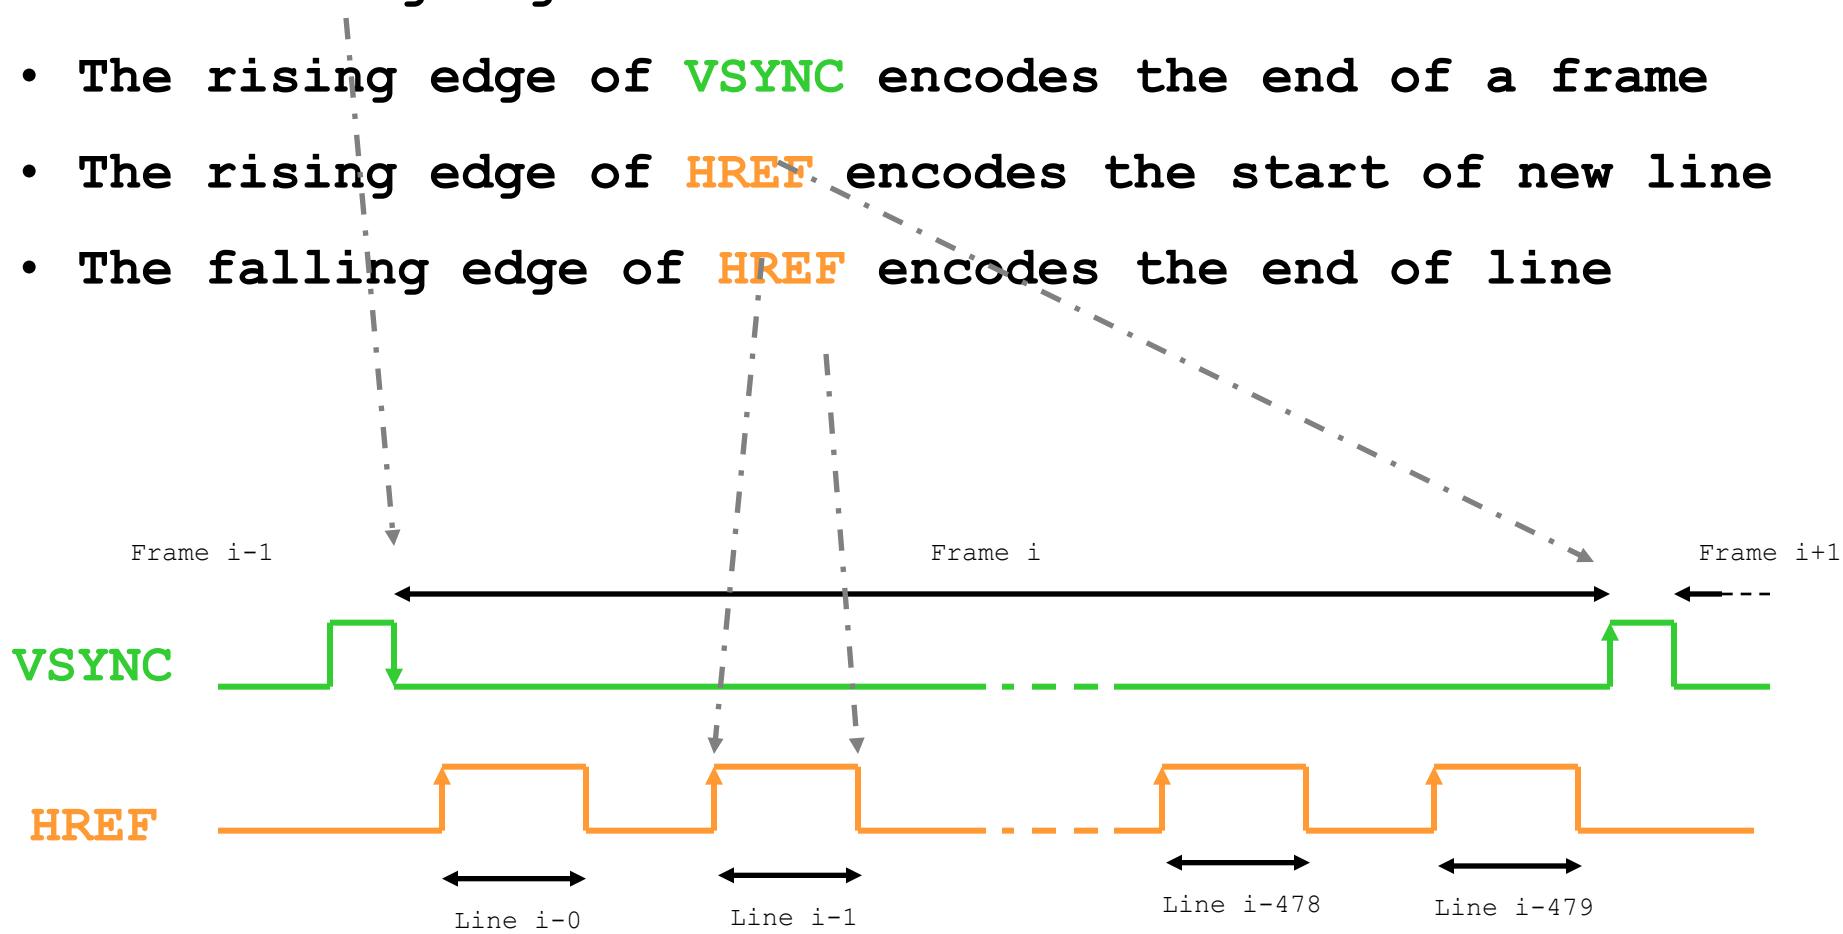


OV 7670: functional description



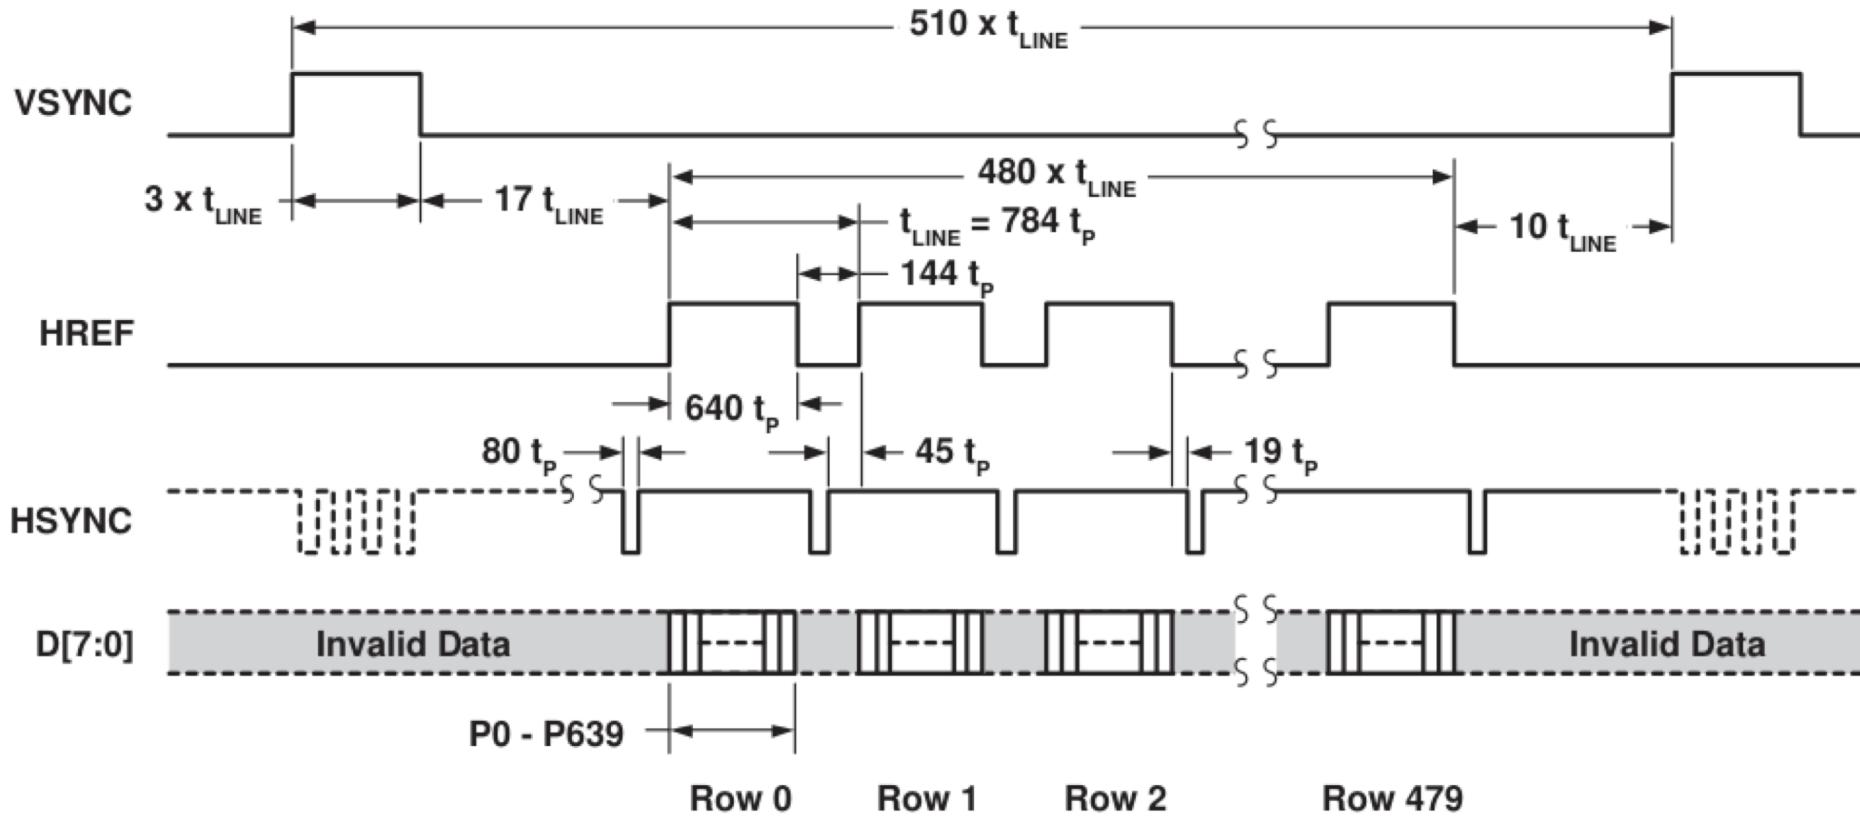
- The OV 7670 takes as input the XCLK clock signal (min frequency 10 MHz, max 48 MHz)
- The output pixel clock (PCLK) is programmable (I2C)
- The synchronous (PCLK) 8 bit data output D[7..0] and VSYNC and HREF signal encode the image content

- The falling edge of **VSYNC** encodes a new frame
- The rising edge of **VSYNC** encodes the end of a frame
- The rising edge of **HREF** encodes the start of new line
- The falling edge of **HREF** encodes the end of line

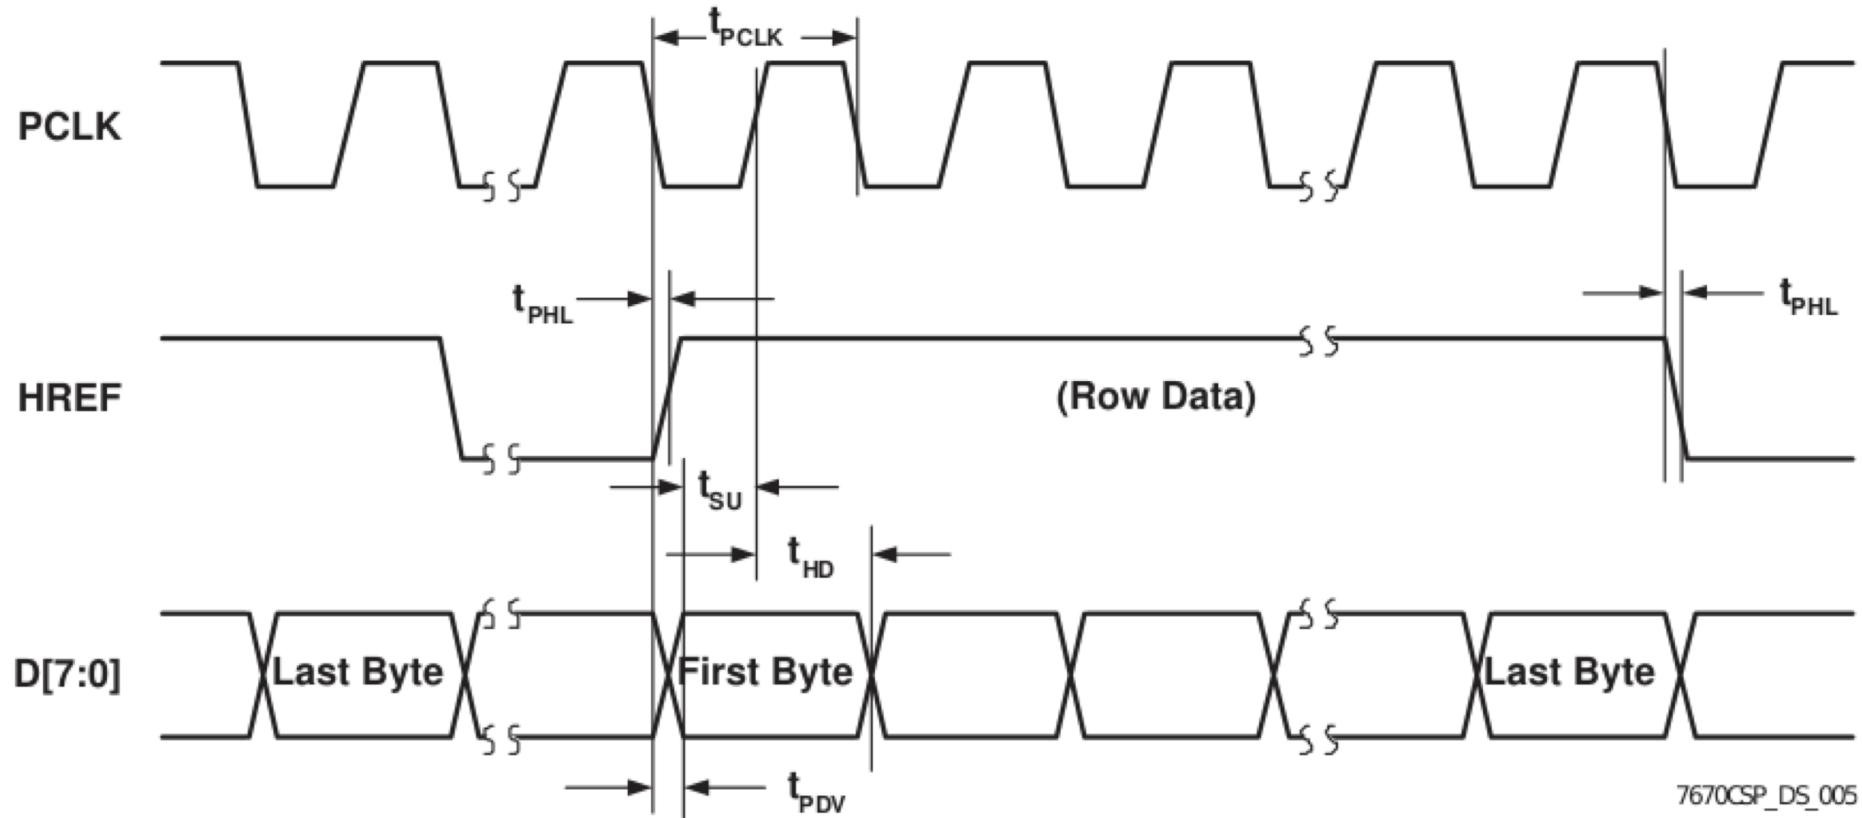


One byte is not a pixel! See YCbCr/YUV section

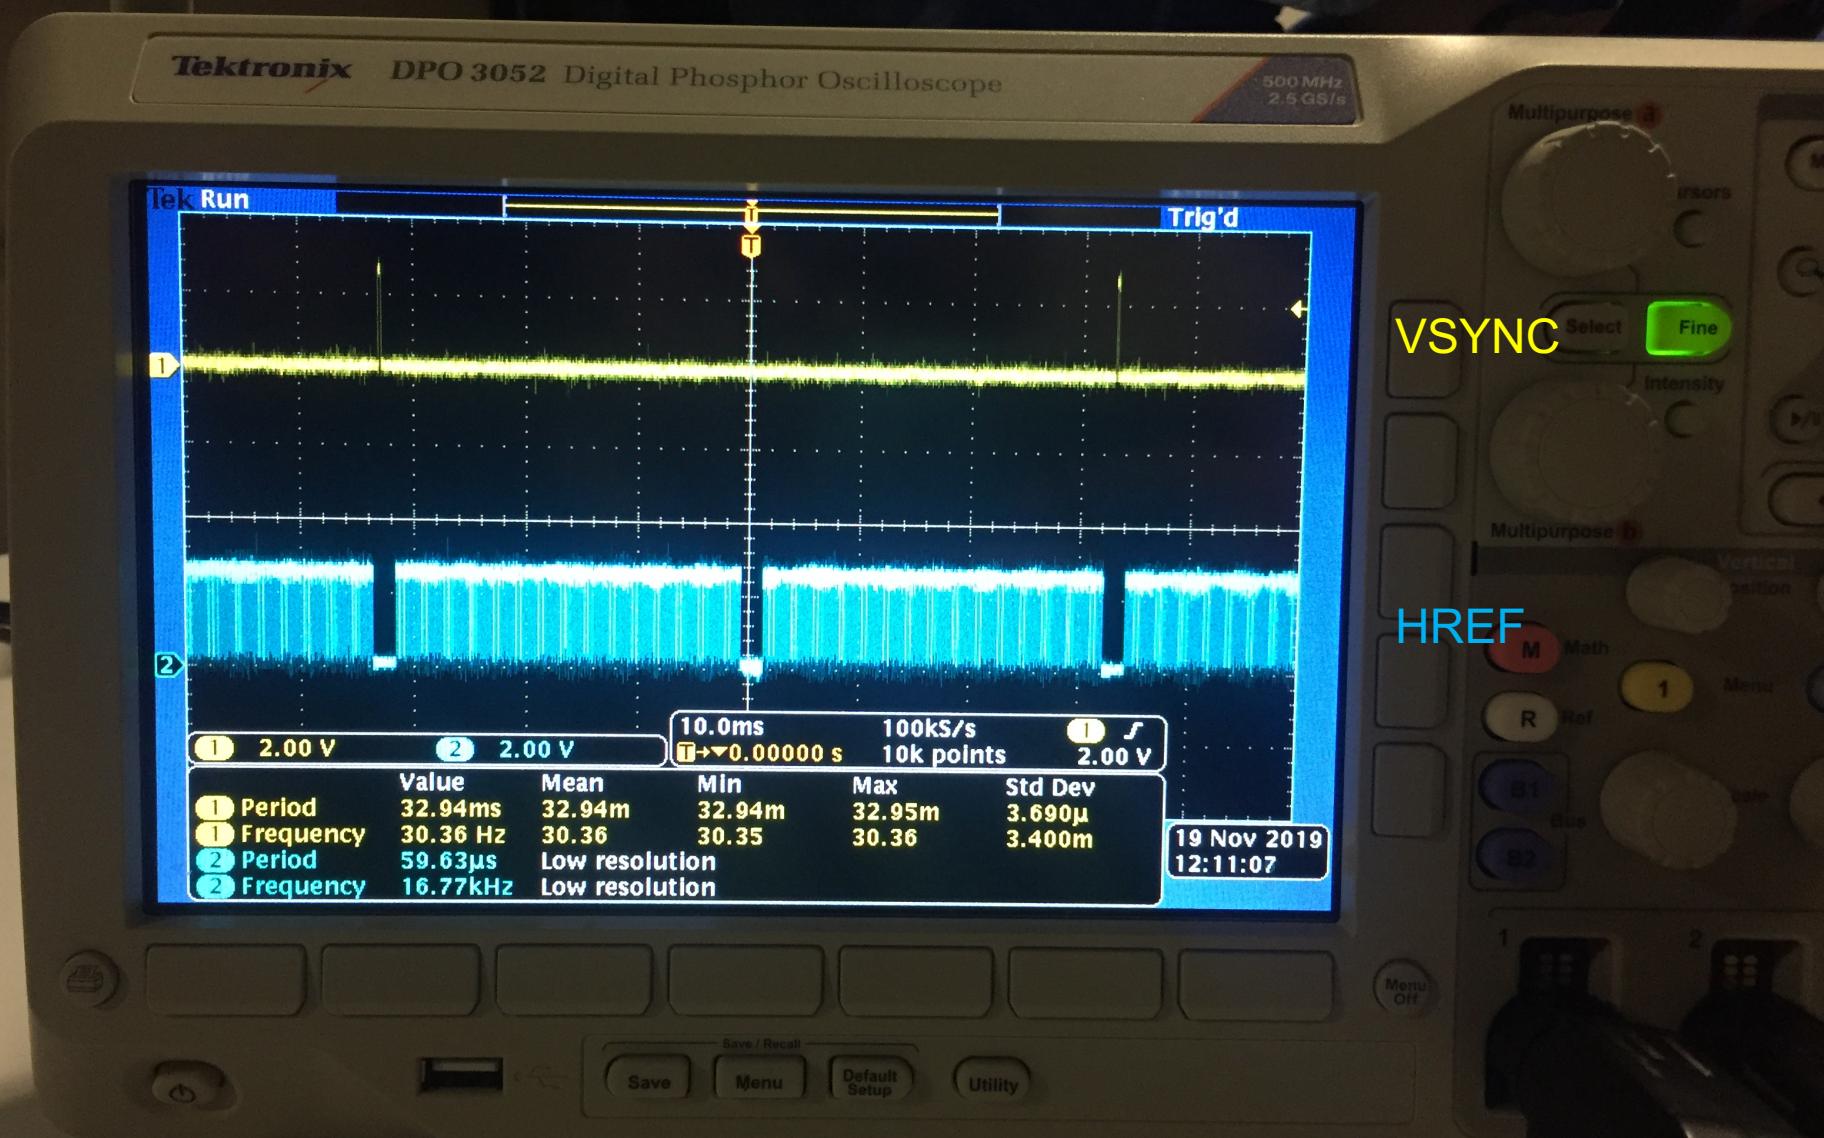
OV 7670: detailed waveforms 1/2



OV 7670: detailed waveforms 2/2



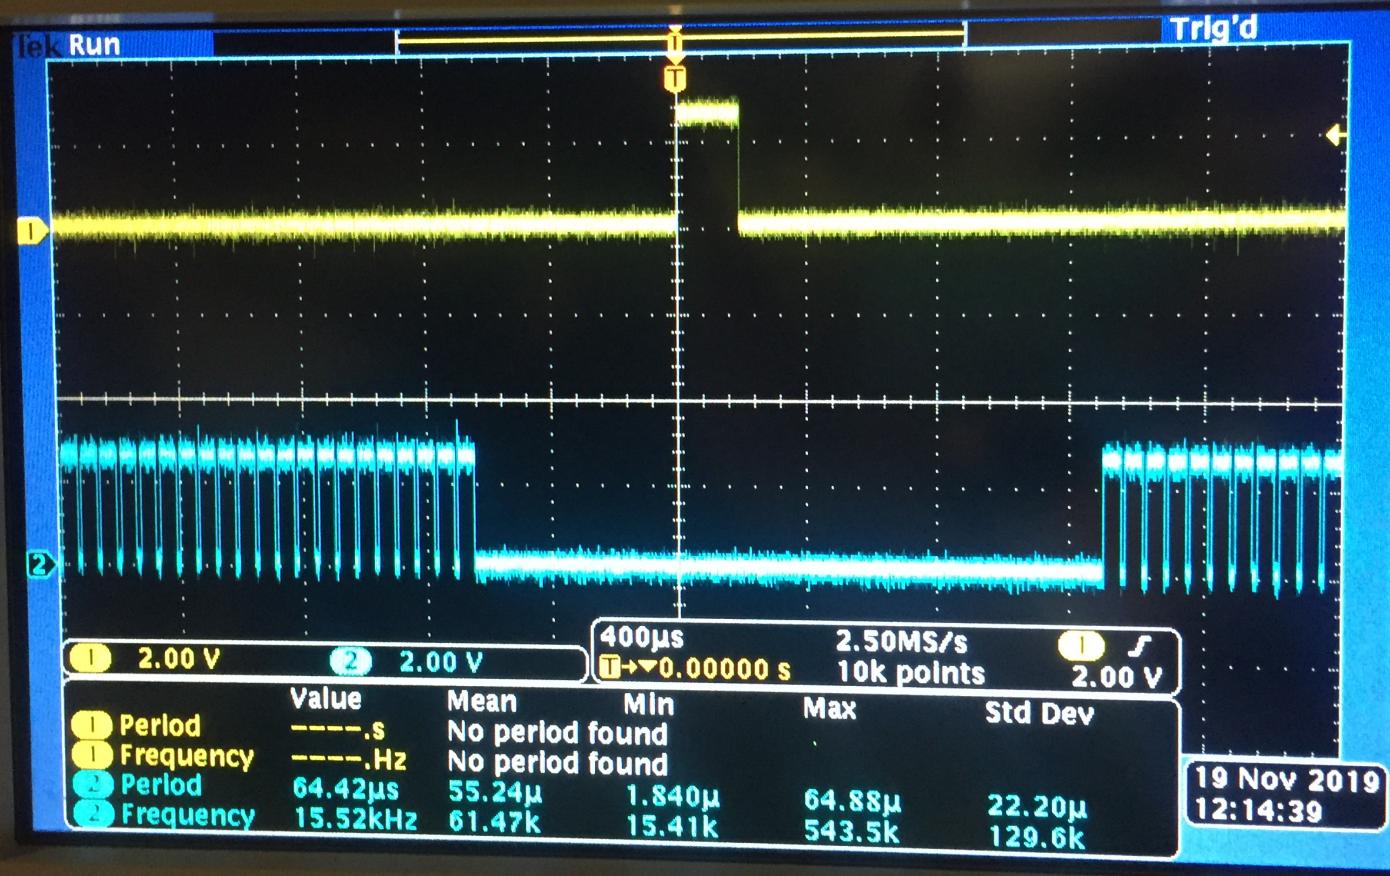
XCLK 24 MHz, no I2C programming 1/3



XCLK 24 MHz, no I2C programming 2/3

Tektronix DPO 3052 Digital Phosphor Oscilloscope

500 MHz
2.5 GS/s

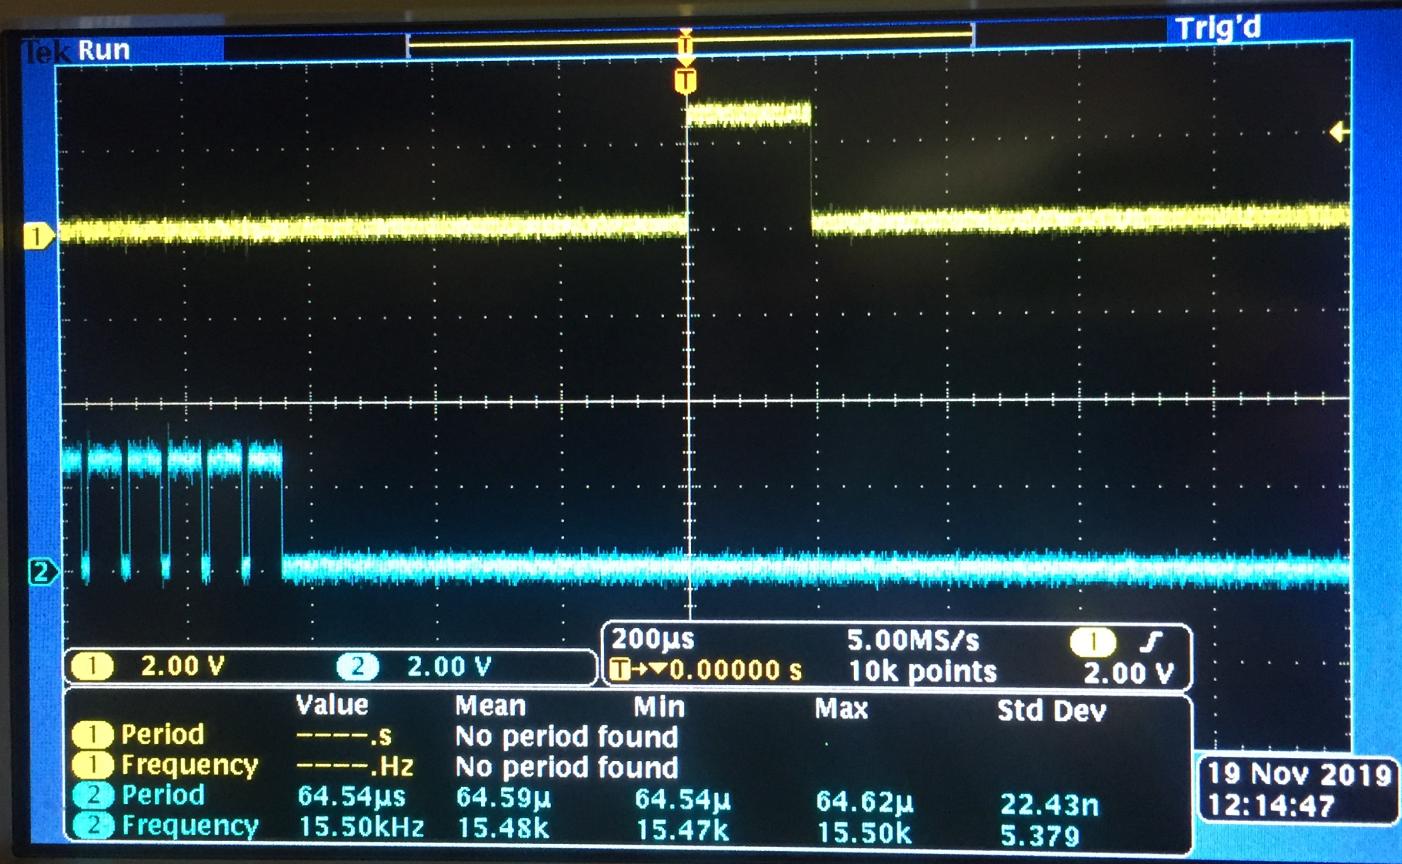


XCLK 24 MHz, no I2C programming 3/3

Tektronix

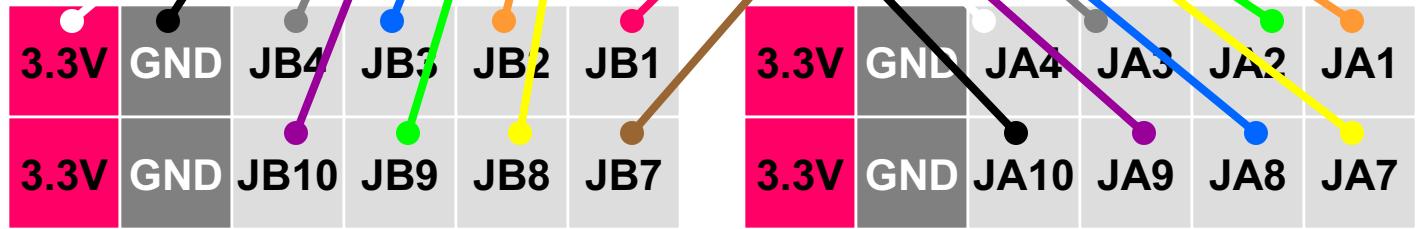
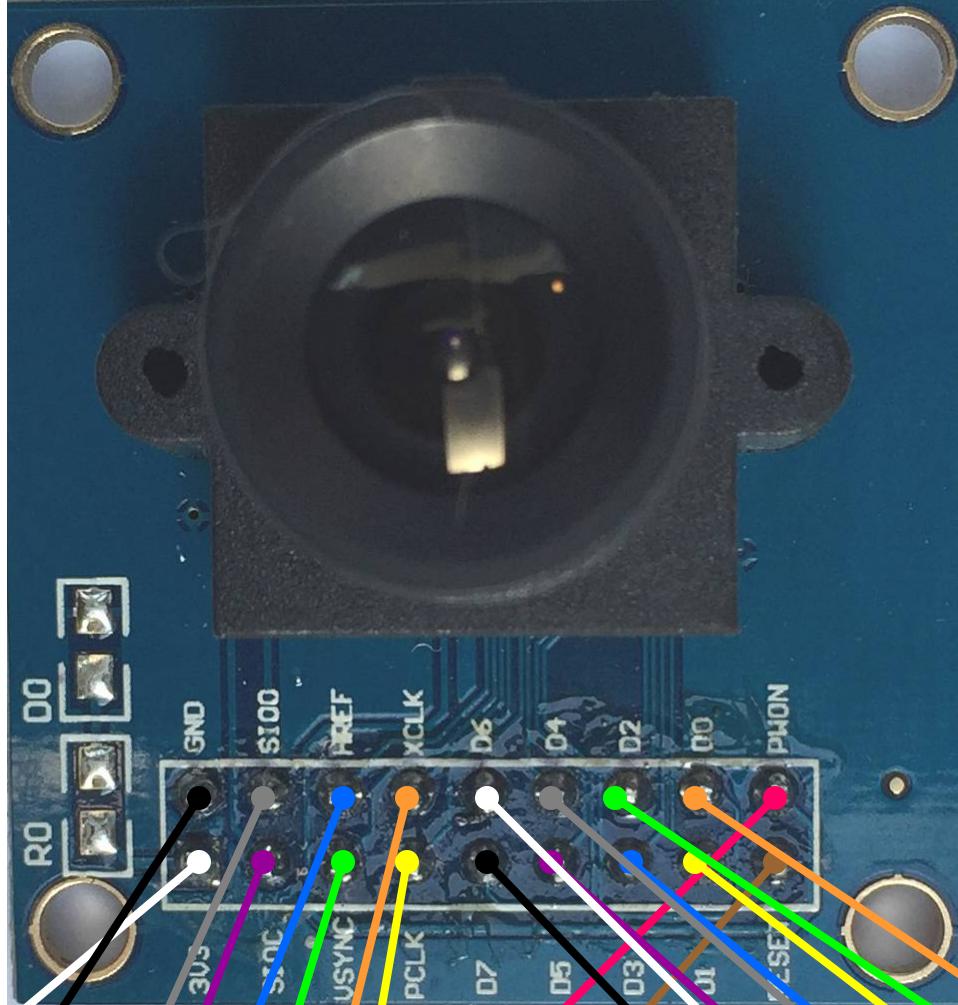
DPO 3052 Digital Phosphor Oscilloscope

500 MHz
2.5 GS/s



VSYNC

HREF



PMOD E PIN

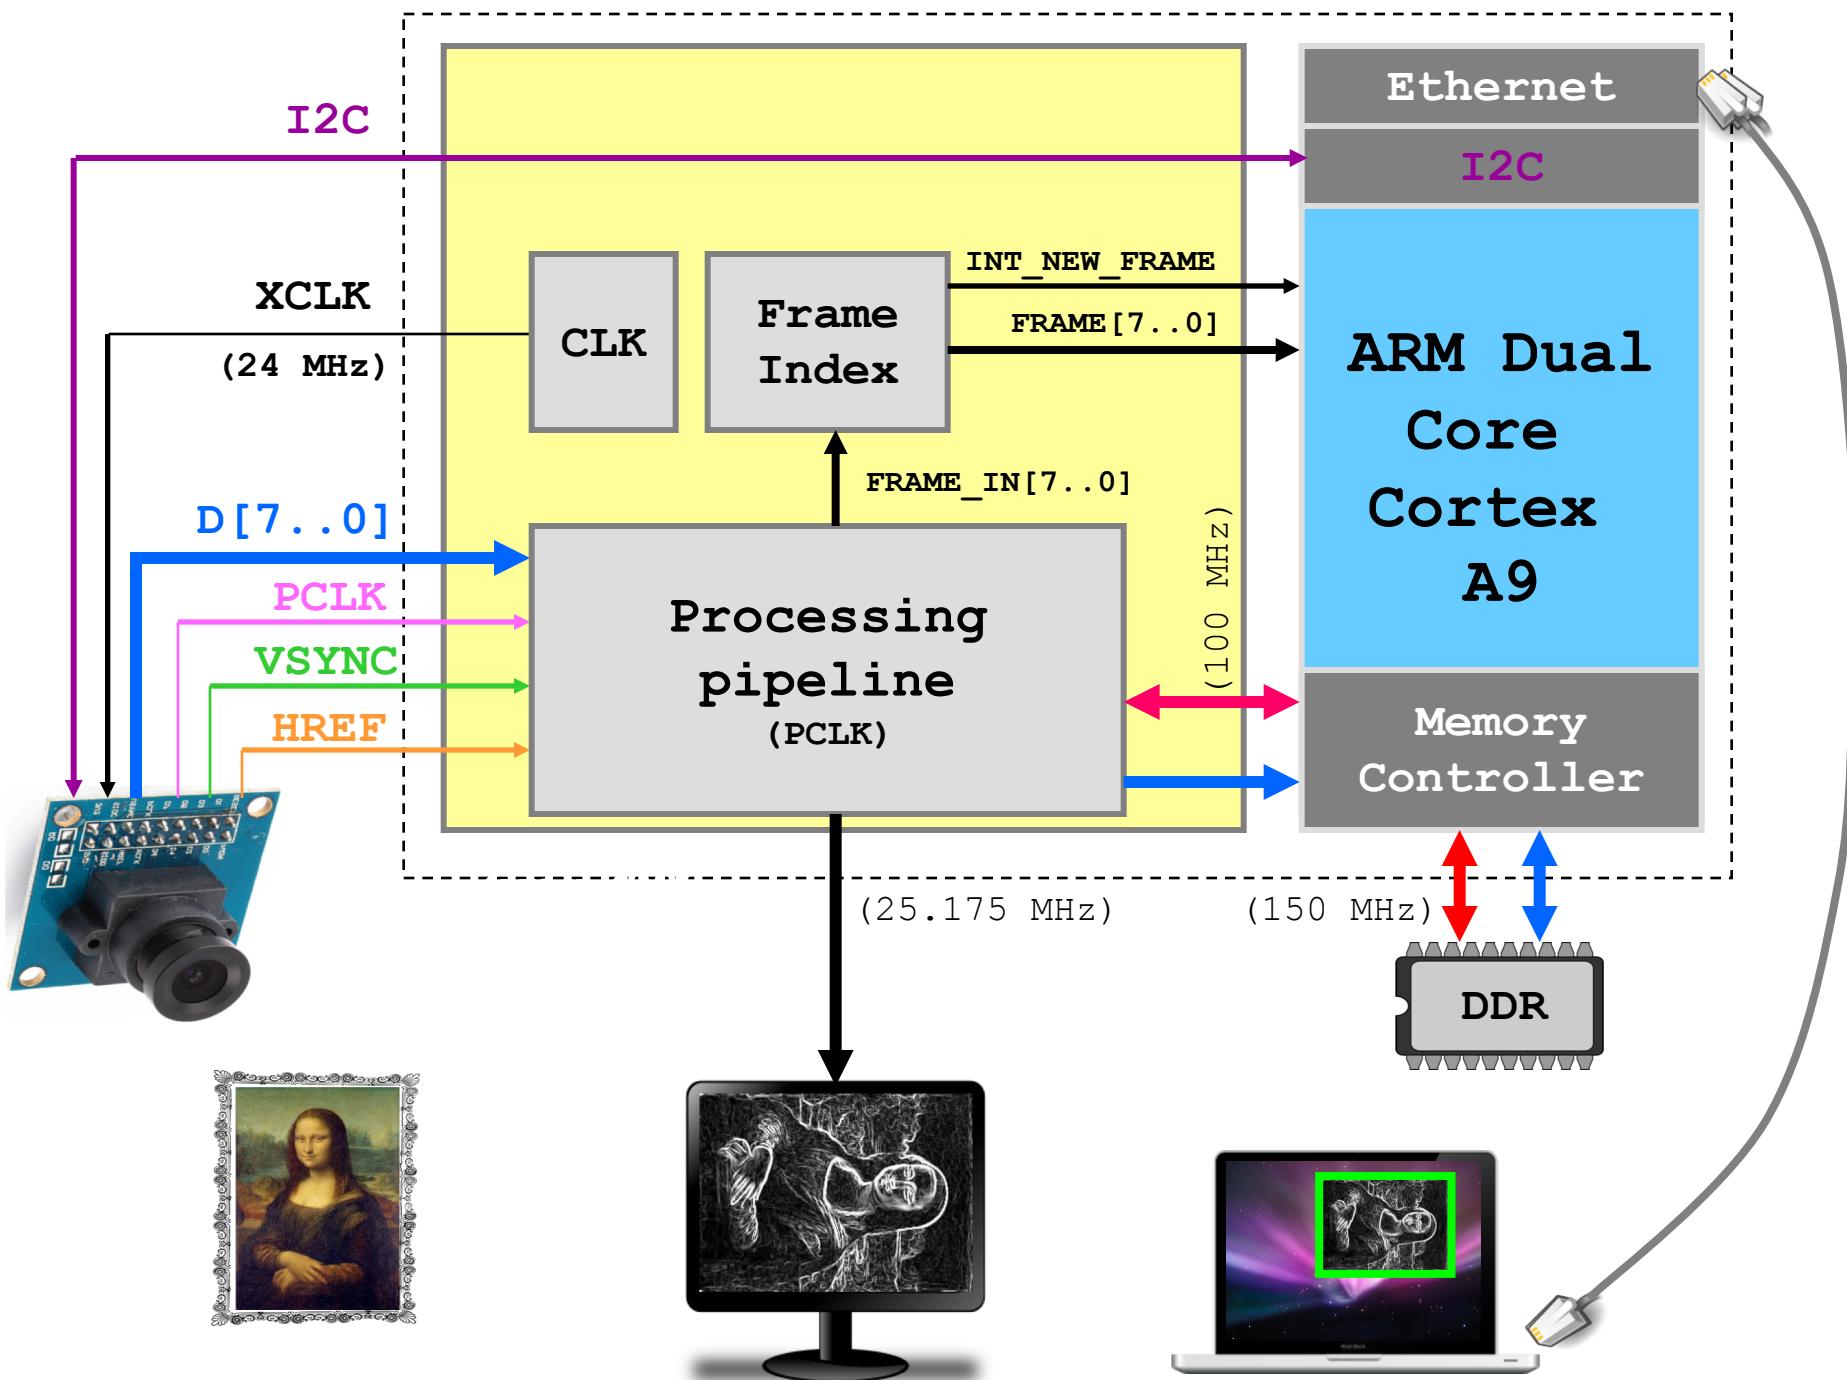
	OV 7670	PMOD A,B PIN	ZYNQ PIN
	3.3 V	VCC/JB6	-
	GND	GND/JB5	-
JE3	SIOD/SDA	JB4	W8
JE2	SIOC/SCL	JB10	V8
	HREF	JB3	V10
	VSYNC	JB9	V9
	XCLK	JB2	W11
	PCLK	JB8	W10
	PWDN	JB1	W12
	RESET*	JB7	V12
	D0	JA1	Y11
	D1	JA7	AB11
	D2	JA2	AA11
	D3	JA8	AB10
	D4	JA3	Y10
	D5	JA9	AB9
	D6	JA4	AA9
	D7	JA10	AA8

Constraints (Vivado):

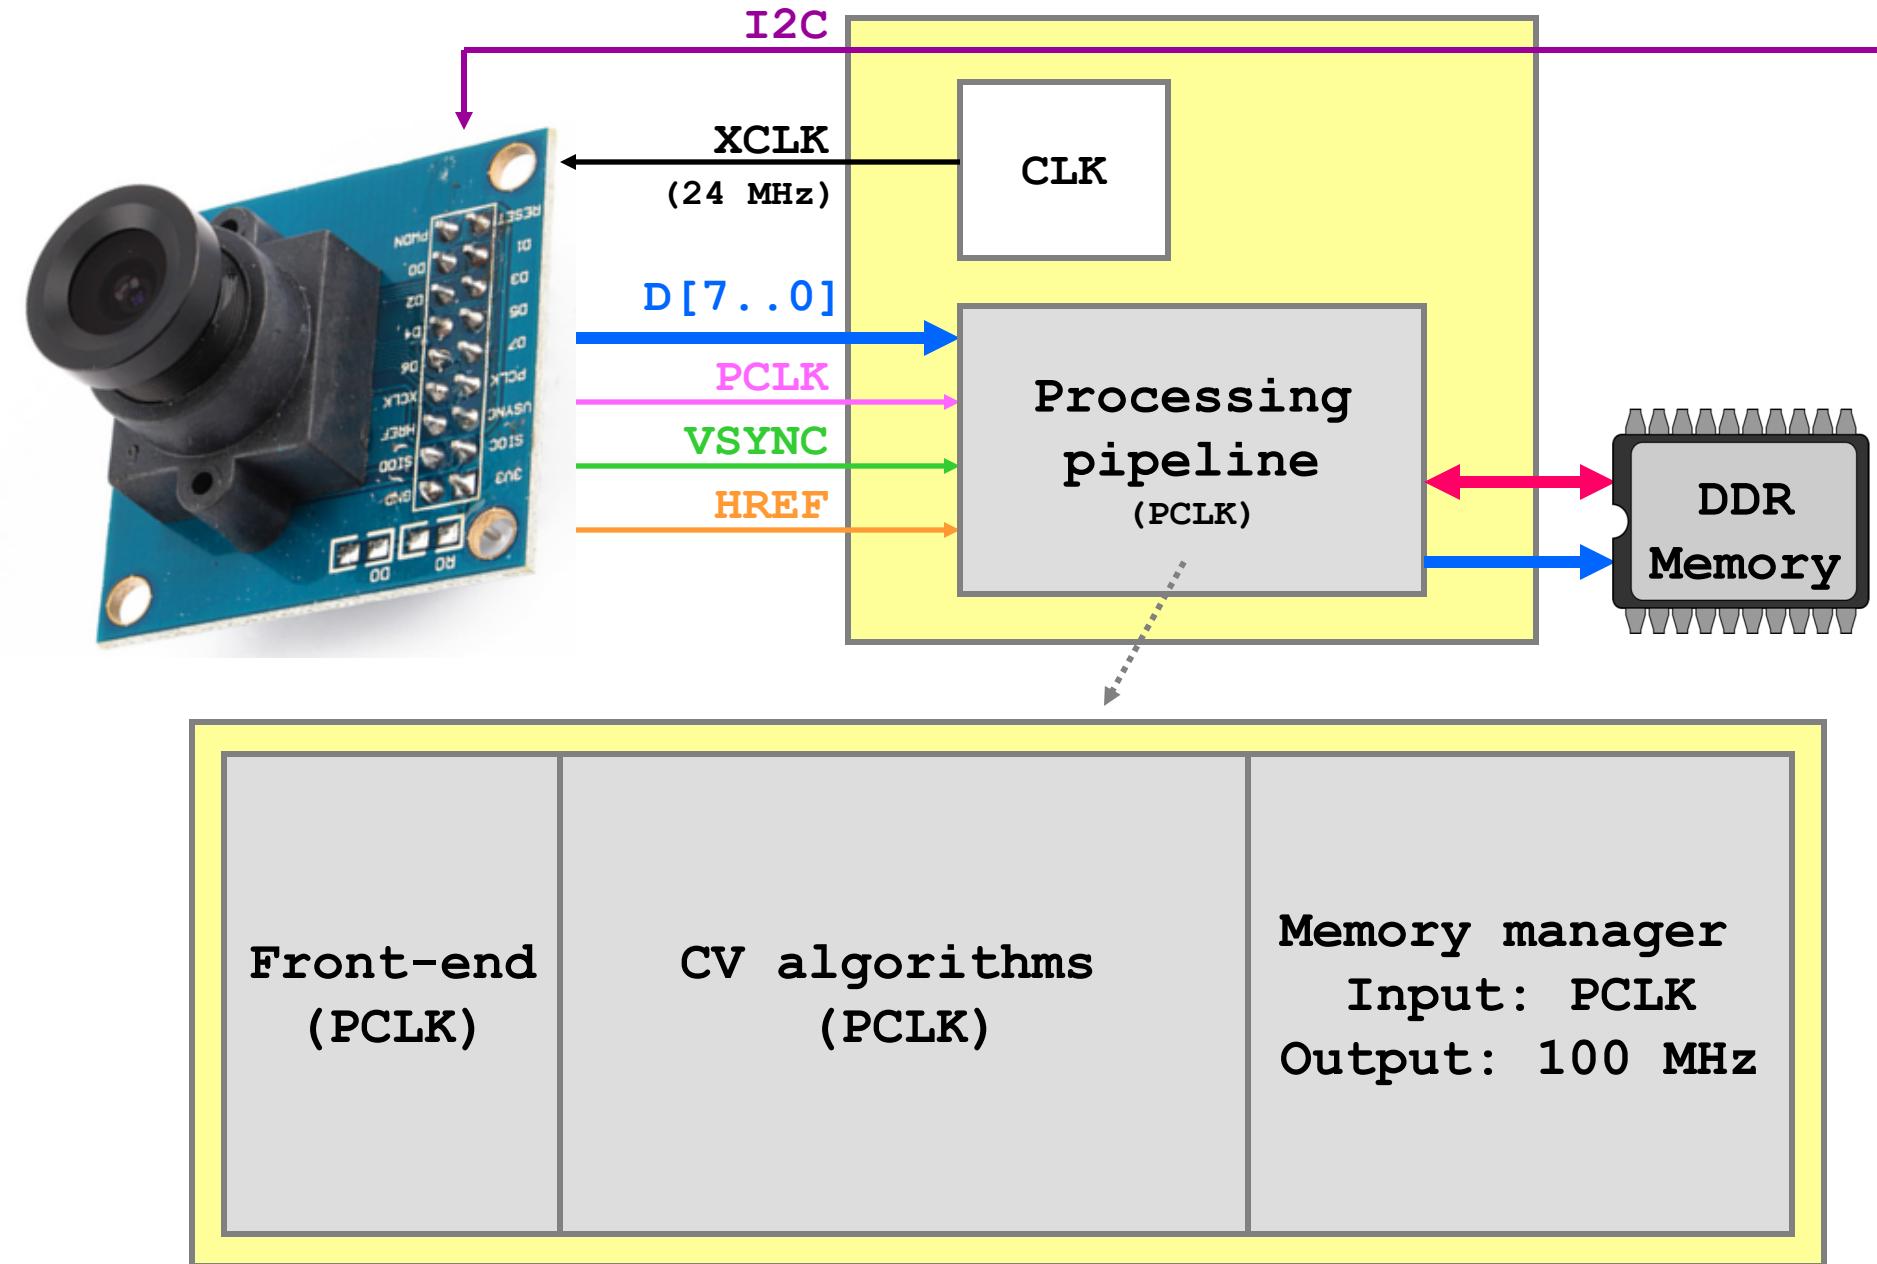
```
set_property IOSTANDARD LVCMOS33 [get_ports XCLK]
set_property PACKAGE_PIN W11 [get_ports XCLK]
.....
.....
.....
.....

set_property IOSTANDARD LVCMOS33 [get_ports PCLK]
set_property PACKAGE_PIN W10 [get_ports PCLK]

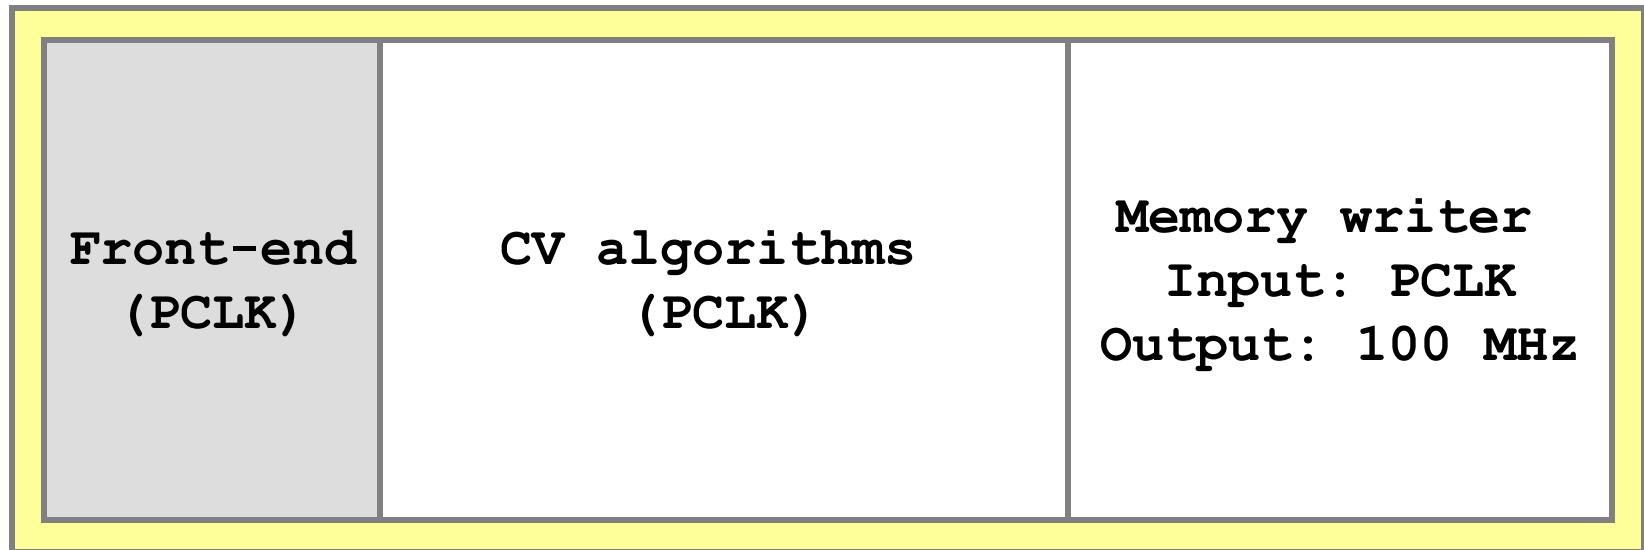
// only for PCLK - the PMOD pin is not CC (clock capable)
set_property CLOCK_DEDICATED_ROUTE FALSE [get_nets PCLK_IBUF]
```

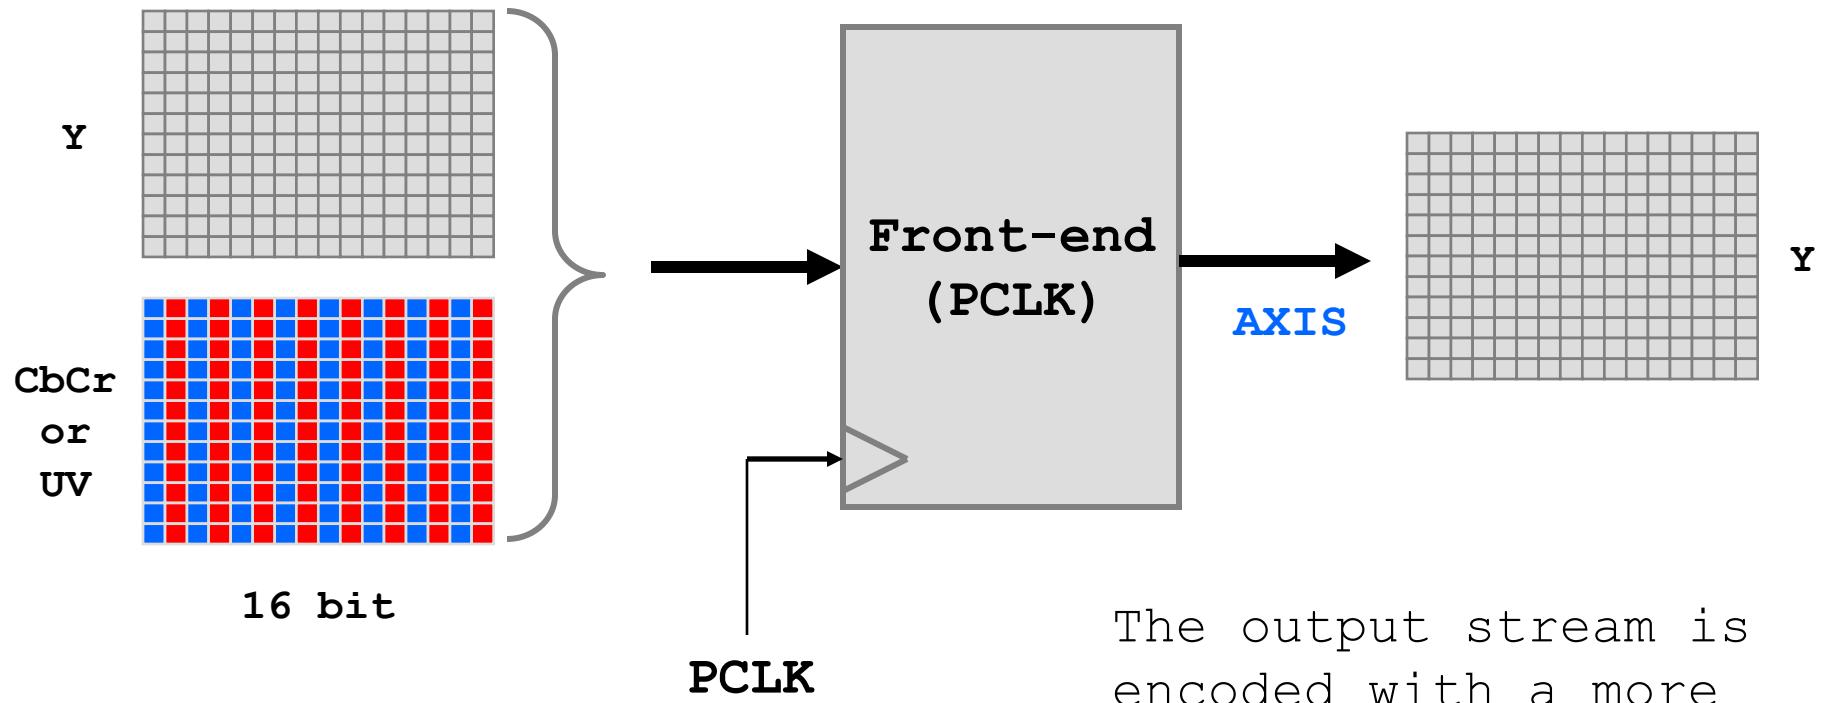


Processing pipeline



Processing pipeline: front-end (OV7670)





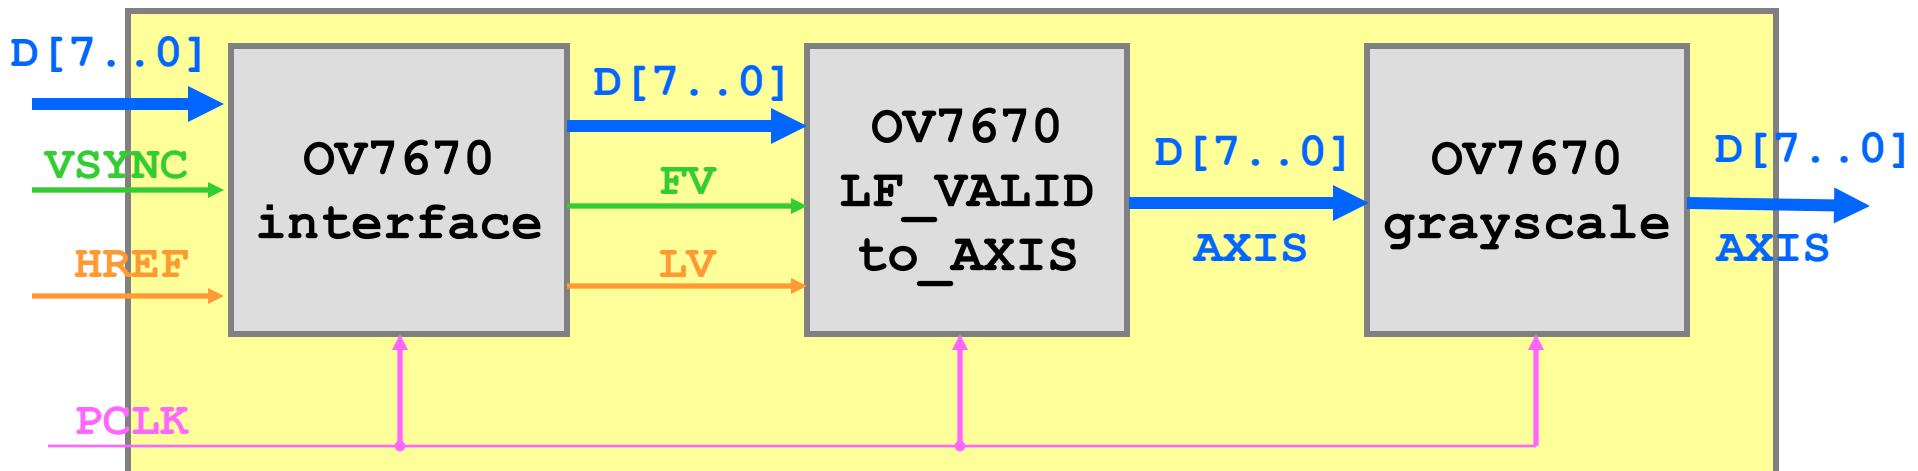
The input stream is encoded with the OV custom/proprietary protocol

The output stream is encoded with a more standard interface (AXIS).
From this point the processing pipeline is agnostic to the specific image sensor

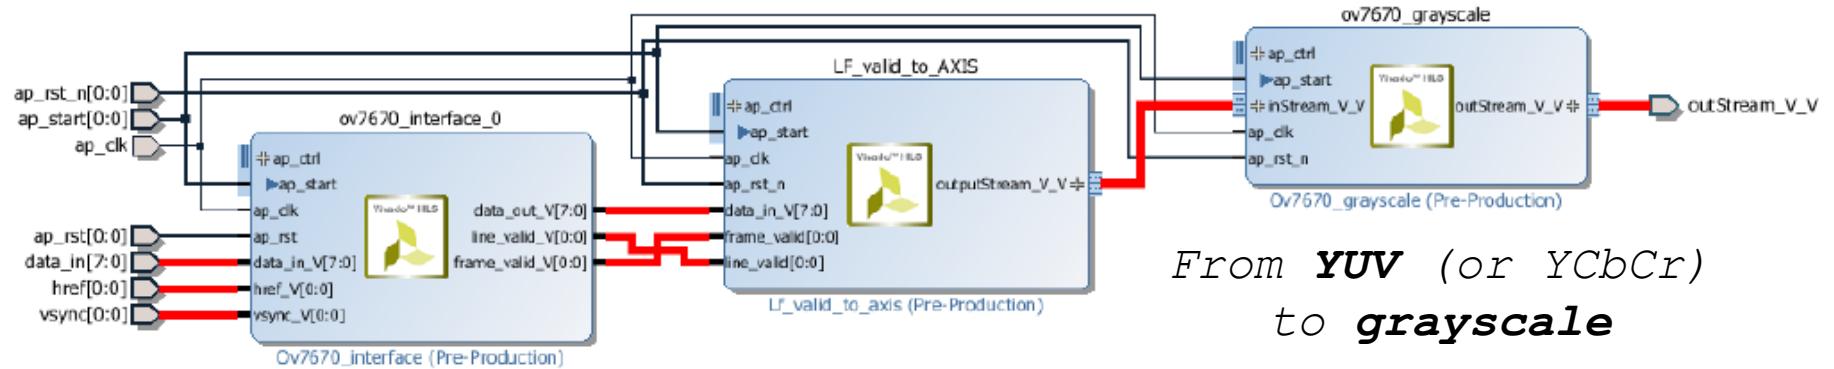
*Transcodes from
HREF-VREF to
LV-FV*

*Converts the LF-FV
format to AXIS*

*Picks only one out
of 2 bytes, the
Luma (Y)*

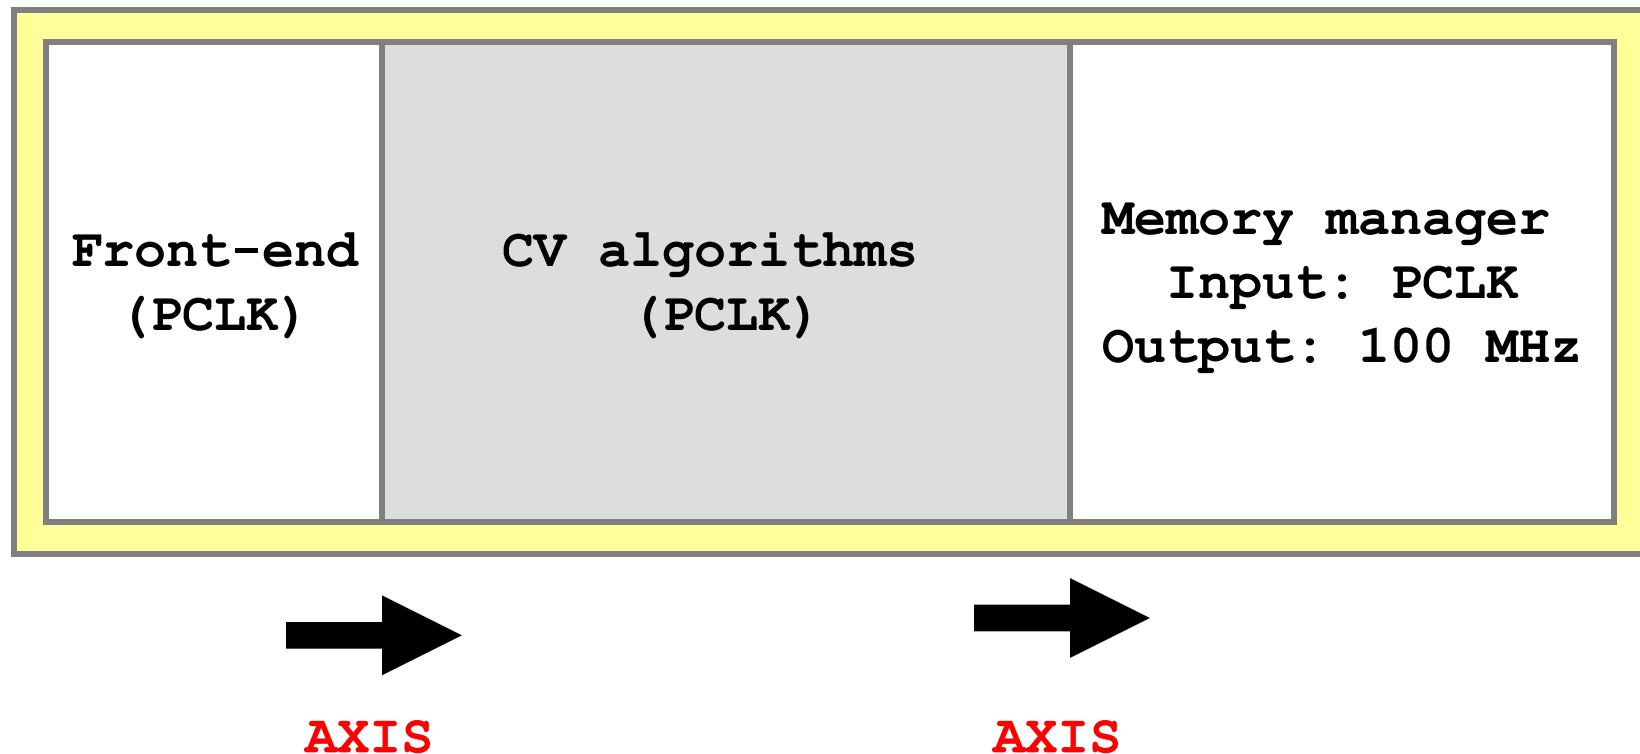


*From here, the encoding is hopefully
agnostic to the imaging sensor interface*



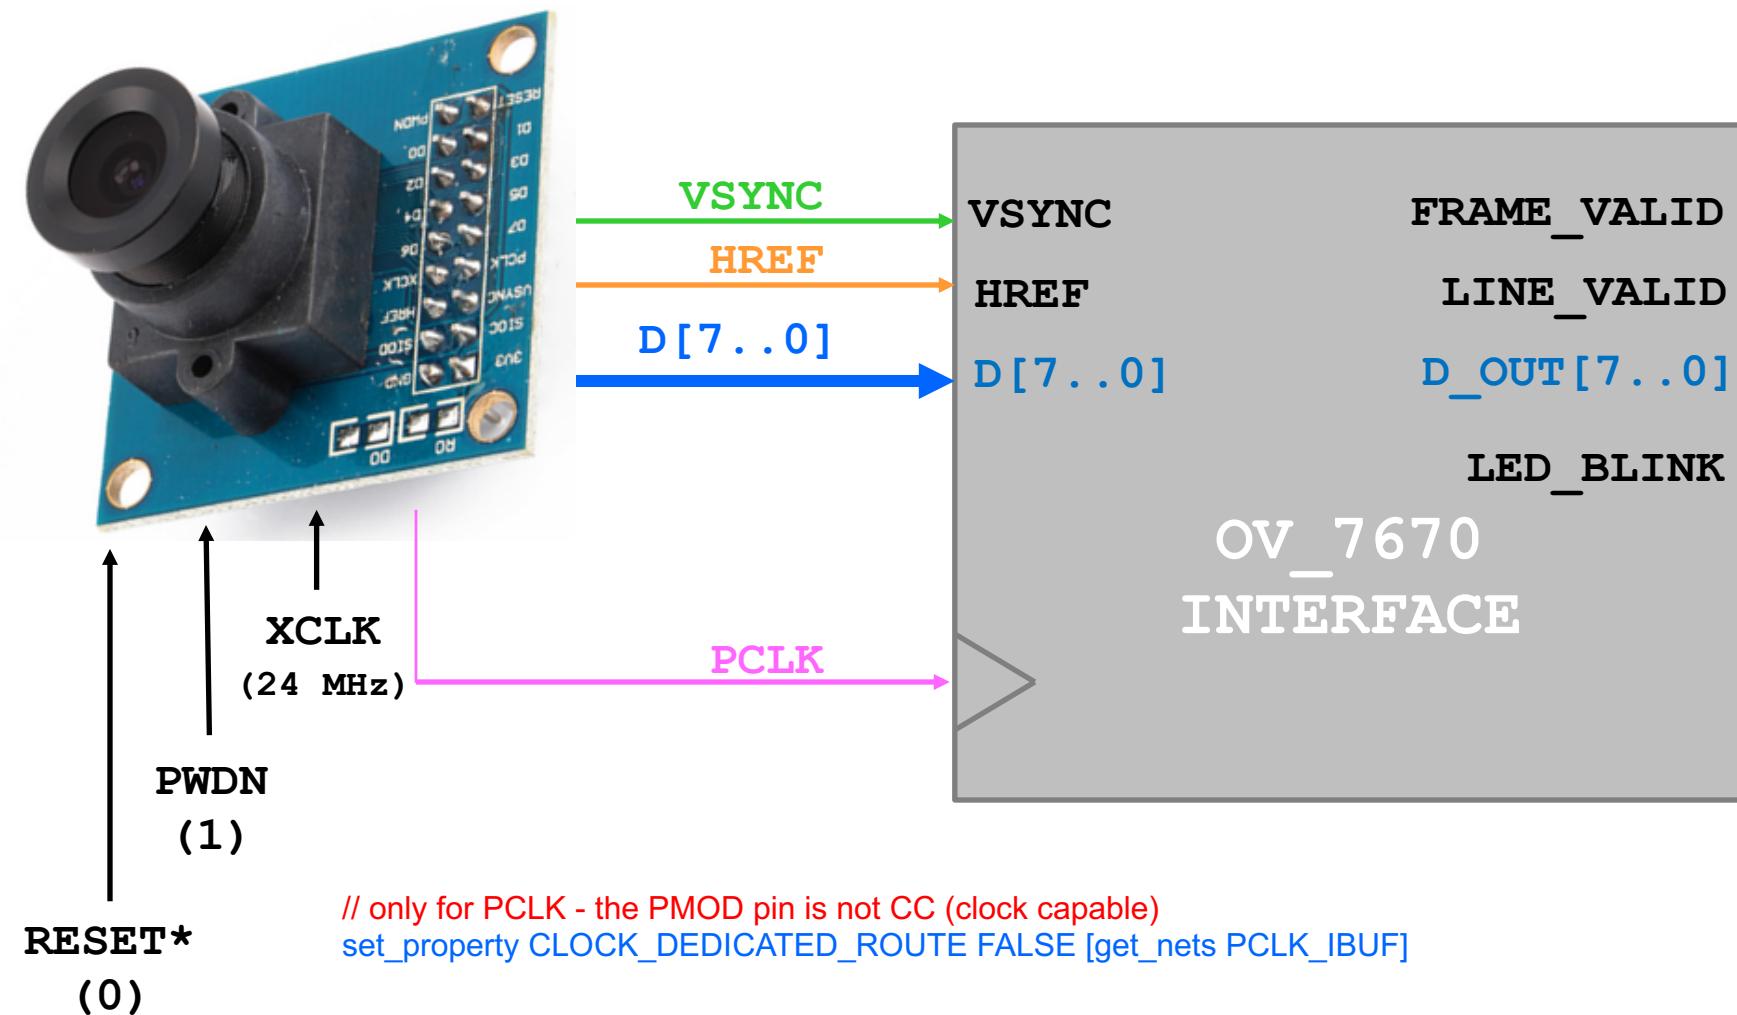
Processing pipeline: CV algorithms

This module is application specific and methodology to implement computer vision algorithms on FPGA will be discussed later.



Exercise

Generate **FRAME_VALID**, **LINE_VALID** and **D[7..0]** signals from an OV7670 video stream. Moreover, **LED_BLINK** inverts its value every 30 images



I2C and OV7670 (ARM side) 1/2

Connections between the PMOD connector JE and the OV7670 camera using the flat cable/wires:

SIOC (clock), purple, connected to **SCL** pin (**JE2**)

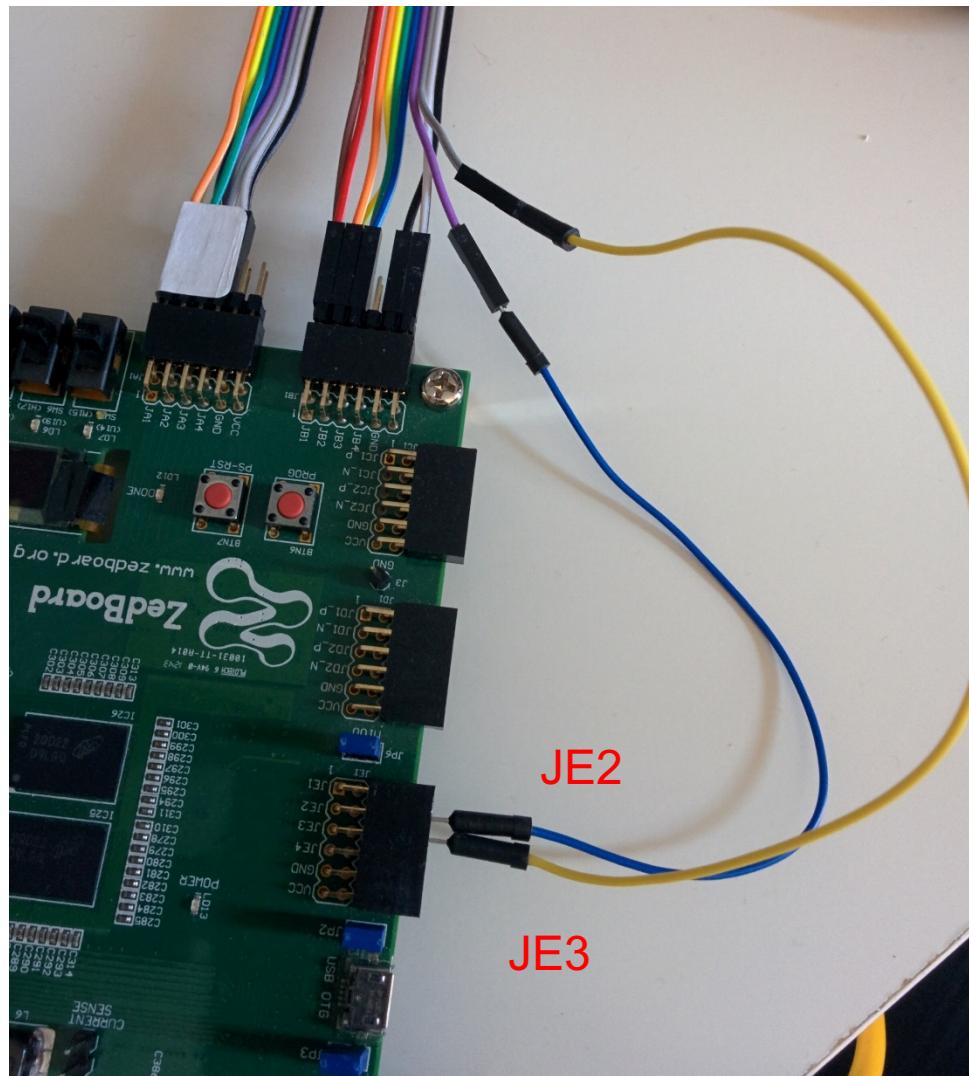
SIOD (data), gray, connected to **SDIOD** pin (**JE3**)

I2C and OV7670 (ARM side) 2/2

ov 7670

SIOD/SDA

SIOC/SCL

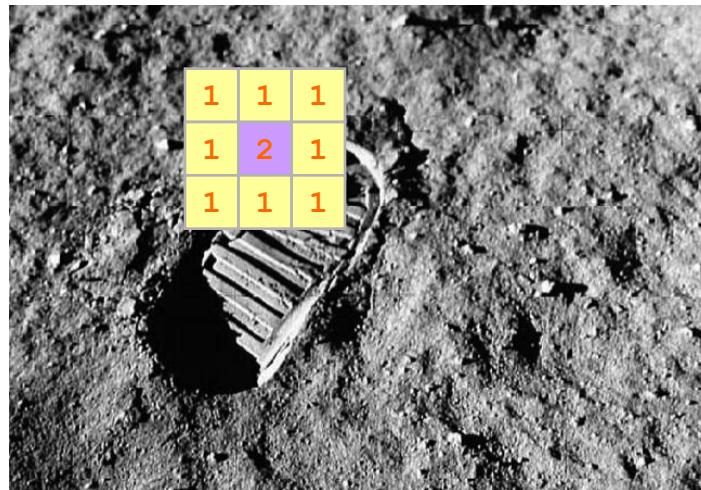


PMOD E PIN

JE3

JE2

Example: convolution filter



$$\begin{array}{ccc|c} 200 & 100 & 100 \\ \hline 100 & 100 & 200 \\ 100 & 200 & 100 \end{array} * \begin{array}{ccc} 1 & 1 & 1 \\ 1 & 2 & 1 \\ 1 & 1 & 1 \end{array} = 130$$

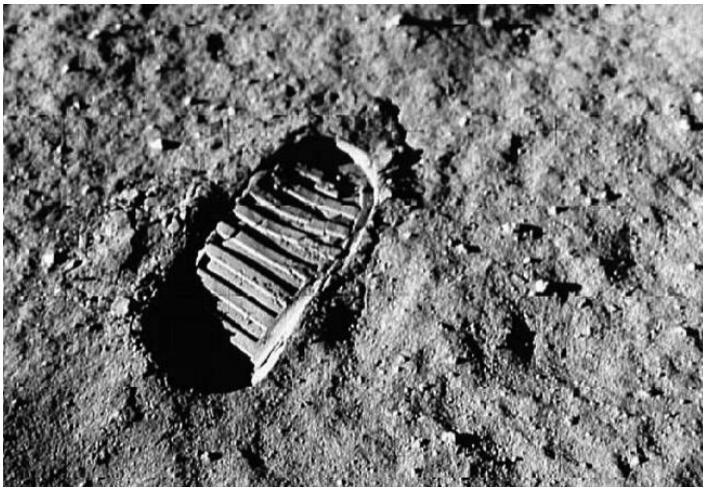
$$\begin{aligned} & (200 \cdot 1 + 100 \cdot 1 + 100 \cdot 1 \\ & + 100 \cdot 1 + 100 \cdot 2 + 200 \cdot 1 \\ & + 100 \cdot 1 + 200 \cdot 1 + 100 \cdot 1) / 10 = 130 \end{aligned}$$

What is a convolution kernel

A convolution kernel is a set of weights that allows obtaining, through the weighted sum of neighboring pixels, the value of the central output one

Precisely as in the previous examples, whose effect on the input image on the left is reported on the right image

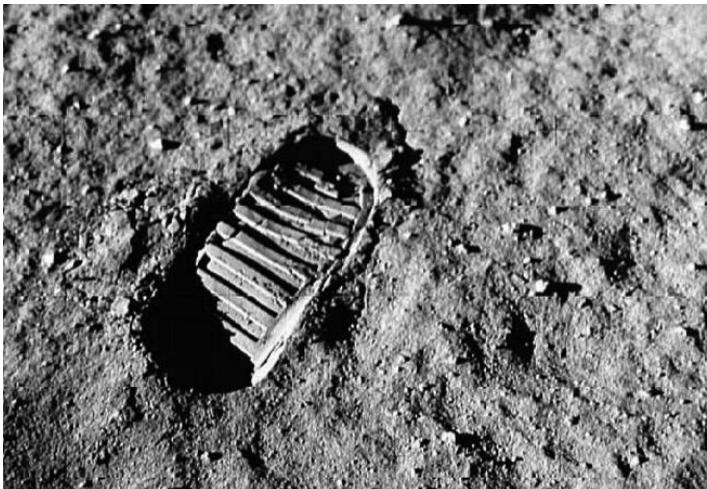
Kernel (weighted average)



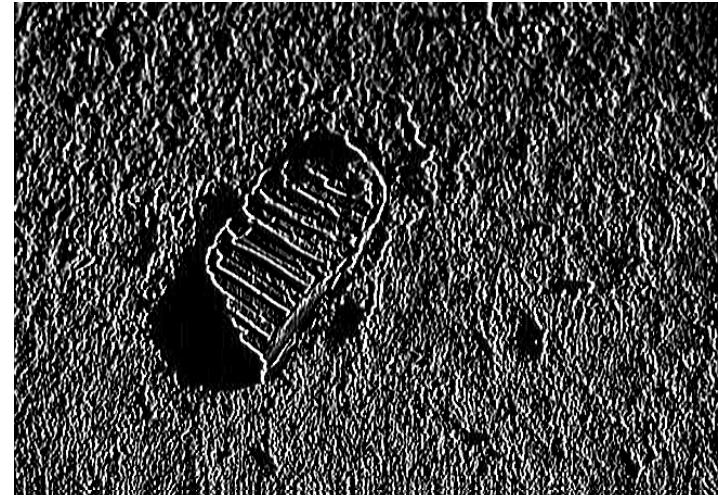
1	1	1
1	2	1
1	1	1



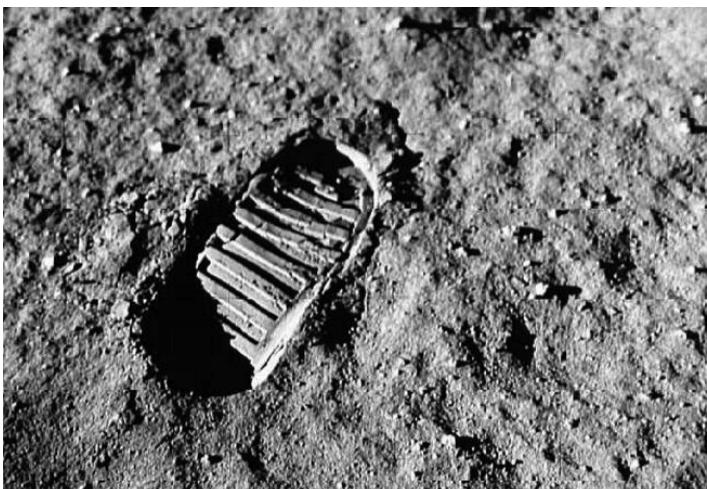
Kernel (Gradient x) : $|G_x|$



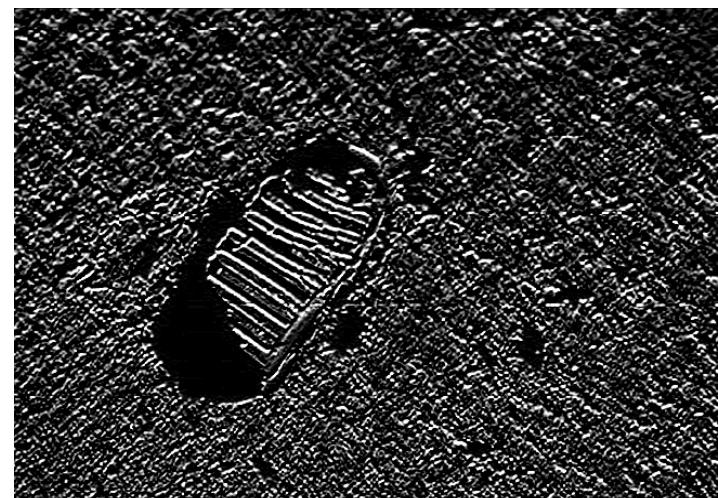
$$\begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$



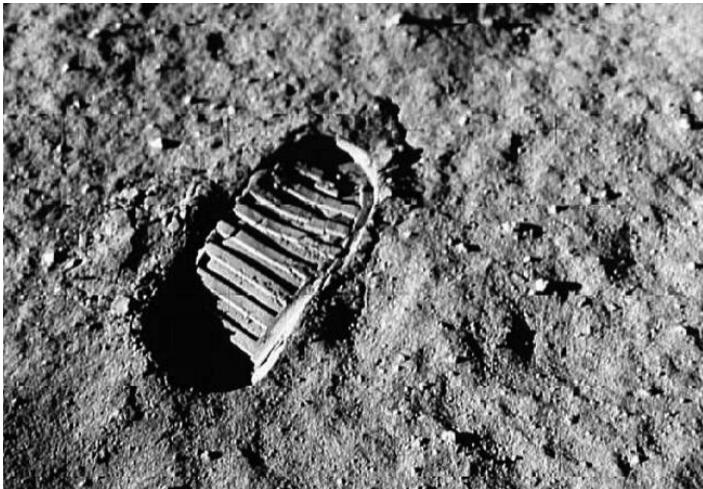
Kernel (Gradient y) : $|G_y|$



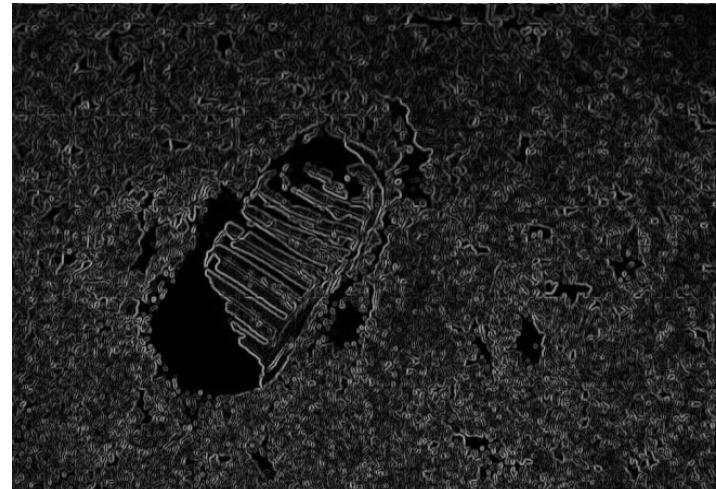
$$\begin{array}{|c|c|c|} \hline -1 & -2 & -1 \\ \hline 0 & 0 & 0 \\ \hline 1 & 2 & 1 \\ \hline \end{array}$$



Sobel (Magnitude)



$$G = \sqrt{(G_x)^2 + (G_y)^2}$$



- Changing the weights, we obtain different results
- In contrast to convolution, we can also perform on the patches non linear operations such as *bilateral filtering*, *median filtering* and so on.

These operations, especially linear ones, are at the core of Convolutional Neural Networks (CNNs). In this case, the weights are not set beforehand but adjusted during the training of the networks.

Example: weighted average 1/10

- Each output pixel is the weighted sum of its nearby points within the patch (3x3 in this example)

Intorno di Pin0

Pin1	Pin2	Pin3
Pin4	Pin0	Pin5
Pin6	Pin7	Pin8

Kernel dei pesi

1	1	1
1	2	1
1	1	1

Concolution
kernel

(weighted average)

DIM_KERNEL(K_H,K_W)

```
Pout0 = (Pin1 + Pin2 + Pin3 + Pin4 + Pin5 +  
Pin6 + Pin5 + Pin8 + (2 * Pin0)) / 10
```

- The HLS design is implemented in a **PIPELINE** fashion, at each clock the modulo reads an input pixel and compute the output (ie, **stream** processing)

Example: weighted average 2/10

- We need to specify which is the interface to handle input and output values (ie, pixels)
 - For instance, we can adopt the AXI stream protocol, deploying the **AXIS** directive in Vivado HLS
-
- It is an interface enabling sequential communication, and it is not *memory mapped* (ie, no address specified)
 - At the lowest level, data is buffered and transmitted using the *handshake* protocol

```
#pragma HLS INTERFACE axis port=in_img  
#pragma HLS INTERFACE axis port=out_img
```



Example: weighted average 3/10

- The “top function” processes the input image and produces the filtered image according to the kernel weights
- Images can be modeled as 2D arrays

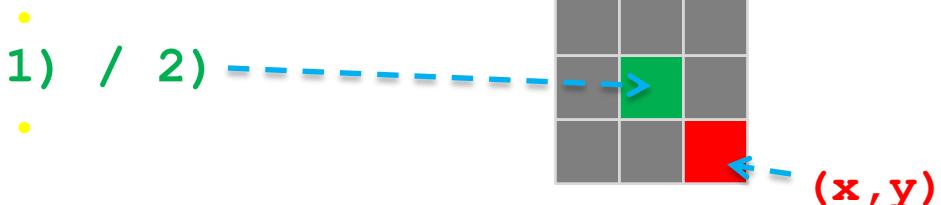
```
void filtroMedia (ptype    in_img[HEIGHT_IMG][WIDTH_IMG],  
                  ptype    out_img[HEIGHT_IMG][WIDTH_IMG]  
                )
```

- The C code structure consists of two for cycles: one (outer) for rows and one (inner) for columns

```
34 Loop_row: for(int row = 0; row < HEIGHT_IMG + Km_H; row++){  
35     Loop_col: for(int col = 0; col < WIDTH_IMG + Km_W; col++)
```

- At each iteration of the inner loop: the module receives the pixel at coordinates **(x,y)** and computes the output pixel at coordinate:

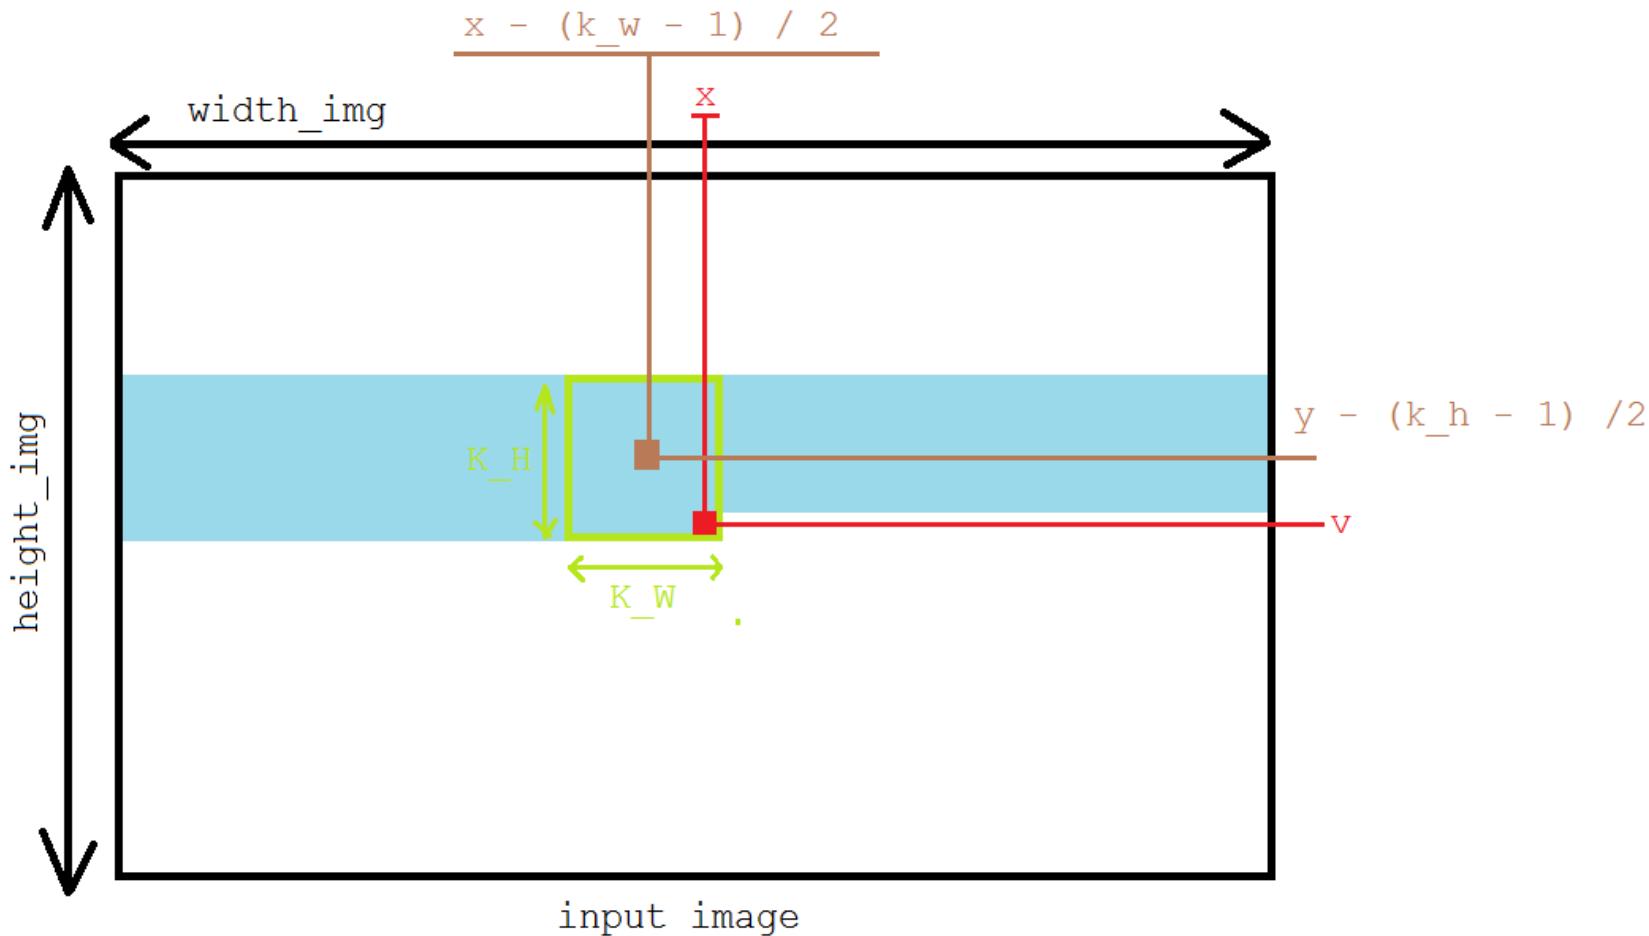
$(x - (k_w - 1)/2, y - (k_h - 1) / 2)$



Example: weighted average 4/10

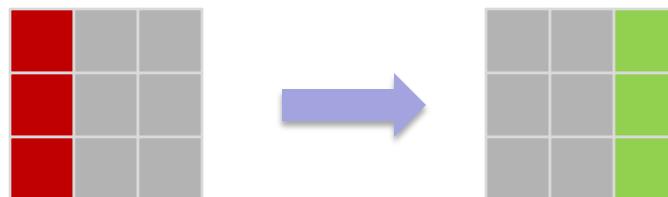
If a subset of previous pixels has been buffered, when the pixel (x, y) is fed to the module, we can process the pixel at

$$(x - (k_w - 1)/2, y - (k_h - 1) / 2)$$



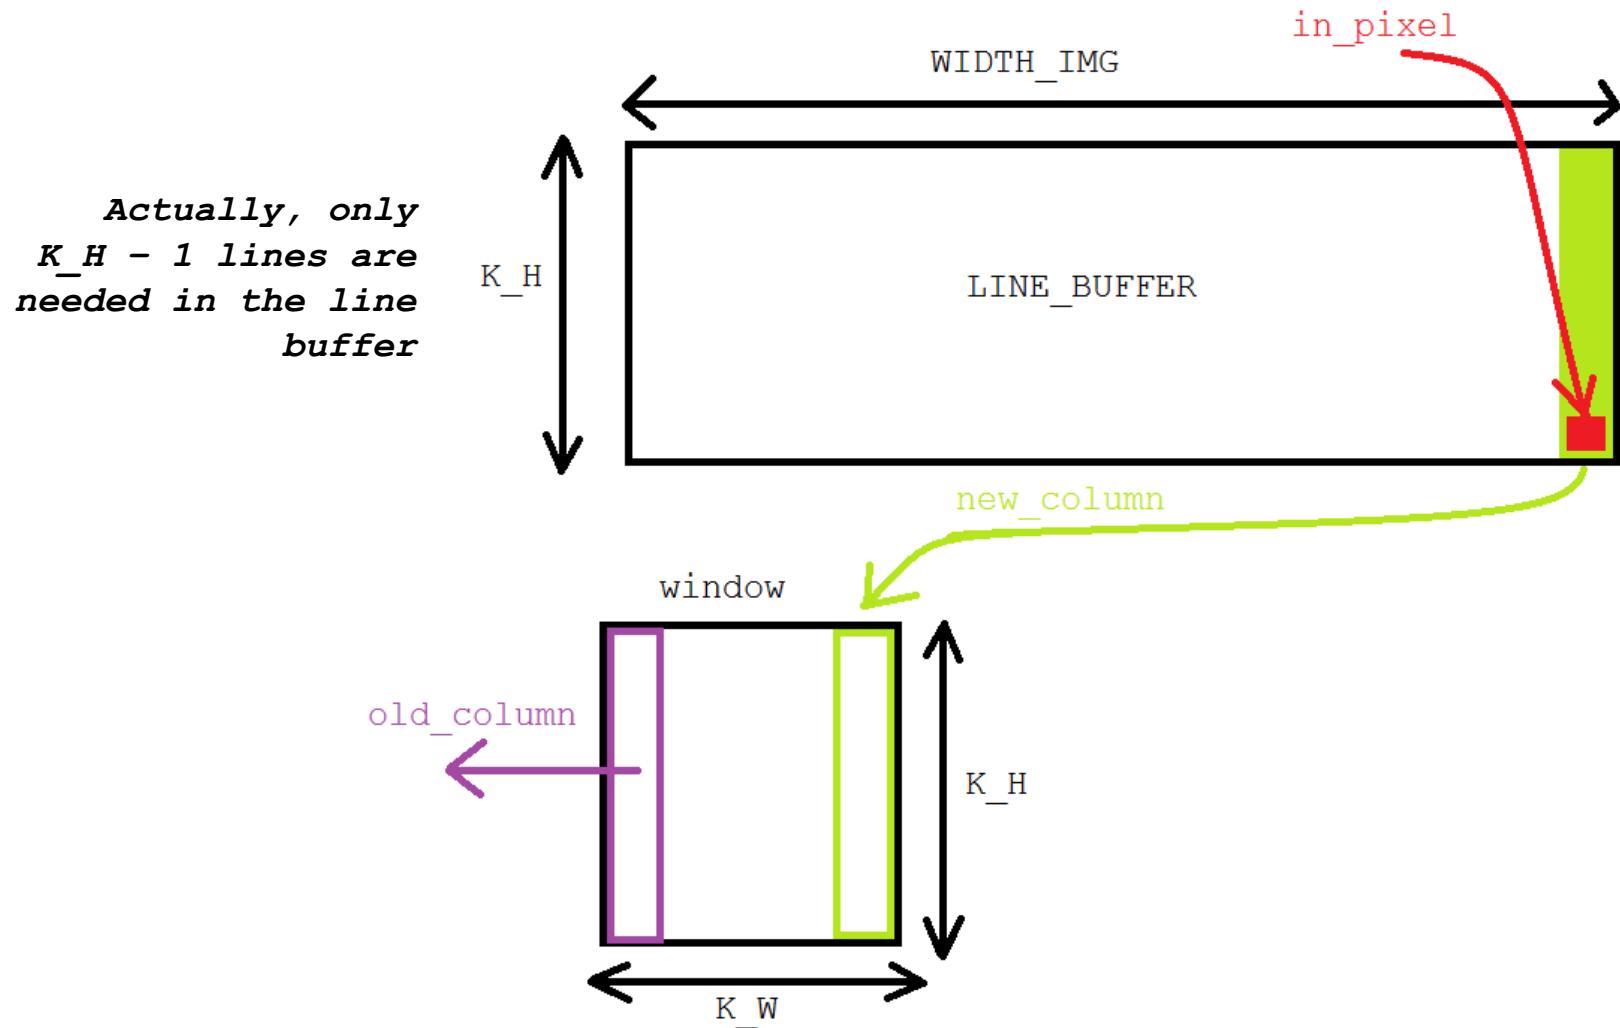
Example: weighted average 5/10

- Each pixel of the input image is received starting from the top left one, line after line
- The lines are buffered in **line buffer** (in C it is a 2D array). In each entry of the array is buffered a line. The size of the line buffer is: **WIDTH_IMG x (K_H - 1)**
- In a data structure of smaller size (**window**) are stored the pixels needed to compute the weighted average. The size of such a window is: **K_H x K_W**
- After each iteration, the least recent column is replaced with a one (from **line buffer**) in the **window** data structure



Example: weighted average 6/10

Data structures line buffer and window:



Example: weighted average 7/10

- The `window` data structure, once initialized, is fed to the `media_pixel` function aimed at computing the weighted average. Purposely, all elements of the `window` data structure need to be accessed concurrently.
- The outcome of such a computation (ie, the weighted average) is the output of the *top function*.

```
ap_uint<8> out_temp = 0;
//function invocation
mediaPixel (window, &out_temp);
//write output
out_img[row - Km_H][col - Km_W] = out_temp(7,0);
```

```
//// kernel function
void mediaPixel (ptype y_window[K_H][K_W], ptype *out)
{
    ap_int<20> out_temp = 0;

    Edge_i: for(int i=0; i < K_H; i++)
        Edge_j: for(int j = 0; j < K_W; j++)
            out_temp = out_temp + y_window[i][j]*M[i][j];

    out_temp = out_temp / SumF;

    *out = out_temp(7,0);

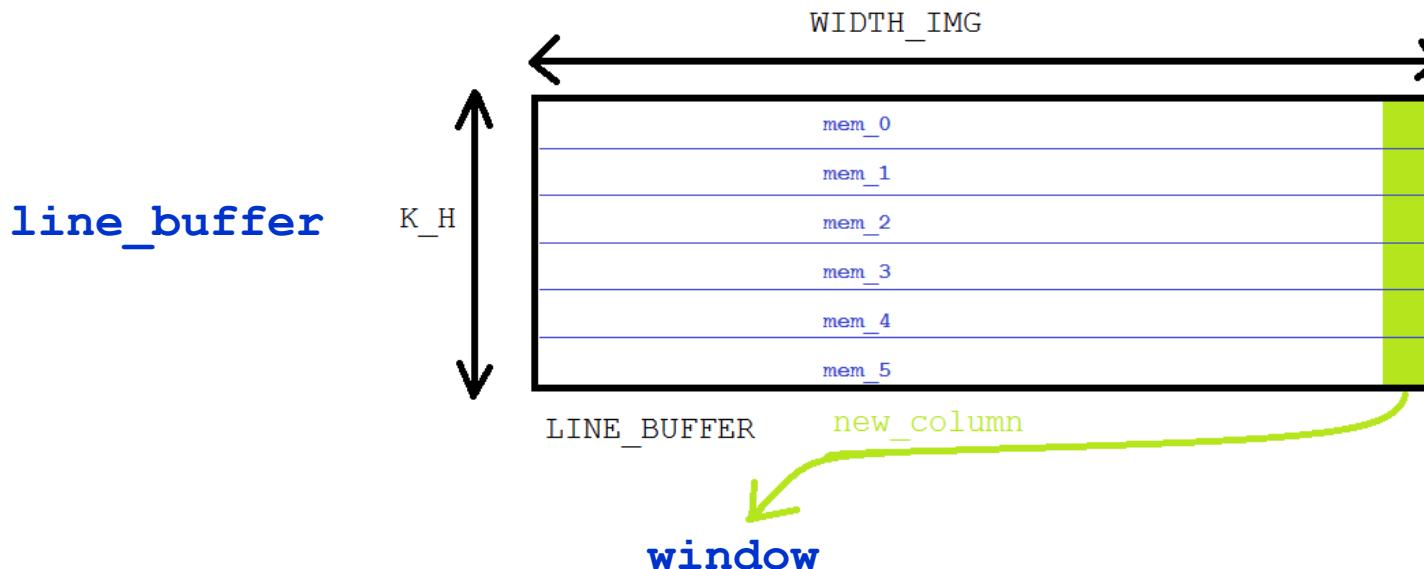
}
```

Example: weighted average 8/10

- For an efficient implementation, at each clock clock all the elements of the `line_buffer` need to be accessed concurrently. Therefore, the `line_buffer`, is partitioned in independent rows (ie, each row is mapped on a different memory device).

```
set_directive_array_partition -type complete -dim 1 "filtroMedia" line_buffer_temp
```

- Following this strategy, in a single clock cycle all the elements of the column in green can be read and then copied in the `window` data structure



Example: weighted average 9/10

- At each clock is computed a weighted average. Therefore, all the elements of the **window** data structure need to be accessed concurrently implying that each element of **window** is a register
- Directive partition, -dim 0 => each dimension of the array is partitioned (hence, each element in a register)

```
set_directive_array_partition -type complete -dim 0 "filtroMedia" window
```

Window			Kernel weights		
FF_0	FF_1	FF_2	1	1	1
FF_3	FF_4	FF_5	1	2	1
FF_6	FF_7	FF_8	1	1	1

```
Pout0 = (Pin1 + Pin2 + Pin3 + Pin4 + Pin5 +  
Pin6 + Pin5 + Pin8 + (2 * Pin0)) / 10
```

Example: weighted average 10/10

- Outcome of the implementation of the weighted average filter on a Spartan 6 45 FPGA:

Area Estimates					
Summary					
	BRAM	DSP48A	FF	LUT	SLICE
Component	-	-	-	-	-
Expression	-	2	0	216	-
FIFO	-	-	-	-	-
Memory	2	-	-	-	-
Multiplexer	-	-	-	73	-
Register	-	-	176	-	-
Total	2	2	176	289	0
Available	116	58	54576	27288	6822
Utilization (%)	1	3	~0	1	0

Details

Memory

	Words	Bits	Banks	W*Bits*Banks	BRAM
line_buffer_0_V_U	640	8	1	5120	1
line_buffer_1_V_U	640	8	1	5120	1
Total	1280	16	2	10240	2

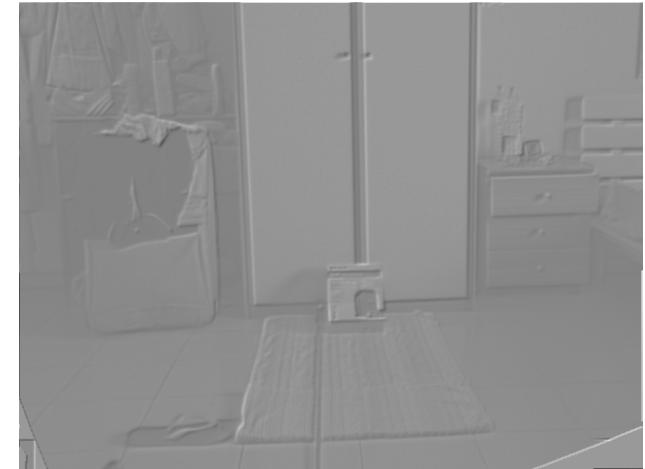
Resource Usage		
	VHDL	Verilog
SLICE	65	-
LUT	143	-
FF	152	-
DSP	2	-
BRAM	2	-

Only 2 BRAM for buffering
the 3 lines

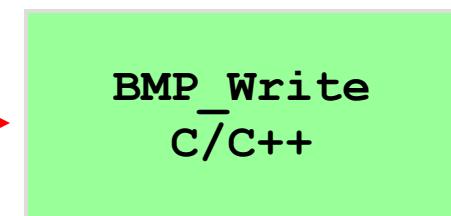
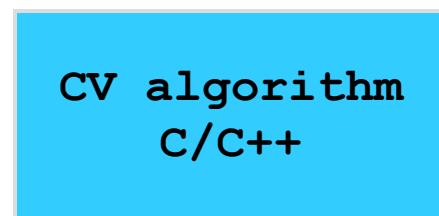
- Further optimization strategies described in the slides «08_Convolution_filters.pdf»

Testbench in C/C++ (HLS)

```
#include "ap_bmp.h"
```



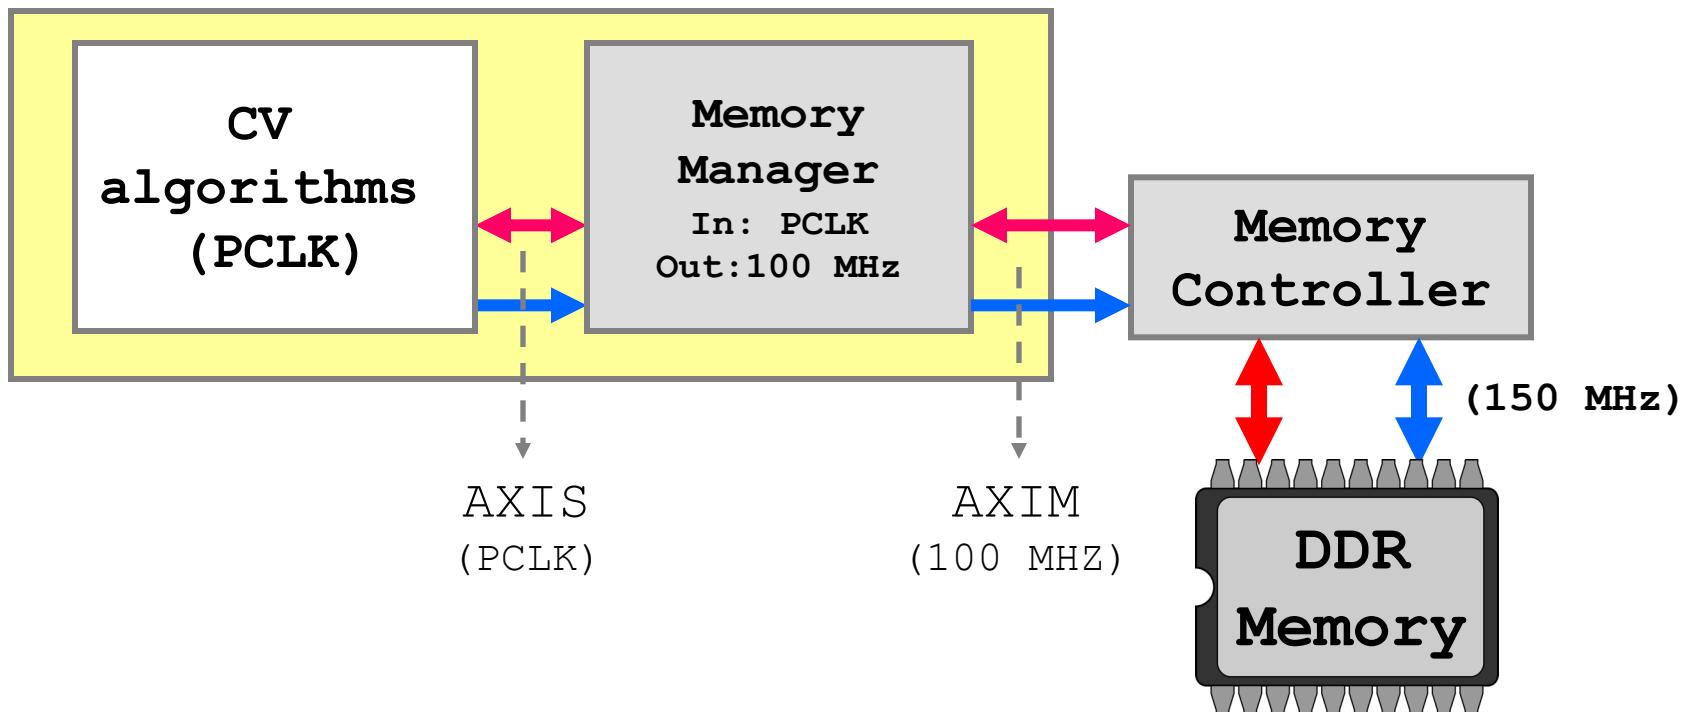
BMP
↓
RGB format



BMP
↑
RGB format

Memory transfer manager

- It is (upstream) connected (AXIS) to the CV pipeline and (downstream) to the memory controller (HP0) of the Zynq/PS with a 64 bit AXI master (AXIM) interface
- Writes data in the frame-buffer (in DDR memory)
- Memory addresses (frame buffer) can be configured by the ARM processor via AXI_LITE
- A valid address for the frame buffer is 0x10000000

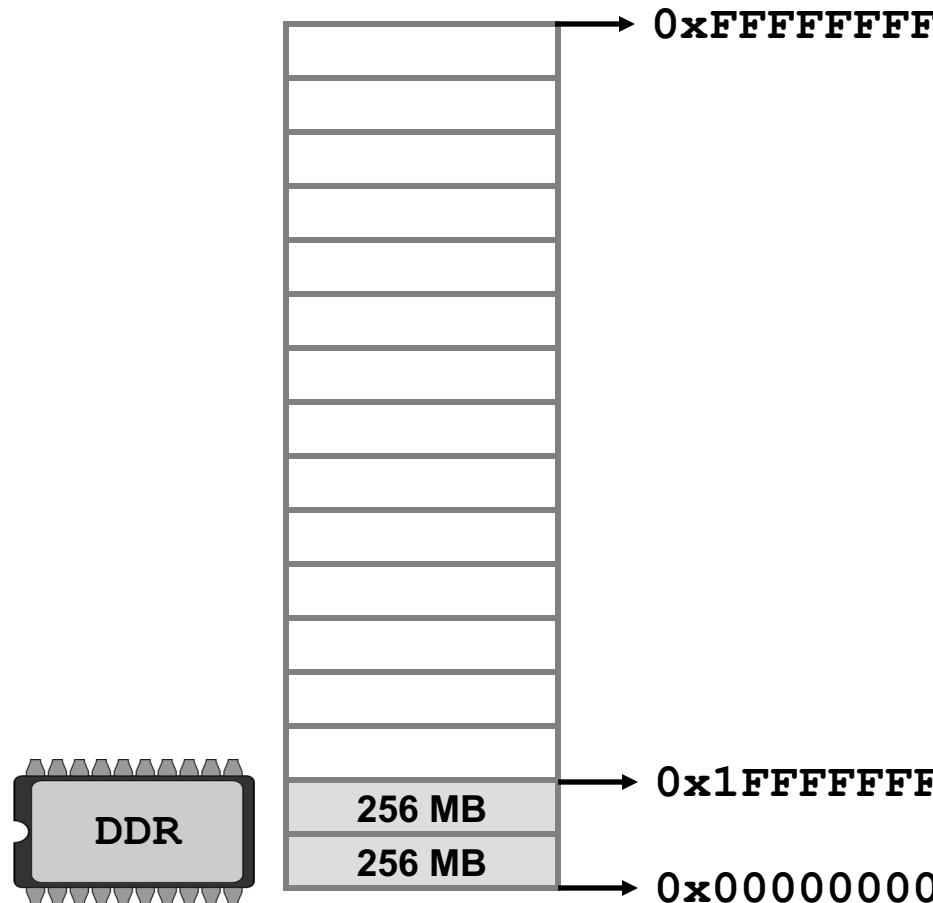


Zedboard: address space and available DDR memory

The Zynq processor has a 4 GB address space

The Zedboard contains 2x256 MB DDR3 memory devices

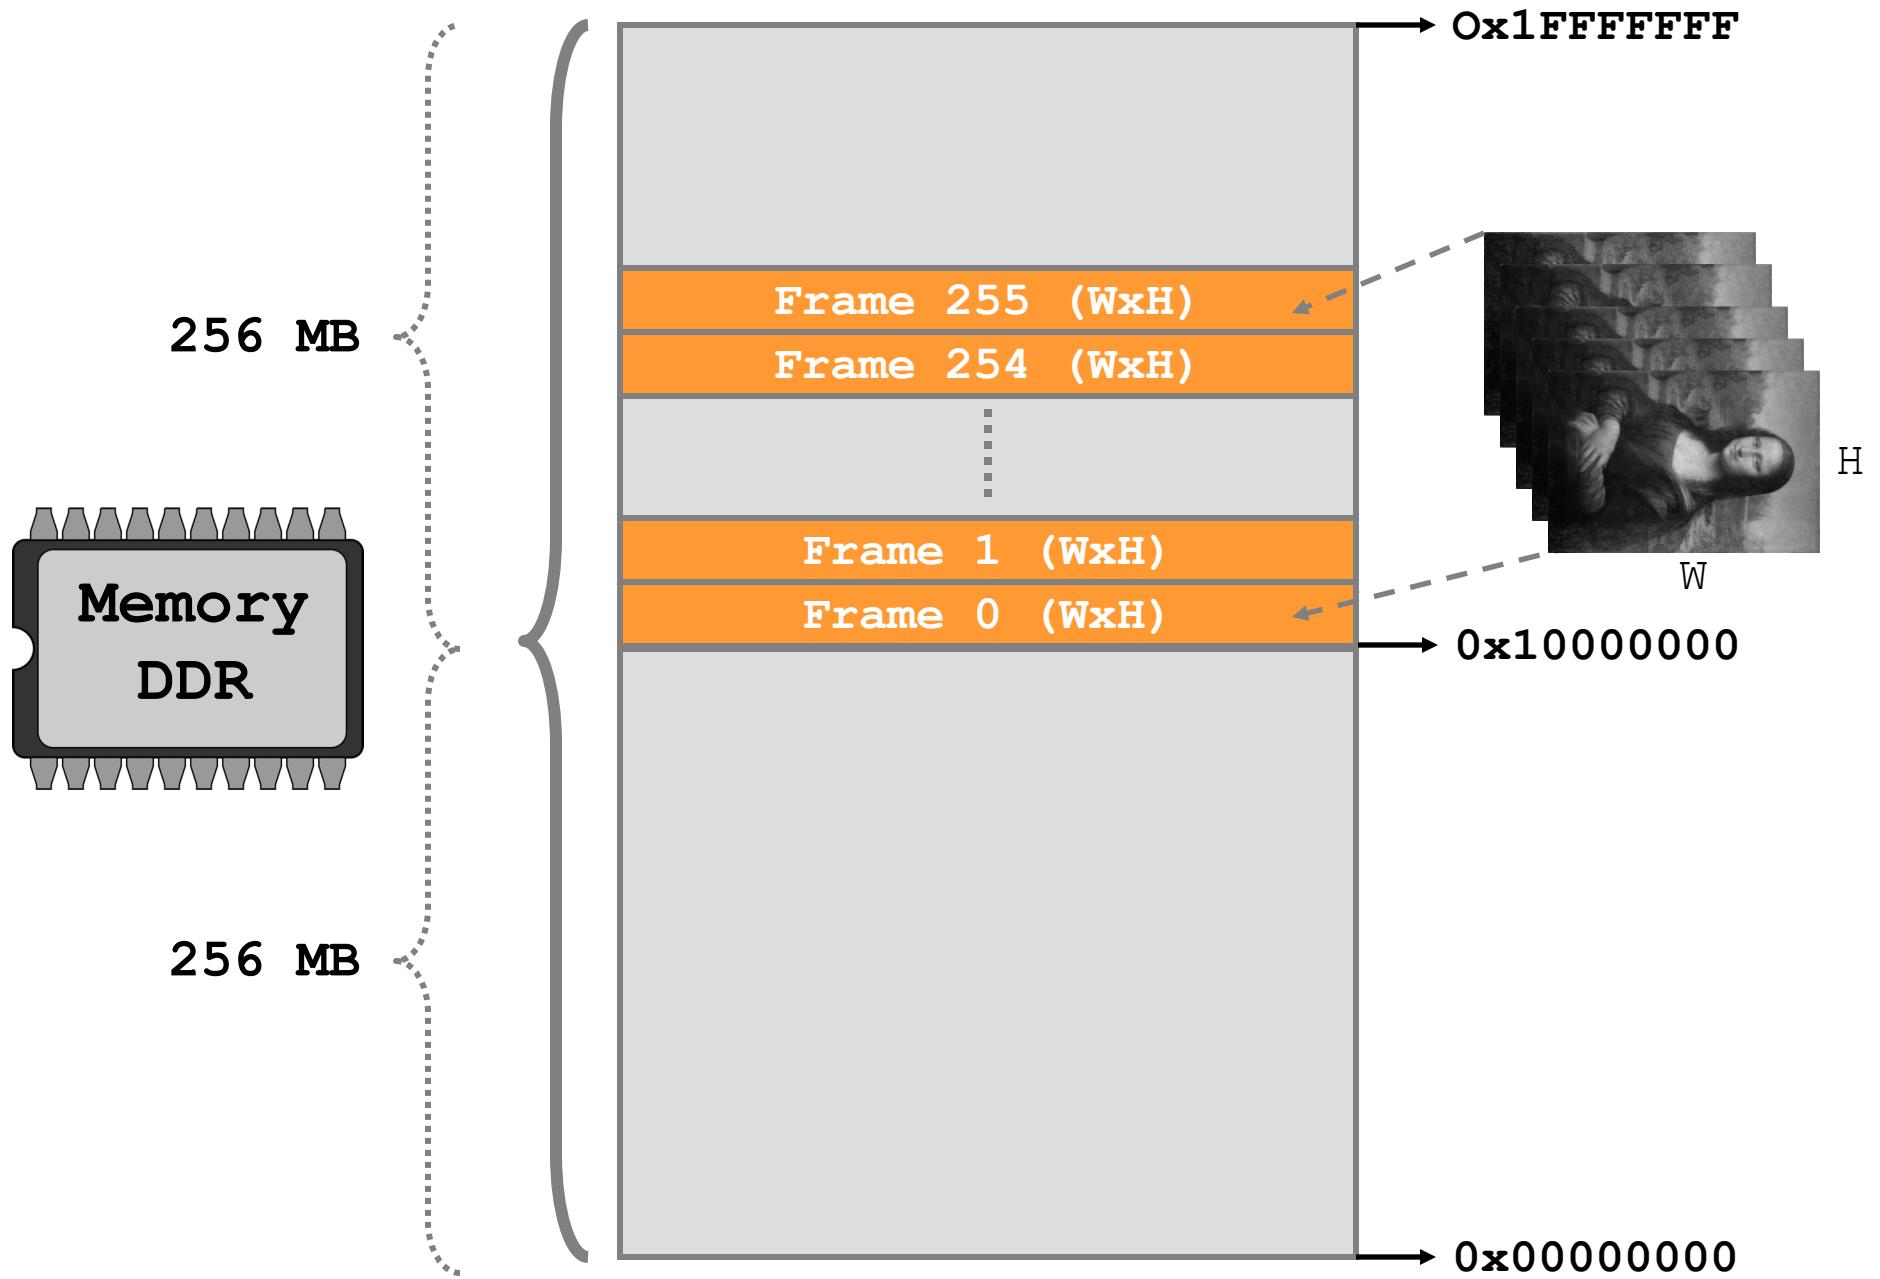
The overall DDR memory available (Zedboard) is 512 MB



The frame-buffer

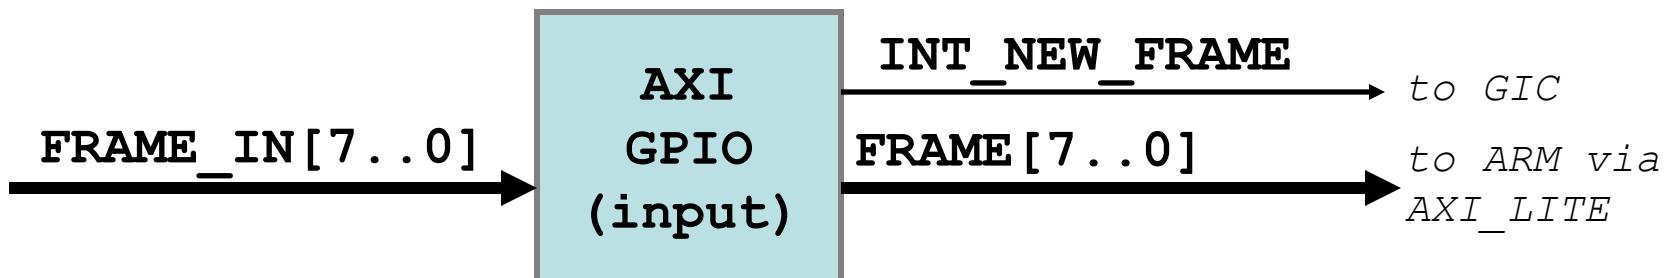
- The design includes a frame-buffer in DDR
- The number of frames is configurable (via AXI-Lite)
- Max 256 frames (typically 8)
- Images are stored in consecutive locations starting from the base of the frame-buffer
- A valid address for the frame buffer is 0x10000000
- Image size is configurable
- The CPU efficiently sends to the host (e.g., via ethernet) the last images stored in the frame buffer by means of an interrupt-driven approach
- The frame-buffer is also crucial for displaying images on a standard monitor (e.g., by means of the standard VGA interface)

The next slide depicts the structure and addresses of the frame-buffer in the DDR memory



The Frame-Index

- This module contains the index of the last frame completely written in memory (i.e., frame-buffer)
- Once an image is completely written in memory, the frame-indexer is filled, by the memory writer, with the index of the frame-buffer element
- This operation triggers an interrupt for the ARM processor
- The interrupt handler reads via AXI-Lite the index and performs appropriate operations accordingly
- The Frame-Index is mapped as peripheral in the address space and based on a standard AXI_GPIO module



Interrupt handling

In the basic embedded CV pipeline outlined the handler manages the pending interrupts according to the strategy configured in the GIC.

In the basic project described so far, for each frame the new images is simply transmitted using the ethernet port

It is worth to note that, with I2C, there is at least another interrupt source

The current project is compatible with *Standalone OS*, often referred to as *Bare-metal OS*

However, the same project can be used with *Linux OS* without any modification to the bitstream

UDP streaming (ARM-Standalone)

Images, are transmitted as UDP packets using the LwIP - Lightweight IP - framework. For this purpose, it is important to start with the "lwIP Echo Server" application template available in Vivado SDK.

In the next slides some images acquired with a software PC client based on OpenCV connected to the Zynq camera via ethernet

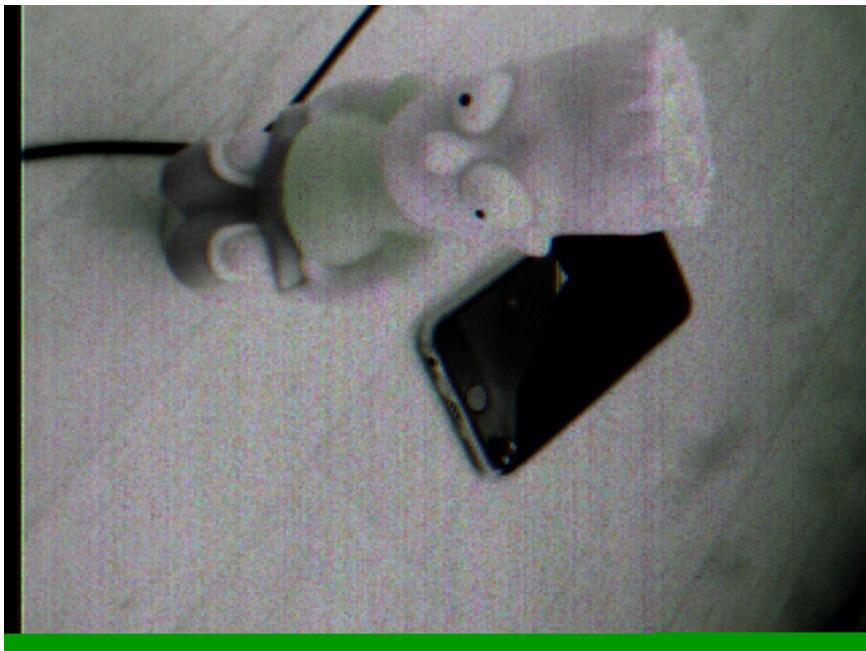
With three images (LUMA, CHROMA and output of the Zynq processing) the bandwidth at VGA resolution is 27 MB/s (640x480x3x30 MB/sec).

Thus, with three images a Gigabit ethernet adapter (or a USB 3.0 to Gigabit ethernet adapter connected to a USB 3.0 port) is required. With a single image at 30 fps (about 9 MB/s) a 100 Mbit (or a USB 2.0 to ethernet adapter) has enough bandwidth.



Luma:

**Left, OV7670 default configuration at startup
Right, after I2C configuration**



Color:

**Left, OV7670 default configuration at startup
Right, after I2C configuration**



Luma (Y)



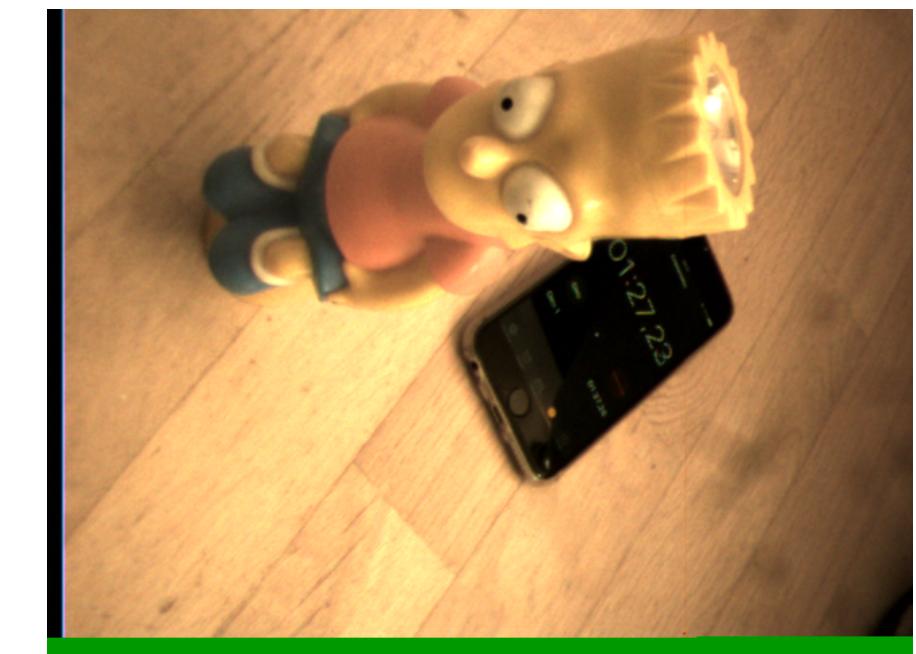
Zynq output

Zynq PL processing:

The OV7670 input stream is processed in real-time by a computer vision pipeline implemented in the PK (i.e., FPGA) with HLS tools (i.e., Vivado HLS)

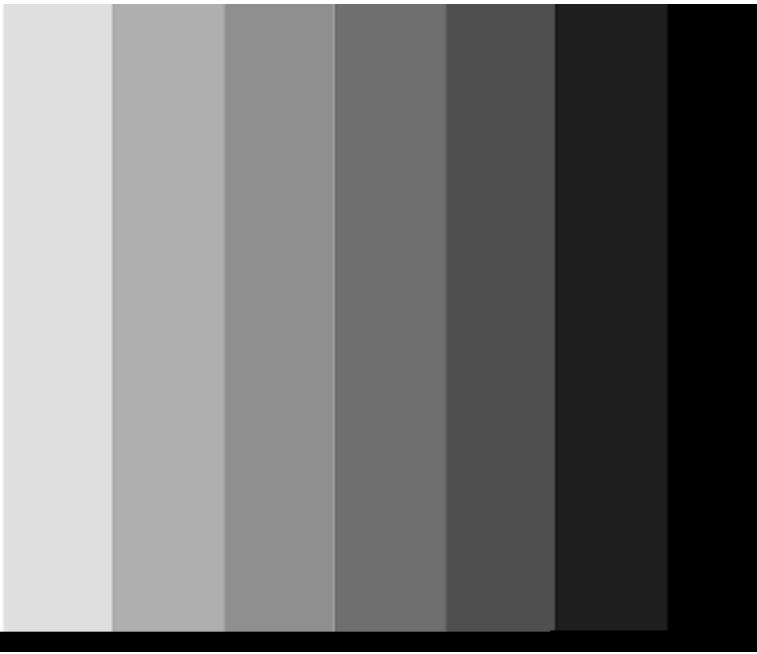


Luma (Y)

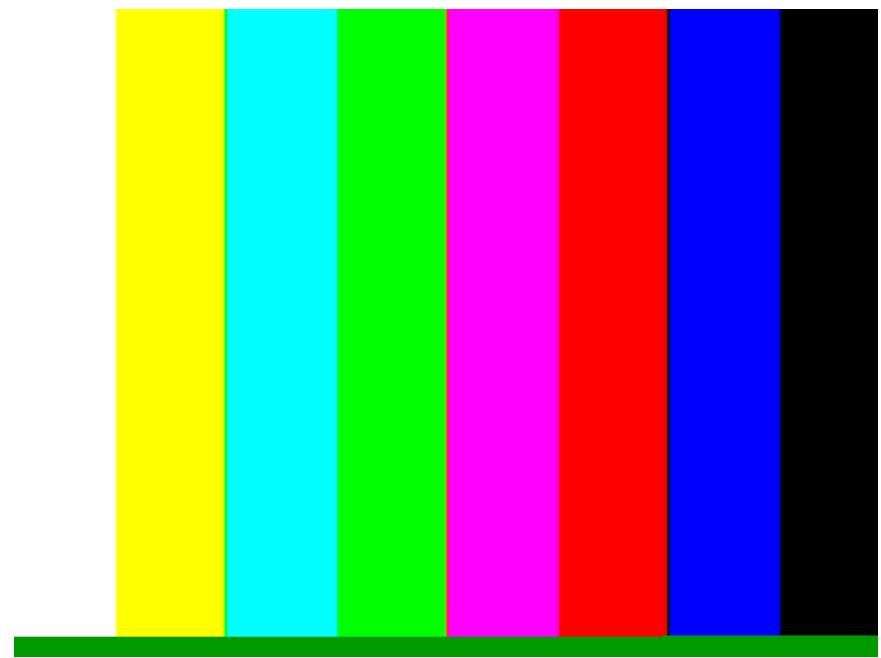


Color (RGB)

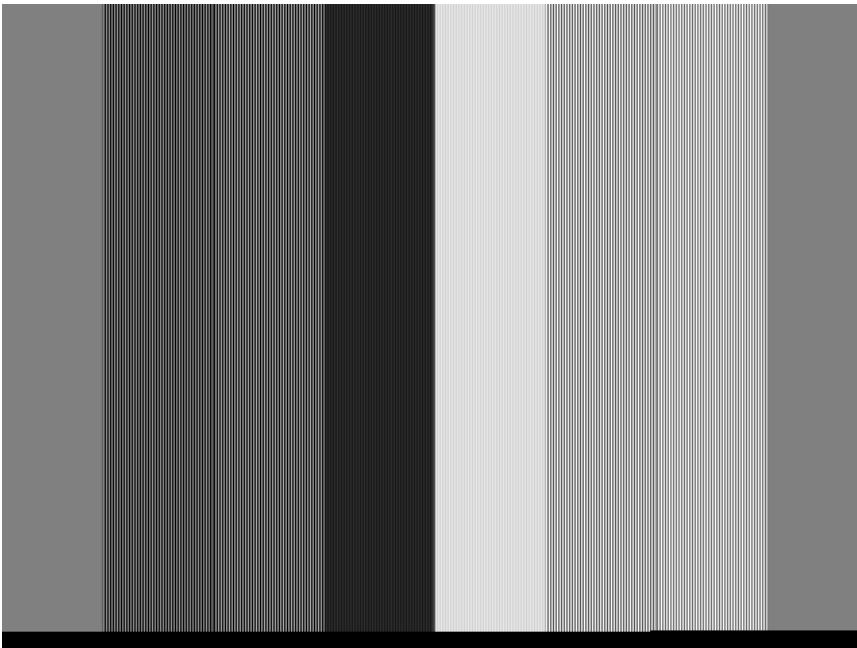
Chroma (UV)



Test pattern Luma (Y)



Test pattern Color (RGB)



Test pattern Chroma (UV)

VGA controller

This (optional) displays on a standard monitor the output video stream of the CV pipeline with (almost) zero latency.

The VGA controller (see the previous project) is fed with data read from the DDR memory using an appropriate READER and the same strategy based on the information provided by the Frame-Index.

In order to account for the critical timing requirements of the VGA video format, clocked at 25.175 MHz at VGA resolution, the READER includes an input buffer (mapped on BRAM).

The VGA module displays color or grayscale (default) images using the VGA output interface available in the Zedboard. Unfortunately this interface has only 4 bits per channel

Nevertheless, external VGA controllers could be easily plugged in the PMOD connectors

Receiver application (OpenCV)

The image stream is sent to a device (typically a PC) connected (point to point) to the ethernet port of the Zedboard.

The image stream is decoded and displayed using a simple OpenCV application running on the receiver device with minimal overhead exploiting threads.

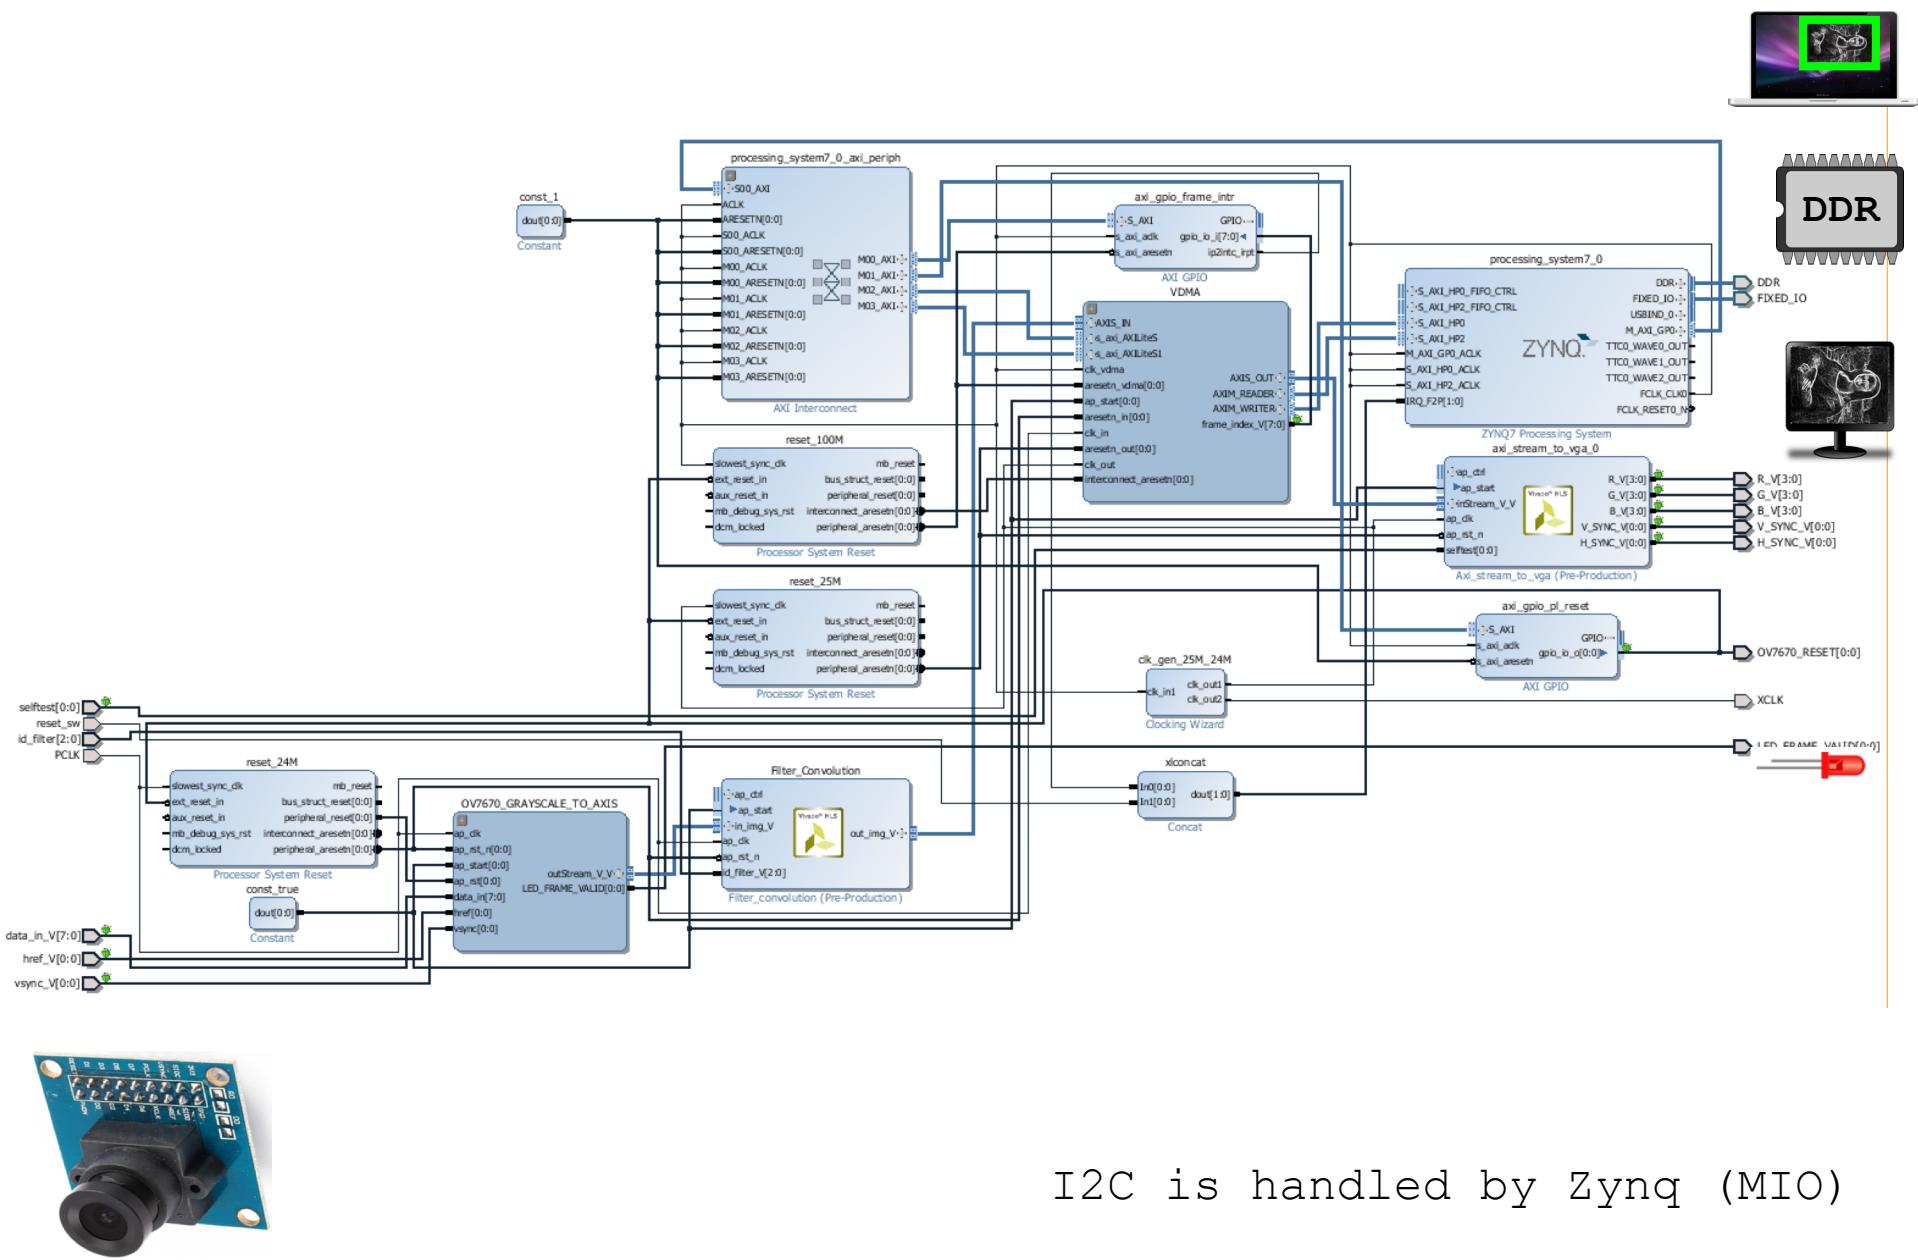
Camera settings:

IP	192.168.1.50
Netmask	255.255.255.0
Gateway	192.168.1.1

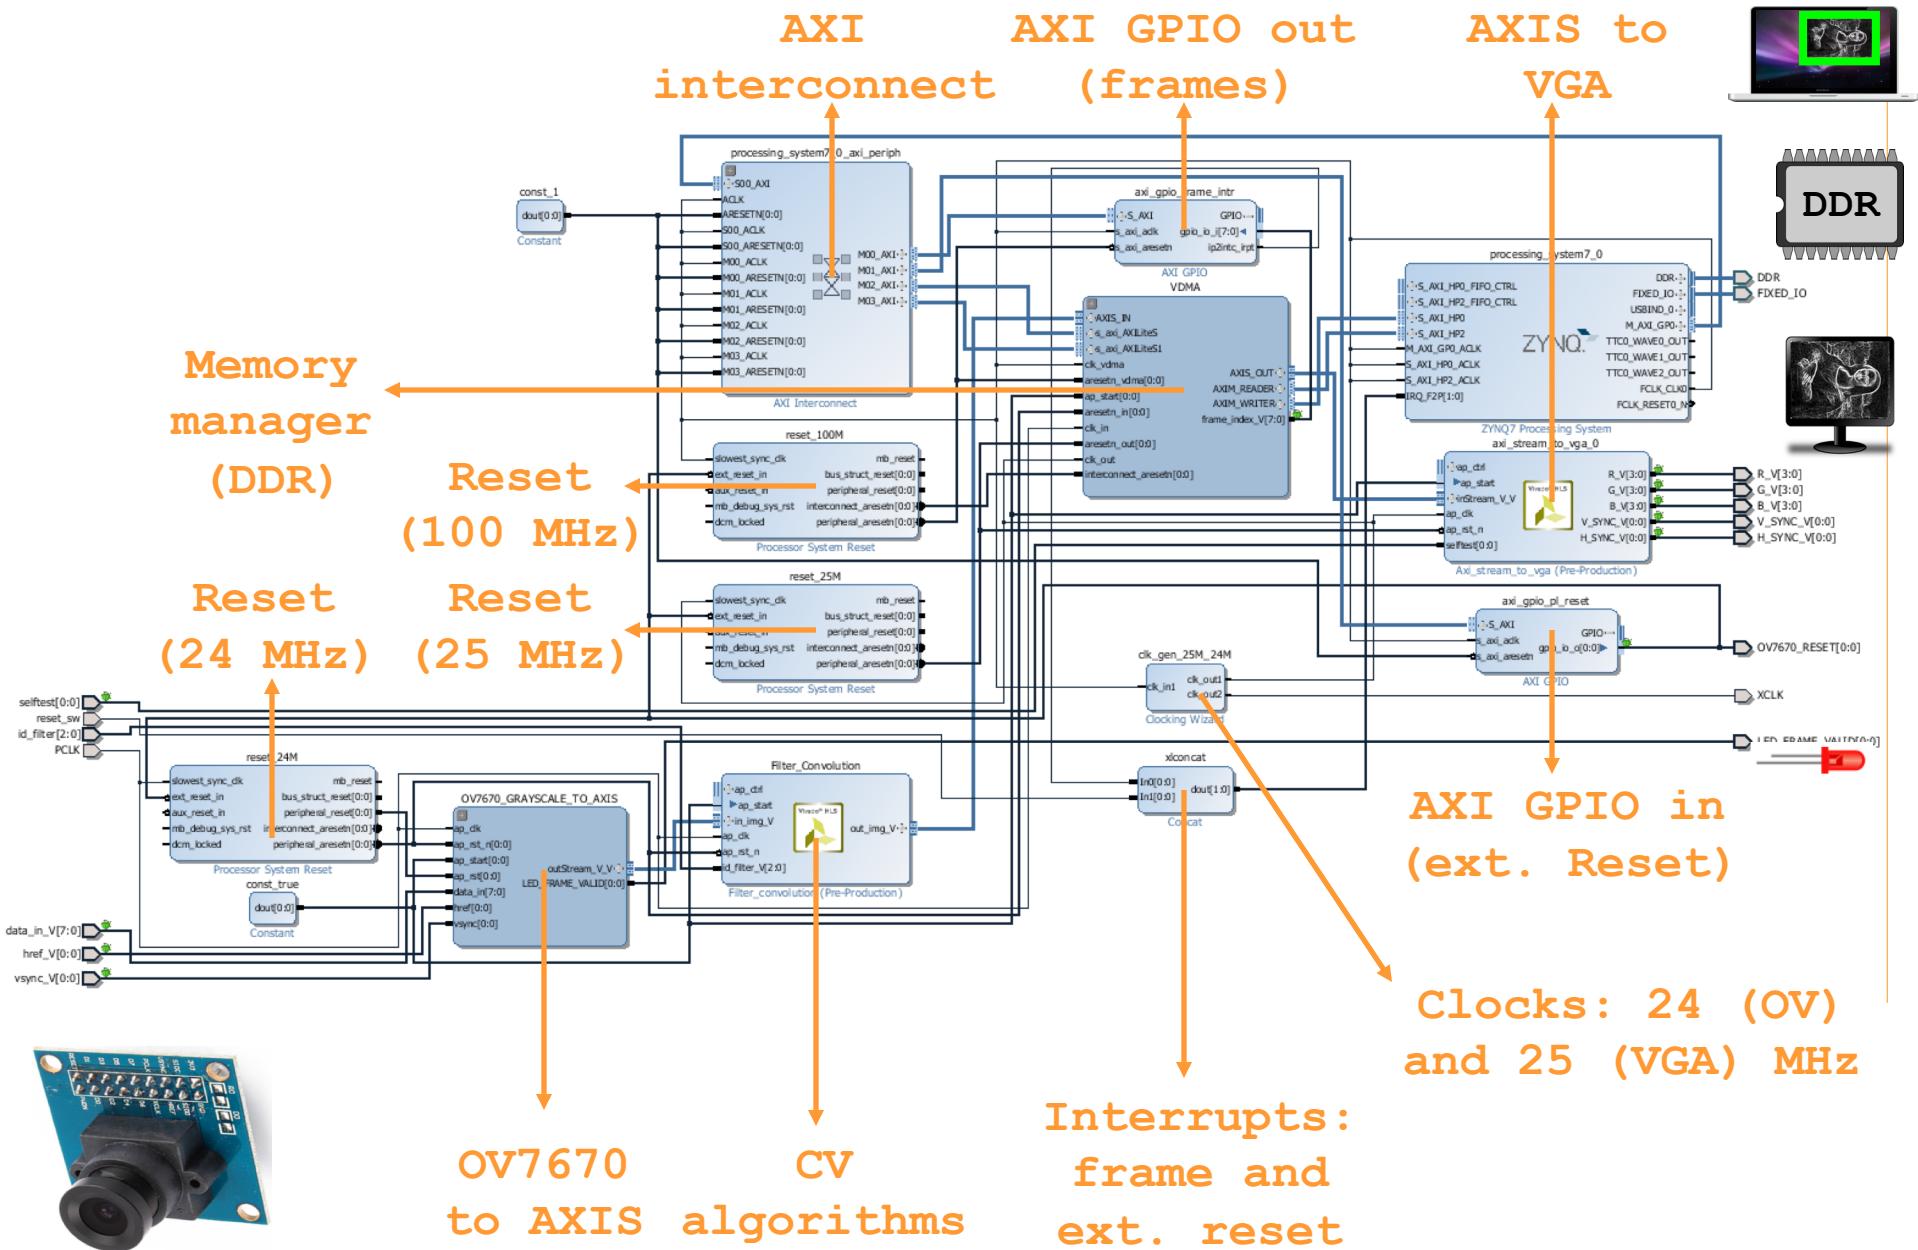
Host/PC settings:

IP	192.168.1.100
----	---------------

Overall design (Zynq PL) 1/2

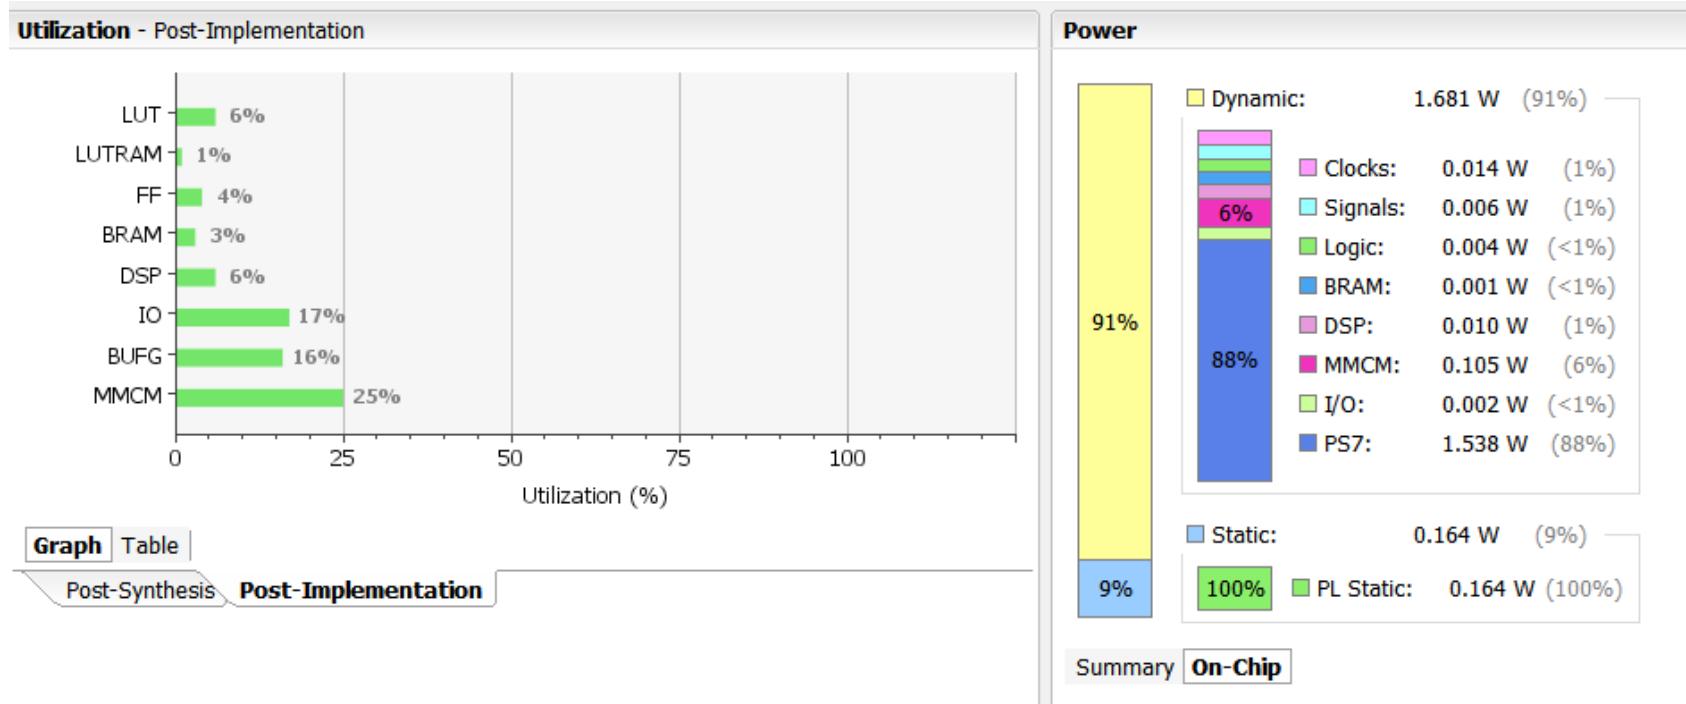


Overall design (Zynq PL) 2/2



The overall design utilizes a negligible portion of the Zynq 7020 device (6% LUT, 1% LUTRAM, 4% FF and 3% BRAM).

Power consumption < 2 W including ARM core (88%)



ARM: memory and devices 1/2

design_1_wrapper_hw_platform_1 Hardware Platform Specification

Design Information

Target FPGA Device: 7z020

Created With: Vivado 2016.2

Created On: Tue Nov 29 22:25:41 2016

Address Map for processor ps7_cortexa9_0

Cell	Base Addr	High Addr	Slave I/f	Mem/Reg
VDMA_axis_to_ddr_writer_0	0x43c00000	0x43c0ffff	s_axi_AXILiteS	REGISTER
VDMA_ddr_to_axis_reader_0	0x43c10000	0x43c1ffff	s_axi_AXILiteS	REGISTER
axi_gpio_frame_intr	0x41200000	0x4120ffff	S_AXI	REGISTER
axi_gpio_pl_reset	0x41210000	0x4121ffff	S_AXI	REGISTER
ps7_afi_0	0xf8008000	0xf8008fff		REGISTER
ps7_afi_1	0xf8009000	0xf8009fff		REGISTER
ps7_afi_2	0xf800a000	0xf800afff		REGISTER
ps7_afi_3	0xf800b000	0xf800bfff		REGISTER
ps7_coresight_comp_0	0xf8800000	0xf88fffff		REGISTER
ps7_ddr_0	0x00100000	0x1fffffff		MEMORY
ps7_ddrc_0	0xf8006000	0xf8006fff		REGISTER
ps7_dev_cfg_0	0xf8007000	0xf80070ff		REGISTER
ps7_dma_ns	0xf8004000	0xf8004fff		REGISTER
ps7_dma_s	0xf8003000	0xf8003fff		REGISTER
ps7_ethernet_0	0xe000b000	0xe000bfff		REGISTER
ps7_globaltimer_0	0xf8f00200	0xf8f002ff		REGISTER
ps7_gpio_0	0xe000a000	0xe000afff		REGISTER
ps7_gpv_0	0xf8900000	0xf89fffff		REGISTER

ARM: memory and devices 1/2

ps7_i2c_0	0xe0004000	0xe0004fff	REGISTER
ps7_intc_dist_0	0xf8f01000	0xf8f01fff	REGISTER
ps7_iop_bus_config_0	0xe0200000	0xe0200fff	REGISTER
ps7_l2cachec_0	0xf8f02000	0xf8f02fff	REGISTER
ps7_ocmc_0	0xf800c000	0xf800cff	REGISTER
ps7_pl310_0	0xf8f02000	0xf8f02fff	REGISTER
ps7_pmu_0	0xf8893000	0xf8893fff	REGISTER
ps7_qspi_0	0xe000d000	0xe000dff	REGISTER
ps7_qspi_linear_0	0xfc000000	0xfcfffff	MEMORY
ps7_ram_0	0x00000000	0x0002fff	MEMORY
ps7_ram_1	0xfffff0000	0xfffffdff	MEMORY
ps7_scuc_0	0xf8f00000	0xf8f000fc	REGISTER
ps7_scugic_0	0xf8f00100	0xf8f001ff	REGISTER
ps7_scutimer_0	0xf8f00600	0xf8f0061f	REGISTER
ps7_scuwdt_0	0xf8f00620	0xf8f006ff	REGISTER
ps7_sd_0	0xe0100000	0xe0100fff	REGISTER
ps7_slcr_0	0xf8000000	0xf8000fff	REGISTER
ps7_ttc_0	0xf8001000	0xf8001fff	REGISTER
ps7_uart_1	0xe0001000	0xe0001fff	REGISTER
ps7_usb_0	0xe0002000	0xe0002fff	REGISTER
ps7_xadc_0	0xf8007100	0xf8007120	REGISTER

Acceleration architectures and tools

