



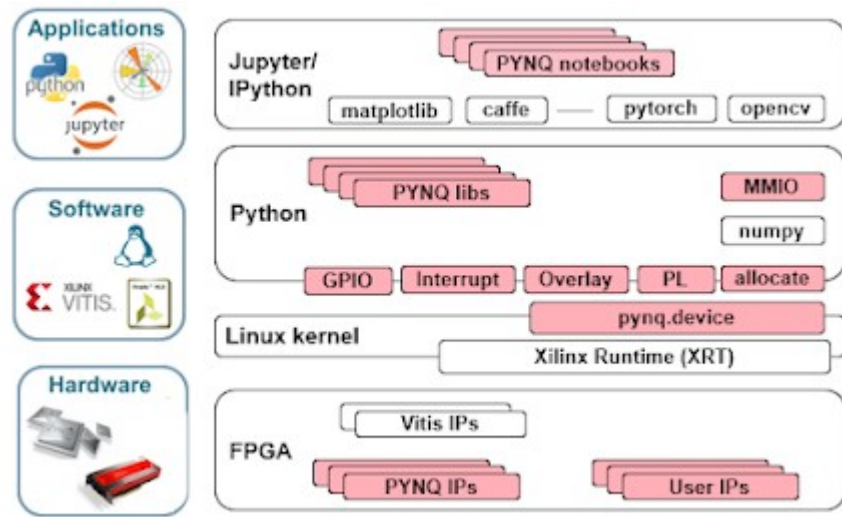
Introduction to PYNQ board

Digital Systems M, Module 2
Matteo Poggi, Università di Bologna

Introduction to PYNQ board

PYNQ is an open-source project from Xilinx that further brings FPGA programming to a higher level

Designed for Python language and libraries, the PYNQ board brings together the benefits of programmable logic and microprocessors to build more capable and exciting digital systems



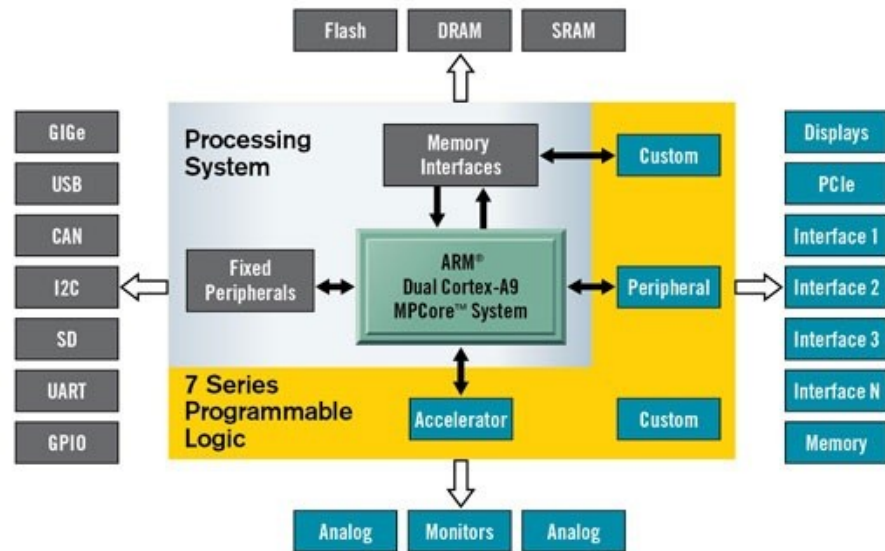
Introduction to PYNQ board

PYNQ boards are very similar to ZYNQ ones

They are made by a dual-core ARM Cortex-A9 processor (often referred to as **Processing System**, PS) and an FPGA (**Programmable Logic**, PL)

The PS can demand the execution of some functions to the PL by means of calls to specific hardware libraries, called **Overlays**

Overlays can both accelerate a software application or customize the hardware platform for a particular application



Introduction to PYNQ board

A software programmer can use an overlay in a similar way to a software library. It is particularly appealing when dealing with problems for which an FPGA can provide extremely good acceleration (e.g., image processing)

Overlays can be loaded to the FPGA dynamically just as software libraries and called when needed


PYNQ provides a Python interface to allow overlays in the **PL** to be controlled from Python running in the **PS**. While overlays are written by hardware specialists, software developers can call them via Python transparently and **without requiring** deep hardware design knowledge



Getting started

Introduction to PYNQ board


The PYNQ board can be easily configured and accessed from a standard PC and a web browser. Let's see how



Introduction to PYNQ board

The board comes with an SD card (3) containing a custom Linux distribution (actually, PYNQ v2.5)

- It can be powered with a regular power supply (6) or via **USB** (4). The modality is selected by a jumper (1)
- To access the board, we can either connect it (6) to a router or to our own PC
- A switch (2) allows to power it up (wait for all leds to light up)



Introduction to PYNQ board

The standard IP assigned to the board is 192.168.2.99

If the board is connected via ethernet to a router, let's browse to

<http://pynq:9090>

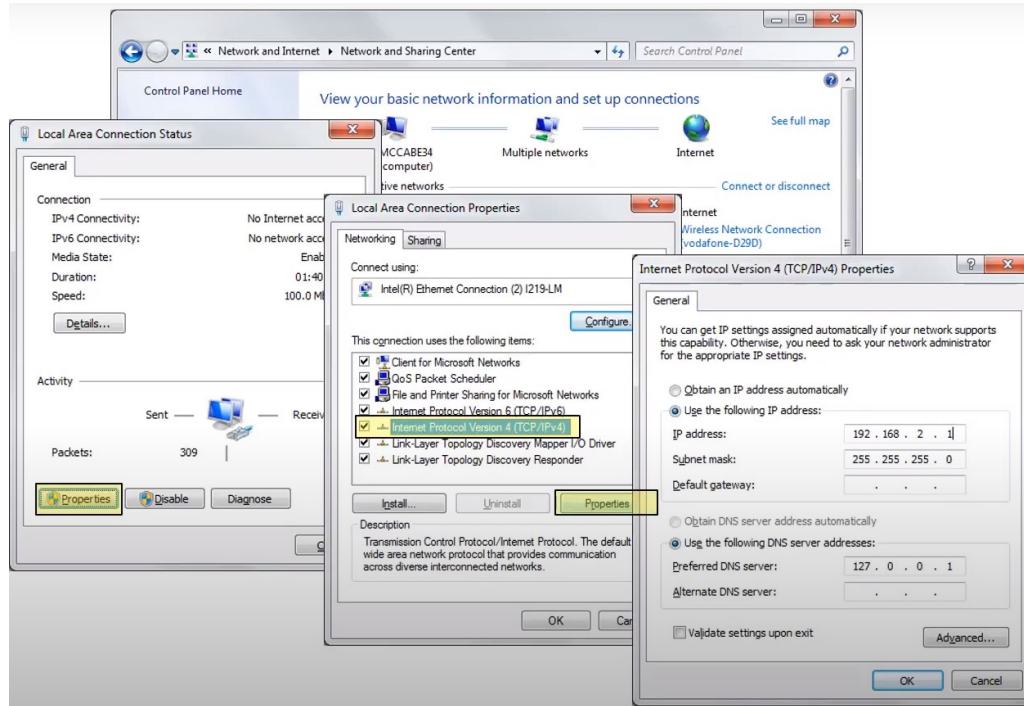
Otherwise, a **static IP** in the same subnetwork (192.168.2.x) is required for our laptop (see next slide). Then, let's browse to

<http://192.168.2.99:9090>

In both cases, the password to access the board is **xilinx**

Introduction to PYNQ board

Windows:

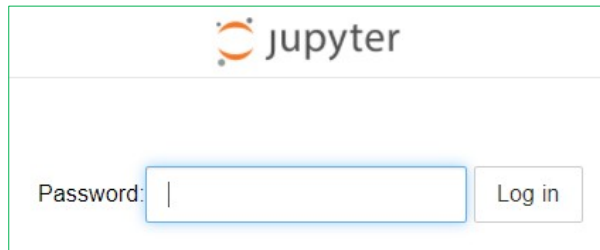


Linux:

```
ifconfig eth0 192.168.2.1 netmask 255.255.255.0 up
```

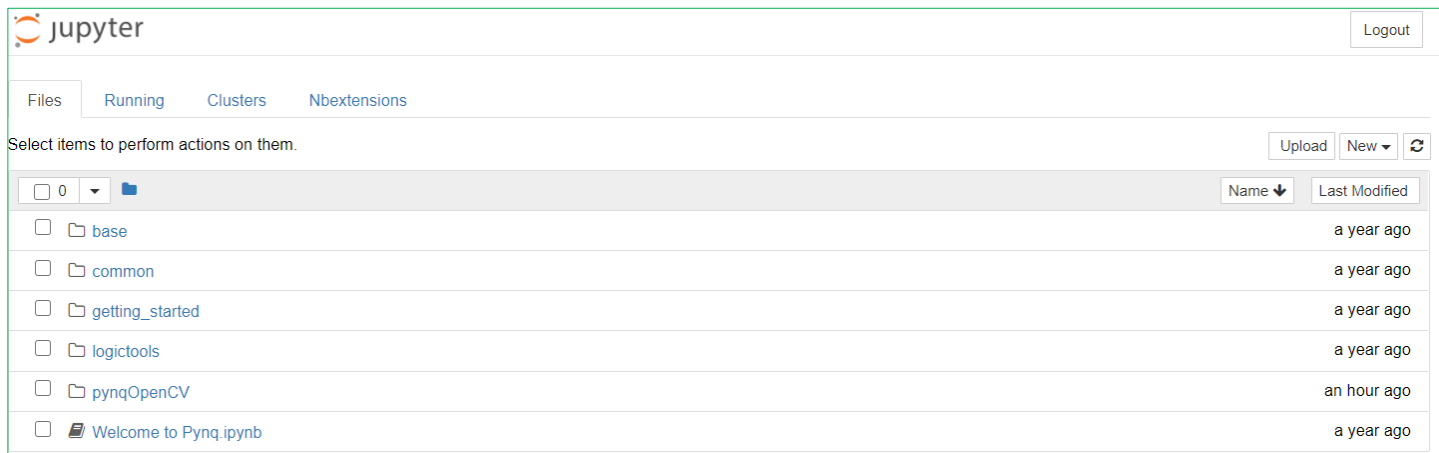
Introduction to PYNQ board

By browsing to the previous links, we should reach this page...



The image shows the Jupyter login interface. At the top is the Jupyter logo. Below it, there is a label "Password:" followed by a text input field and a "Log in" button.

... and then be ready to play!



The image shows the Jupyter file browser interface. At the top is the Jupyter logo and a "Logout" button. Below the logo are tabs for "Files", "Running", "Clusters", and "Nbextensions". The "Files" tab is active. Below the tabs, there is a message "Select items to perform actions on them." and buttons for "Upload", "New", and a refresh icon. Below this is a table of files and folders. The table has two columns: "Name" and "Last Modified". The first row shows a folder icon, a dropdown menu with "0", and a folder icon. The subsequent rows show folders named "base", "common", "getting_started", "logictools", "pynqOpenCV", and a file named "Welcome to Pynq.ipynb". The "Last Modified" column shows the time since each item was last modified.

Name	Last Modified
0	
base	a year ago
common	a year ago
getting_started	a year ago
logictools	a year ago
pynqOpenCV	an hour ago
Welcome to Pynq.ipynb	a year ago



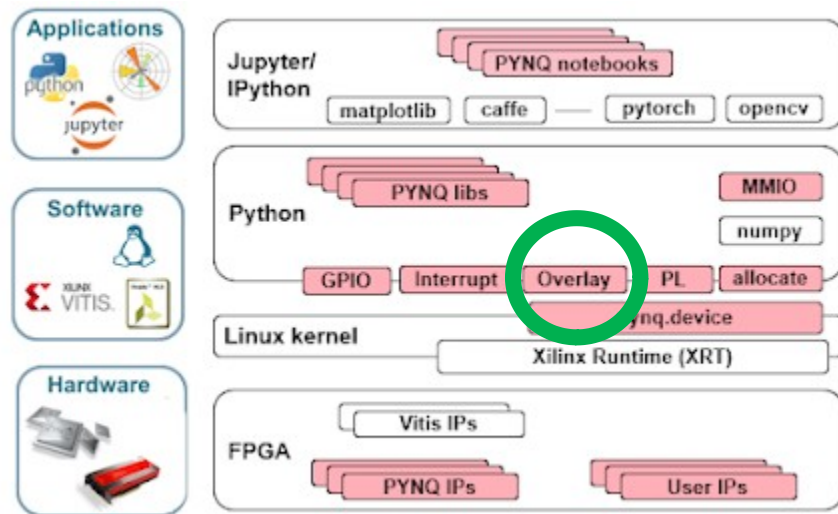
Overlays

Introduction to PYNQ board

Overlays allows for easy and transparent use of low-level IPs at Python level

An overlay **loads** a bitstream and give access to its functionalities

This is done by means of **drivers** objects



Introduction to PYNQ board

Example: call a custom IP, **scalar_add**, written in HLS to be run on Python
(more examples at https://github.com/PeterOgden/overlay_tutorial)

Let's consider a simple HLS function that returns the sum of two 32-bit integers

```
void add(int a, int b, int& c) {  
    #pragma HLS INTERFACE ap_ctrl_none port=return  
    #pragma HLS INTERFACE s_axilite port=a  
    #pragma HLS INTERFACE s_axilite port=b  
    #pragma HLS INTERFACE s_axilite port=c  
  
    c = a + b;  
}
```

After designing this IP, we can load it on Python side by means of an Overlay

Introduction to PYNQ board

From the **pynq** package (pre-installed on PYNQ SD image), we can import Overlay and create an overlay object from the bitstream containing the custom IP, then we can retrieve the specific IP

```
In [1]: from pynq import Overlay  
  
        overlay = Overlay('/home/xilinx/tutorial_1.bit')
```

```
In [3]: add_ip = overlay.scalar_add  
        help(add_ip)
```

The **help** function print the IP object extracted by the overlay

Introduction to PYNQ board

The IP core is loaded by a DefaultIP object, a driver exposing **read** and **write** APIs

According to the documentation generated by HLS, to use the IP core we need to write the two arguments to offset 0x10 and 0x18, then read the result back from 0x20

```
In [4]: add_ip.write(0x10, 4)
        add_ip.write(0x18, 5)
        add_ip.read(0x20)
```

```
Out[4]: 9
```

```
class DefaultIP(builtins.object)
    Driver for an IP without a more specific driver

    This driver wraps an MMIO device and provides a base class
    for more specific drivers written later. It also provides
    access to GPIO outputs and interrupts inputs via attributes. More specific
    drivers should inherit from 'DefaultIP' and include a
    'bindto' entry containing all of the IP that the driver
    should bind to. Subclasses meeting these requirements will
    automatically be registered.

    Attributes
    -----
    mmio : pynq.MMIO
        Underlying MMIO driver for the device
    _interrupts : dict
        Subset of the PL.interrupt_pins related to this IP
    _gpio : dict
        Subset of the PL.gpio_dict related to this IP

    Methods defined here:

    __init__(self, description)
        Initialize self. See help(type(self)) for accurate signature.

    read(self, offset=0)
        Read from the MMIO device

        Parameters
        -----
        offset : int
            Address to read

    write(self, offset, value)
        Write to the MMIO device

        Parameters
        -----
        offset : int
            Address to write to
        value : int or bytes
            Data to write

    -----
    Data descriptors defined here:

    __dict__
        dictionary for instance variables (if defined)

    __weakref__
        list of weak references to the object (if defined)
```


Introduction to PYNQ board

The default driver is not user-friendly. A better solution would be a custom IP driver exposing a single **add** function. We can do this by extending the DefaultIP class

This driver wraps an MMIO device and provides a base class for more specific drivers written later. It also provides access to GPIO outputs and interrupts inputs via attributes. More specific drivers should inherit from `DefaultIP` and include a `bindto` entry containing all of the IP that the driver should bind to. Subclasses meeting these requirements will automatically be registered.

```
In [5]: from pynq import DefaultIP

class AddDriver(DefaultIP):
    def __init__(self, description):
        super().__init__(description=description)

    bindto = ['xilinx.com:hls:add:1.0']

    def add(self, a, b):
        self.write(0x10, a)
        self.write(0x18, b)
        return self.read(0x20)
```

Introduction to PYNQ board

Calling **help** on the overlay using the DefaultIP driver:

```
class Overlay(pynq.pl.Bitstream)
|   Default documentation for overlay /home/xilinx/tutorial_1.bit. The following
|   attributes are available on this overlay:
|
|   IP Blocks
|   -----
|   scalar_add          : pynq.overlay.DefaultIP
```

Calling **help** on the overlay using a custom driver:

```
class Overlay(pynq.pl.Bitstream)
|   Default documentation for overlay /home/xilinx/tutorial_1.bit. The following
|   attributes are available on this overlay:
|
|   IP Blocks
|   -----
|   scalar_add          : __main__.AddDriver
```

Introduction to PYNQ board

Now, we can access to the IP core and use a single **add** API, made available by our custom driver

```
In [7]: overlay.scalar_add.add(15,20)
```

```
Out[7]: 35
```

With the overlays mechanism we can easily reuse IPs, implement IP Hierarchies, custom overlays and more

Introduction to PYNQ board

The pynq library includes a number of drivers as part of the **pynq.lib** package.

These include

- AXI GPIO
- AXI DMA (simple mode only)
- AXI VDMA
- AXI Interrupt Controller (internal use)
- Pynq-Z1 Audio IP
- Pynq-Z1 HDMI IP
- Color convert IP
- Pixel format conversion
- HDMI input and output frontends
- Pynq Microblaze program loading
- ...



PYNQ's Hello LED

Introduction to PYNQ board

All the peripherals on the PYNQ board are connected to PL. This means controllers must be implemented in an overlay before these peripherals can be used. The base overlay contains controllers for all the peripherals.

(Full tutorial at https://pynq.readthedocs.io/en/v1.3/5_programming_onboard.html)

By loading the base overlay, we can import specific classes for each peripherals from the **pynq.overlays.base** module

```
import time
from pynq.overlays.base import BaseOverlay
base = BaseOverlay("base.bit")
```

Introduction to PYNQ board

We can easily instantiate objects and manipulate the corresponding peripherals. For instance, we can access to a single LED and turn it **on/off**

```
led0 = base.leds[0] #Corresponds to LED LD0  
led1 = base.leds[1] #Corresponds to LED LD1  
led2 = base.leds[2] #Corresponds to LED LD2  
led3 = base.leds[3] #Corresponds to LED LD3
```

```
led0.on()
```

```
led0.off()
```


Introduction to PYNQ board

We can also change the current status of a LED with the **toggle()** method

```
for i in range(20):  
    led0.toggle()  
    time.sleep(0.1)
```

Introduction to PYNQ board

We can use **switches** or **buttons** available on the board to turn **on/off** the LEDs

We first get access to LEDs/switches/buttons available

```
MAX_LEDS = 4
MAX_SWITCHES = 2
MAX_BUTTONS = 4

leds = [0] * MAX_LEDS
switches = [0] * MAX_SWITCHES
buttons = [0] * MAX_BUTTONS

for i in range(MAX_LEDS):
    leds[i] = base.leds[i]
for i in range(MAX_SWITCHES):
    switches[i] = base.switches[i]
for i in range(MAX_BUTTONS):
    buttons[i] = base.buttons[i]
```

Introduction to PYNQ board

Let's first reset the LEDs status by implementing a simple **clear_LEDs** function

```
In [7]: # Helper function to clear LEDs
def clear_LEDs(LED_nos=list(range(MAX_LEDS))):
    """Clear LEDs LD3-0 or the LEDs whose numbers appear in the list"""
    for i in LED_nos:
        leds[i].off()

clear_LEDs()
```

Then, we can implement the following policy:

- If Switch(0) is on, LED(2) and LED(0) will be on
- If Switch(1) is on, LED(3) and LED(1) will be on

```
In [8]: clear_LEDs()

for i in range(MAX_LEDS):
    if switches[i%2].read():
        leds[i].on()
    else:
        leds[i].off()
```

Introduction to PYNQ board

Let's try with another policy:

- Toggle LED(i) if Button(i) is pressed while Switch(0) is switched on

```
In [9]: import time

        clear_LEDs()

        while switches[0].read():
            for i in range(MAX_LEDS):
                if buttons[i].read():
                    leds[i].toggle()
                    time.sleep(.1)

        clear_LEDs()
```



Example: HDMI stream

Introduction to PYNQ board

Let's start with a simple pipeline to acquire an input video stream from a and show it into a monitor over HDMI ports

Source:


http://pynq:9090/notebooks/getting_started/5_base_overlay_video.ipynb

(other examples – not dealing with image processing – are at

http://pynq:9090/notebooks/getting_started/)

Introduction to PYNQ board

The PYNQ board is equipped with 2x HDMI ports, respectively dedicated to handle **input** and **output** streams



Introduction to PYNQ board

The PYNQ board provides drivers to handle the HDMI ports

We can simply load the `pynq.lib.video` module

```
In [1]: from pynq.overlays.base import BaseOverlay
        from pynq.lib.video import *

        base = BaseOverlay("base.bit")
        hdmi_in = base.video.hdmi_in
        hdmi_out = base.video.hdmi_out
```

Introduction to PYNQ board

We can easily configure the in/out HDMI handlers with few lines:

```
In [2]: hdmi_in.configure()  
        hdmi_out.configure(hdmi_in.mode)  
  
        hdmi_in.start()  
        hdmi_out.start()
```

At first, we can just replicate the input stream on the output HDMI port with the **tie** method

```
In [3]: hdmi_in.tie(hdmi_out)
```

Introduction to PYNQ board

Equivalently, we can use the **readframe** and **writeframe** methods, respectively on the input and outputports, to grab a frame and forward it towards the output HDMI port

```
In [4]: import time

numframes = 600
start = time.time()

for _ in range(numframes):
    f = hdmi_in.readframe()
    hdmi_out.writeframe(f)

end = time.time()
print("Frames per second: " + str(numframes / (end - start)))
```

Frames per second: 34.12674016266714

Introduction to PYNQ board

This allows to **process** the frames before showing them on screen

```
In [5]: import cv2
import numpy as np

numframes = 10
grayscale = np.ndarray(shape=(hdmi_in.mode.height, hdmi_in.mode.width),
                        dtype=np.uint8)
result = np.ndarray(shape=(hdmi_in.mode.height, hdmi_in.mode.width),
                     dtype=np.uint8)

start = time.time()

for _ in range(numframes):
    inframe = hdmi_in.readframe()
    cv2.cvtColor(inframe, cv2.COLOR_BGR2GRAY, dst=grayscale)
    inframe.freebuffer()
    cv2.Laplacian(grayscale, cv2.CV_8U, dst=result)

    outframe = hdmi_out.newframe()
    cv2.cvtColor(result, cv2.COLOR_GRAY2BGR, dst=outframe)
    hdmi_out.writeframe(outframe)

end = time.time()
print("Frames per second: " + str(numframes / (end - start)))
```

Frames per second: 3.5044367097311837

Introduction to PYNQ board

Don't forget to close the stream interfaces once we are done

```
In [6]: hdmi_out.close()  
        hdmi_in.close()
```

More options:

- Cachable frames
 - Frames are cached to speed-up software processing (ARM). This slows HDMI processing and can be disabled with *hdmi_in.cacheable_frames = False*
- Gray-scale
 - We can delegate conversion to HW with *hdmi_in.configure(PIXEL_GRAY)*
- Other color-spaces
 - *hdmi_in.colospace = COLOR_IN_YCBCR*



Example: HW
accelerated filtering

Introduction to PYNQ board

In this example, the previous acquisition/visualization pipeline is extended with **hardware accelerated** image processing

Source: <https://github.com/Xilinx/PYNQ-ComputerVision.git>



To run it, we first need to upgrade Cython and install the PYNQ-ComputerVision package

```
pip3 install --upgrade Cython
```

```
pip3 install git+https://github.com/Xilinx/PYNQ-ComputerVision.git
```


Introduction to PYNQ board

We now have a new folder, **pynqOpenCV**, containing our examples

<input type="checkbox"/>  base	a year ago
<input type="checkbox"/>  common	2 hours ago
<input type="checkbox"/>  getting_started	22 minutes ago
<input type="checkbox"/>  logictools	2 hours ago
<input type="checkbox"/>  pynqOpenCV	2 days ago
<input type="checkbox"/>  Welcome to Pynq.ipynb	2 months ago

Source: <http://pynq:9090/notebooks/pynqOpenCV/filter2d.ipynb>

Introduction to PYNQ board

As usual, we first import the Overlay module. This time, we load a **custom** bitstream containing the 2D filter module

We also load the **xlnk memory manager**, a python class in charge of allocating **continuous memory** that is more efficient for PL IPs. It is compatible with numpy arrays

```
In [1]: # Load filter2D + dilate overlay
from pynq import Overlay
bareHDMI = Overlay("/usr/local/lib/python3.6/dist-packages/"
                  "pynq_cv/overlays/xv2Filter2DDilate.bit")
import pynq_cv.overlays.xv2Filter2DDilate as xv2

# Load xlnk memory mangager
from pynq import Xlnk
Xlnk.set_allocator_library("/usr/local/lib/python3.6/dist-packages/"
                          "pynq_cv/overlays/xv2Filter2DDilate.so")

mem_manager = Xlnk()

hdmi_in = bareHDMI.video.hdmi_in
hdmi_out = bareHDMI.video.hdmi_out
```

Introduction to PYNQ board

After configuring the HDMI input and output ports, we can run some 2D filtering on software side (nothing new here)

```
In [4]: import numpy as np
import time
import cv2

#Sobel Vertical filter
kernelF = np.array([[1.0,0.0,-1.0],[2.0,0.0,-2.0],[1.0,0.0,-1.0]],np.float32)

numframes = 20

start = time.time()
for _ in range(numframes):
    inframe = hdmi_in.readframe()
    outframe = hdmi_out.newframe()
    cv2.filter2D(inframe, -1, kernelF, dst=outframe)
    inframe.freebuffer()
    hdmi_out.writeframe(outframe)
end = time.time()
print("Frames per second: " + str(numframes / (end - start)))

Frames per second: 3.3585366809661044
```

Introduction to PYNQ board

We might define a set of filters and **interactively** choose the one we want to apply using **ipywidgets**

```
In [7]: from ipywidgets import interact, interactive, fixed, interact_manual
        from ipywidgets import IntSlider, FloatSlider
        import ipywidgets as widgets

        #Sobel Vertical filter
        kernel_g = np.array([[1.0,0.0,-1.0],[2.0,0.0,-2.0],[1.0,0.0,-1.0]],np.float32)

        def setKernelAndFilter3x3(kernelName):
            global kernel_g

            kernel_g = {
                'Laplacian high-pass': np.array([[0.0,1.0,0.0],[1.0,-4.0,1.0],
                                                  [0.0,1.0,0.0]],np.float32),
                'Gaussian high-pass': np.array([[ -0.0625,-0.125,-0.0625],
                                                  [-0.125,0.75,-0.125],
                                                  [-0.0625,-0.125,-0.0625]],np.float32),
                'Average blur': np.ones((3,3),np.float32)/9.0,
                'Gaussian blur': np.array([[0.0625,0.125,0.0625],
                                                  [0.125,0.25,0.125],
                                                  [0.0625,0.125,0.0625]],np.float32),
                'Sobel ver': np.array([[1.0,0.0,-1.0],[2.0,0.0,-2.0],
                                       [1.0,0.0,-1.0]],np.float32),
                'Sobel hor': np.array([[1.0,2.0,1.0],[0.0,0.0,0.0],
                                       [-1.0,-2.0,-1.0]],np.float32)
            }.get(kernelName, np.ones((3,3),np.float32)/9.0)

        interact(setKernelAndFilter3x3, kernelName
                 = ['Sobel ver','Sobel hor','Laplacian high-pass','Gaussian high-pass','Average blur',
                   'Gaussian blur',]);
```

Introduction to PYNQ board

Let's now call the 2D filter (implemented as IP core) through the Overlay

In order to allow kernel redefinition on the fly, subsequent function calls are run as **threads** thanks to the **Thread** class

```
In [8]: import numpy as np
import cv2
from threading import Thread

def loop_hw2_app():
    global kernel_g

    numframes = 600

    start=time.time()
    for _ in range(numframes):
        outframe = hdmi_out.newframe()
        inframe = hdmi_in.readframe()
        xv2.filter2D(inframe, -1, kernel_g, dst=outframe, borderType=cv2.BORDER_CONSTANT)
        hdmi_out.writeframe(outframe)
        inframe.freebuffer()
    end=time.time()
    print("Frames per second: " + str(numframes / (end - start)))

t = Thread(target=loop_hw2_app)
t.start()

Frames per second: 43.98824715869524
```



Existing IP cores and open-source projects

Introduction to PYNQ board

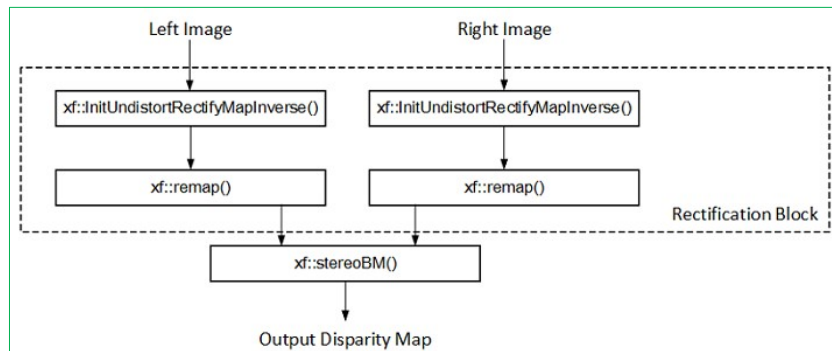
Xilinx provides plenty of IP cores implementing OpenCV functions!

More than 60 kernels at

<https://github.com/Xilinx/xfopencv>

https://github.com/Xilinx/Vitis_Libraries/tree/master/vision

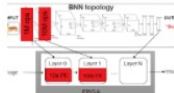
e.g. stereo matching pipeline



Introduction to PYNQ board

Finally, a list of PYNQ open-source projects is available at <http://www.pynq.io/examples>

Machine Learning on Xilinx FPGAs with FINN



PYNQ has been widely used for machine learning research and prototyping.

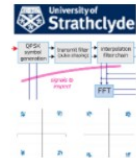
FINN, an experimental framework from Xilinx Research Labs to explore deep neural network inference on FPGAs. It specifically targets quantized neural networks, with emphasis on generating dataflow-style architectures customized for each network.

FINN makes extensive use of PYNQ as a prototyping platform.
For more information see pynq.io/ml

PYNQ community projects


PYNQ RFSoC

University Strathclyde
QPSK demo on ZCU111




PYNQ-PRIO

BYU
Partial reconfiguration Input/Output




PYNQ Hello World

Hardware accelerated image resizer example




RISC-V on PYNQ

UCSD



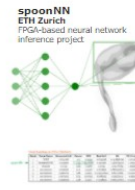
Extended Kalman filter

University Sydney



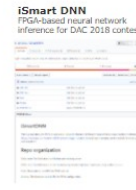
spoonNN

ETH Zurich
FPGA-based neural network inference project




iSmart DNN

FPGA-based neural network inference for DAC 2018 contest




TCIIF

FPGA-based neural network inference for DAC 2018 contest




cv2PYNQ

ETH Leuven
Accelerated OpenCV image filtering library




Video processing

KU Leuven
Hardware accelerated video




ZipML-PYNQ

ETH Zurich
Hardware accelerated compression




PYNQ bot

ET Tallgate
Control of robotic car from PYNQ




PYNQ LED cube

Fudan University, Xilinx China
Controlling an LED cube from PYNQ




SPYN

Xilinx TSM, Trezz electronics
Industrial motor control




PYNQ networking

Xilinx labs
On-chip with network analysis capability




QNN

Xilinx labs
Quantized neural network




LTSM

TU Kaiserslautern
Quantized LSTM on PYNQ




BNN

University Sydney, Xilinx labs
Binarized neural network




Video filters with PR

BYU
Video filtering with partial reconfiguration



PYNQ computer vision

Xilinx labs
Build a vision processing pipeline from v4OpenCV



GZip on PYNQ

University Bucharest
GZip compression with DFLATE, compatible data, and fixed Huffman coding



SPynq

NTUA
Apache Spark on PYNQ



FIR filter example

CU Boulder
FIR filter example



CNN on PYNQ

Imperial College London



References

- PYNQ: Python productivity - <http://www.pynq.io/>
- Python productivity for Zynq (Pynq) - <https://pynq.readthedocs.io/en/latest/index.html>
- PYNQ Computer Vision - <https://github.com/Xilinx/PYNQ-ComputerVision>